


Introduction

We model communication:



Examples: optical signals, electrical telegraph, SMS (compression), postcodes, CDs (error correction), zip/gz files (compression).

Given a source and a channel, modelled probabilistically, the basic problem is to design an encoder and decoder to transmit messages economically (noiseless coding; compression) and reliably (noisy coding).

Examples:

- Noiseless coding: Morse code: common letters are assigned shorter code-words, e.g. $A \mapsto \bullet-$, $E \mapsto \bullet$, $Q \mapsto --\bullet-$, $S \mapsto \bullet\bullet\bullet$, $O \mapsto --$, $Z \mapsto --\bullet\bullet$. Noiseless coding is adapted to source.
- Noisy coding: Every book has an ISBN $a_1, a_2, \dots, a_9, a_{10}$, $a_i \in \{0, 1, \dots, 9\}$ for $1 \leq i \leq 9$ and $a_{10} \in \{0, 1, \dots, 9, X\}$ with $\sum_{j=1}^{10} ja_j \equiv 0 \pmod{11}$. This detects common errors - e.g one incorrect digit, transposition of two digits. Noisy coding is adapted to the channel.

Plan:

- (I) Noiseless coding - entropy
- (II) Error correcting codes - noisy channels
- (III) Information theory - Shannon's theorems
- (IV) Examples of codes
- (V) Cryptography

Books: [GP], [W], [CT], [TW], Buchmann, Körner. Online notes: Carne, Körner.

Basic Definitions

Definition (Communication channel). A *communication channel* accepts symbols from a alphabet $\mathcal{A} = \{a_1, \dots, a_r\}$ and it outputs symbols from alphabet $\mathcal{B} = \{b_1, \dots, b_s\}$. Channel modelled by the probabilities $\mathbb{P}(y_1 \dots y_n \text{ received} | x_1 \dots x_n \text{ sent})$. A *discrete memoryless channel* (DMC) is a channel with

$$p_{ij} = \mathbb{P}(b_j \text{ received} | a_i \text{ sent})$$

the same for each channel use and independent of all past and future uses. The channel matrix is $P = (b_{ij})$, a $r \times s$ stochastic matrix.

Definition (Binary symmetric channel). The *binary symmetric channel* (BSC) with error probability $p \in [0, 1)$ from $\mathcal{A} = \mathcal{B} = \{0, 1\}$. The channel matrix is

$$\begin{pmatrix} 1-p & p \\ p & 1-p \end{pmatrix}.$$

A symbol is transmitted correctly with probability $1 - p$. Usually assume $p < 1/2$.

The *binary erasure channel* (BEC) has $\mathcal{A} = \{0, 1\}$, $\mathcal{B} = \{0, 1, *\}$. The channel matrix is

$$\begin{pmatrix} 1-p & 0 & p \\ 0 & 1-p & p \end{pmatrix}.$$

So $p = \mathbb{P}(\text{symbol can't be read})$.

Definition. We model n uses of a channel by the n th extension, with input alphabet \mathcal{A}^n and output alphabet \mathcal{B}^n . A *code* C of length n is a function $\mathcal{M} \rightarrow \mathcal{A}^n$ where \mathcal{M} is the set of possible messages. Implicitly we also have a decoding rule $\mathcal{B}^n \rightarrow \mathcal{M}$. The *size* of C is $m = |\mathcal{M}|$. The *information rate* is $\rho(C) = \frac{1}{n} \log_2 m$. The *error rate* is $\hat{e}(C) = \max_{x \in \mathcal{M}} \mathbb{P}(\text{error} | x \text{ sent})$.

Remark. For the remainder of the course we write \log instead of \log_2 .

Definition. A channel can *transmit reliably at rate* R if there exists $(C_n)_{n=1}^\infty$ with each C_n a code of length n such that

$$\lim_{n \rightarrow \infty} \rho(C_n) = R \text{ \& } \lim_{n \rightarrow \infty} \hat{e}(C_n) = 0.$$

The *capacity* is the supremum of all reliable transmission rates. We'll see in Chapter 9 that a BSC with error probability $p < 1/2$ has non-zero capacity.

1 Noiseless coding

1.1 Prefix-free codes

For an alphabet \mathcal{A} , $|\mathcal{A}| < \infty$, let $\mathcal{A}^* = \bigcup_{n \geq 0} \mathcal{A}^n$, the set of all finite strings from \mathcal{A} . The *concatenation* of strings $x = x_1 \dots x_r$ and $y = y_1 \dots y_s$ is $xy = x_1 \dots x_r y_1 \dots y_s$.

Definition. Let \mathcal{A}, \mathcal{B} be alphabets. A code is a function $c : \mathcal{A} \rightarrow \mathcal{B}^*$. The strings $c(a)$ for $a \in \mathcal{A}$ are called *codewords* or *words* (CWS).

Example 1.1 (Greek fire code). $\mathcal{A} = \{\alpha, \beta, \dots, \omega\}$ (greek alphabet), $\mathcal{B} = \{1, 2, 3, 4, 5\}$, $c : \alpha \mapsto 11, \beta \mapsto 12, \dots, \psi \mapsto 53, \omega \mapsto 54$. xy means hold up x torches and another y torches nearby.

Example 1.2. \mathcal{A} = words in a dictionary, $\mathcal{B} = \{A, B, \dots, Z, \omega\}$. $c : \mathcal{A} \rightarrow \mathcal{B}$ splits the word and follows with a space. Send message $x_1 \dots x_n \in \mathcal{A}^*$ as $c(x_1) \dots c(x_n) \in \mathcal{B}^*$. So c extends to a function $c^* : \mathcal{A}^* \rightarrow \mathcal{B}^*$.

Definition. c is said to be *decipherable* if the induced map c^* (as in the previous example) is injective. In other words, each string from \mathcal{B} corresponds to at most one message.

Clearly if c is decipherable, it is necessary for c to be injective. However it is not sufficient:

Example 1.3. $\mathcal{A} = \{1, 2, 3, 4\}$, $\mathcal{B} = \{0, 1\}$. Define $c : 1 \mapsto 0, 2 \mapsto 1, 3 \mapsto 00, 4 \mapsto 01$. Then $c^*(114) = 0001 = c^*(312) = c^*(144)$ yet c is injective.

Notation: $|\mathcal{A}| = m$, $|\mathcal{B}| = a$, call c an a -ary code of size m . For example a 2-ary code is a binary one, and a 3-ary code is a ternary code.

Our aim is to construct decipherable codes with short word lengths. Assuming c is injective, the following codes are always decipherable:

- (i) A block code has all codewords of the same length (e.g Greek fire code);
- (ii) A comma code reserves a letter from \mathcal{B} to signal the end of a word (e.g Example 1.2);
- (iii) A prefix-free code is a code where no codeword is a prefix of any other distinct word (if $x, y \in \mathcal{B}^*$ then x is a prefix of y if $y = xz$ for some string $z \in \mathcal{B}^*$).

(i) and (ii) are special cases of (iii). As we can decode the message as it is received, prefix-free codes are sometimes called *instantaneous*.

Exercise: find a decipherable code which is not prefix-free.

Definition (Kraft's inequality). $|\mathcal{A}| = m$, $|\mathcal{B}| = a$, $c : \mathcal{A} \rightarrow \mathcal{B}^*$ has word lengths l_1, \dots, l_m . Then Kraft's inequality is

$$\sum_{i=1}^m a^{-l_i} \leq 1. \quad (*)$$

Theorem 1.1. A prefix-free code exists if and only if Kraft's inequality $(*)$ holds.

Proof. Rewrite $(*)$ as

$$\sum_{l=1}^s n_l a^{-l} \leq 1, \quad (**)$$

where n_l is the number of codewords with length l , and $s = \max_{1 \leq i \leq m} l_i$.

Now if $c : \mathcal{A} \rightarrow \mathcal{B}^*$ is prefix-free,

$$n_1 a^{s-1} + n_2 a^{s-2} + \dots + n_{s-1} a + n_s \leq a^s.$$

Indeed the LHS is the number of strings of length s in B with some codeword of c as a prefix, and the RHS is the total number of strings of length S . Dividing through by a^s we get (**).

Now given n_1, \dots, n_s satisfying (**), we try to construct a prefix-free code c with n_l codewords of length l , $\forall l \leq s$. Proceed by induction on s , $s = 1$ is clear (since (**) gives $n_1 \leq a$ so can construct code).

By the induction hypothesis, there exists a prefix-code \hat{c} with n_l codewords of length l for all $l \leq s - 1$. Then (**) implies

$$n_1 a^{s-1} + n_2 a^{s-2} + \dots + n_{s-1} a + n_s \leq a^s.$$

The first $s - 1$ terms on the LHS sum to the number of strings of length s with a codeword of \hat{c} as a prefix and the RHS is the number of strings of length s . Hence we can add at least n_s new codewords of length s to \hat{c} and maintain the prefix-free property. □

Remark. This proof is constructive: just choose codewords in order of increasing length, ensuring that no previous codeword is a prefix.

Theorem 1.2 (McMillan). *Any decipherable code satisfies Kraft's inequality.*

Proof (Karush, 1961). Let $c : \mathcal{A} \rightarrow \mathcal{B}^*$ be a decipherable code with word lengths l_1, \dots, l_m . Set $s = \max_{1 \leq i \leq m} l_i$. For $R \in \mathbb{N}$

$$\left(\sum_{i=1}^m a^{-l_i} \right)^R = \sum_{l=1}^{Rs} b_l a^{-l}, \quad (\dagger)$$

where b_l is the number of ways of choosing R codewords of total length l . Since c is decipherable, any string of length l formed from codewords must correspond to at most one sequence of codewords, i.e $b_l \leq |\mathcal{B}^l| = a^l$. Subbing this into (†)

$$\left(\sum_{i=1}^m a^{-l_i} \right)^R \leq \sum_{l=1}^{Rs} a^l a^{-l} = Rs,$$

so

$$\sum_{i=1}^m a^{-l_i} \leq (Rs)^{1/R} \rightarrow 1 \text{ as } R \rightarrow \infty.$$

Hence $\sum_{i=1}^m a^{-l_i} \leq 1$. □

Corollary 1.3. *A decipherable code with prescribed word lengths exists if and only if a prefix-free code with the same word lengths exists.*

Proof. Combine previous two theorems. □

Therefore we can restrict our attention to prefix-free codes.

2 Shannon's Noiseless Coding Theorem

Entropy is a measure of 'randomness' or 'uncertainty'. Suppose we have a random variable X taking a finite set of values x_1, \dots, x_n with probabilities p_1, \dots, p_n respectively. The *entropy* $H(X)$ of X is the expected number of fair coin tosses needed to simulate X (roughly speaking).

Example 2.1. Suppose $p_1 = p_2 = p_3 = p_4 = 1/4$. Identify (x_1, x_2, x_3, x_4) with (HH, HT, TH, TT) . Then the entropy is 2.

Example 2.2. Suppose $(p_1, p_2, p_3, p_4) = (1/2, 1/4, 1/8, 1/8)$. Identify (x_1, x_2, x_3, x_4) with (H, TH, TTH, TTT) . Then the entropy is

$$\frac{1}{2} \times 1 + \frac{1}{4} \times 2 + \frac{1}{8} \times 3 + \frac{1}{8} \times 3 = \frac{7}{4}.$$

In a sense, the previous example (2.1) was 'more random' than this.

Definition (Entropy). The *entropy* of X is

$$H(X) = - \sum_{i=1}^b p_i \log p_i.$$

(Recall that $\log =: \log_2$ here.) Note $H(X) \geq 0$. It is measured in *bits* (binary digits). Conventionally, we take $0 \log 0 = 0$.

Example 2.3. Take a biased coin $\mathbb{P}(H) = p, \mathbb{P}(T) = 1 - p$. Write $H(p, 1 - p) := H(p)$. Then

$$H(p) = -p \log p - (1 - p) \log(1 - p).$$

Note that $H'(p) = \log \frac{1-p}{p}$. Hence the entropy is maximised for $p = 1/2$ (giving entropy 1).

Proposition 2.1 (Gibbs' inequality). *Let $(p_1, \dots, p_n), (q_1, \dots, q_n)$ be probability distributions. Then*

$$- \sum_{i=1}^n p_i \log p_i \leq - \sum_{i=1}^n p_i \log q_i.$$

(The RHS is sometimes called the *cross entropy* or *mixed entropy*) Furthermore we have equality iff $p_i = q_i$ for all i .

Proof. Since $\log x = \frac{\ln x}{\ln 2}$, we may replace \log with \ln . Put $I = \{1 \leq i \leq n : p_i \neq 0\}$. Now $\ln x = x - 1$ for all $x > 0$ with equality iff $x = 1$. Hence $\ln \frac{q_i}{p_i} \leq \frac{q_i}{p_i} - 1$ for all $i \in I$. So

$$\begin{aligned} \sum_{i \in I} p_i \ln \frac{q_i}{p_i} &\leq \underbrace{\sum_{i \in I} q_i}_{\leq 1} - \underbrace{\sum_{i \in I} p_i}_{=1} \leq 0 \\ \implies - \sum_{i \in I} p_i \ln p_i &\leq - \sum_{i \in I} p_i \ln q_i \end{aligned}$$

$$\implies -\sum_{i=1}^n p_i \ln p_i \leq -\sum_{i=1}^n p_i \ln q_i.$$

If equality holds, then $\sum_{i \in I} q_i = 1$ and $\frac{p_i}{q_i} = 1$ for all $i \in I$. So $q_i = p_i$ for all $1 \leq i \leq n$. \square

Corollary 2.2. $H(p_1, p_2, \dots, p_n) \leq \log n$ with equality iff $p_1 = p_2 = \dots = p_n = 1/n$.

Proof. Take $q_1 = q_2 = \dots = q_n = 1/n$ in Gibbs' inequality. \square

Let $\mathcal{A} = \{\mu_1, \dots, \mu_m\}$, $|\mathcal{B}| = a$ ($m, n \geq 2$). The random variable X takes values μ_1, \dots, μ_m with probabilities p_1, \dots, p_m .

Definition. If $c : \mathcal{A} \rightarrow \mathcal{B}^*$ is a code, we say it is *optimal* if it has the smallest possible expected word length. i.e. $\mathbb{E}S := \sum_{i=1}^m p_i l_i$ is minimal amongst all decipherable codes.

Theorem 2.3 (Shannon's Noiseless Coding Theorem). *The expected word length $\mathbb{E}S$ of an optimal code satisfies*

$$\frac{H(X)}{\log a} \leq \mathbb{E}S < \frac{H(X)}{\log a} + 1.$$

Remark. The lower bound is actually true for any decipherable code.

Proof. We first get the lower bound. Let $c : \mathcal{A} \rightarrow \mathcal{B}^*$ be decipherable with word lengths l_1, \dots, l_m . Let $q_i = \frac{a^{-l_i}}{D}$ where $D = \sum_{i=1}^m a^{-l_i}$. Note $\sum_{i=1}^m q_i = 1$. By Gibbs' inequality

$$\begin{aligned} H(X) &\leq -\sum_{i=1}^m p_i \log q_i \\ &= -\sum_{i=1}^m p_i (-l_i \log a - \log D) \\ &= \left(\sum_{i=1}^m p_i l_i \right) \log a + \log D. \end{aligned}$$

By McMillan, $D \leq 1$ so $\log D \leq 0$. Hence

$$H(X) \leq \left(\sum_{i=1}^m p_i l_i \right) \log a \implies \frac{H(X)}{\log a} \leq \mathbb{E}S.$$

And we have equality iff $p_i = a^{-l_i}$ for some integers l_1, \dots, l_m . Note we have only used decipherability so far.

Now we get the upper bound. Take $l_i = \lceil -\log_a p_i \rceil$. Then

$$-\log_a p_i \leq l_i < -\log_a p_i + 1.$$

Hence $\log_a p_i \geq -l_i$, so $p_i \geq a^{-l_i}$. Therefore $\sum_{i=1}^m a^{-l_i} \leq \sum_{i=1}^m p_i = 1$. By Kraft's inequality, there exists a prefix-free code c with word lengths l_1, \dots, l_m . c has expected word length

$$\mathbb{E}S = \sum_{i=1}^m p_i l_i < \sum_{i=1}^m p_i (-\log_a p_i + 1) = \frac{H(X)}{\log a} + 1.$$

□

Example 2.4 (Shannon-Fano Coding). We mimic the above proof: given p_1, \dots, p_m , set $l_i = \lceil -\log_a p_i \rceil$. Construct a prefix-free code with word lengths l_i by choosing codewords in order of increasing length, ensuring any new codeword has no previous codeword as a prefix (Kraft's inequality ensures we can do this).

Example 2.5. Take $a = 2, m = 5$.

i	p_i	$\lceil -\log_2 p_i \rceil$	code
1	0.4	2	00
2	0.2	3	010
3	0.2	3	011
4	0.1	4	1000
5	0.1	4	1001

Then $\mathbb{E}S = \sum_{i=1}^m p_i l_i = 2.8$, $H = H/\log a = 2.12$. [See also Carne p13.]

3 Huffman Coding

How to construct an optimal code? Take $\mathcal{A} = \{\mu_1, \dots, \mu_m\}$, $p_i = \mathbb{P}(X = \mu_i)$. For simplicity take $|\mathcal{B}| = a = 2$. Without loss of generality $p_1 \geq p_2 \geq \dots \geq p_m$. Huffman gave an inductive definition of codes that we can prove are optimal. If $m = 2$, we take codewords 0, 1. If $m > 2$, first take the Huffman code for messages $\mu_1, \dots, \mu_{m-2}, \nu$ with probabilities $p_1, \dots, p_{m-2}, p_{m-1} + p_m$. Then append 0 (respectively 1) to the codeword for ν to give a codeword for μ_{m-1} (respectively μ_m).

Notes.

- Huffman codes are prefix-free;
- Huffman codes are not unique: choice is needed if some of the p_i are equal.

Example 3.1. Revisit Example 2.5. We have

i	p_i	$c^{(1)}$	$p_i^{(2)}$	$c^{(2)}$	$p_i^{(3)}$	$c^{(3)}$	$p_i^{(4)}$	$c^{(4)}$
1	0.4	1	0.4	1	0.4	1	0.6	0
2	0.2	01	0.2	01	0.4	00	0.4	1
3	0.2	000	0.2	000	0.2	01		
4	0.1	0010	0.2	001				
5	0.1	0011						

Theorem 3.1. *Huffman codes are optimal (Huffman, 1952).*

Proof. We show by induction on m that Huffman codes of size $m = |\mathcal{A}|$ are optimal.

$m = 2$: codewords are 0, 1 - clearly optimal.

$m > 2$: let c_m be a Huffman code for X_m , which takes values μ_1, \dots, μ_m with probabilities $p_1 \geq p_2 \geq \dots \geq p_m$; each c_m is constructed from Huffman code c_{m-1} for X_{m-1} which takes values $\mu_1, \dots, \mu_{m-2}, \nu$ with probabilities $p_1, \dots, p_{m-2}, p_{m-1} + p_m$. Then the expected word length is

$$\mathbb{E}S_m = \mathbb{E}S_{m-1} + p_{m-1} + p_m. \quad (*)$$

Let c'_m be an optimal code for X_m . Wlog c'_m is still prefix-free. Wlog the last two codewords of c'_m have maximal length and differ only in the final position (see next lemma). Say

$$c'_m(\mu_{m-1}) = y0, \quad c'_m(\mu_m) = y1 \text{ for some } y \in \{0, 1\}^*.$$

Let c'_{m-1} be some prefix-free code for X_{m-1} , given by

$$c'_{m-1}(\mu_i) = \begin{cases} c'_m(\mu_i) & 1 \leq i \leq m-2 \\ c'_m(\mu_{m-1}) = y0 & \\ c'_m(\mu_m) = y1 & \end{cases}.$$

Then the expected word length satisfies

$$\mathbb{E}S'_m = \mathbb{E}S'_{m-1} + p_{m-1} + p_m. \quad (**)$$

By the inductive hypothesis, c_{m-1} is optimal, so $\mathbb{E}S_{m-1} \leq \mathbb{E}S'_{m-1}$. By (*) and (**) this implies $\mathbb{E}S_m \leq \mathbb{E}S'_m$. \square

Lemma 3.2. Suppose letters μ_1, \dots, μ_m in \mathcal{A} are sent with probabilities p_1, p_2, \dots, p_m . Let c be an optimal (prefix-free) code with word lengths l_1, \dots, l_m . Then

- (i) If $p_i > p + j$, then $l_i \leq l_j$;
- (ii) Amongst all codewords of maximal length there exist two that differ only in the final digit.

Proof. (i) is obvious. For (ii), could otherwise just delete the final digit of the codeword of maximal length (since prefix-free). \square

Remark. Note not all optimal codes are Huffman (look at the case $m = 4$).

Our main result says that if we have a prefix-free optimal code with word lengths l_1, \dots, l_m and associated probabilities p_1, \dots, p_m , then there is a Huffman code with these word lengths.

4 Joint Entropy

If X, Y are random variables with values in \mathcal{A} and \mathcal{B} respectively, then (X, Y) is a random variable with values in $\mathcal{A} \times \mathcal{B}$, and the *entropy* $H(X, Y)$ is called the joint entropy, given by

$$H(X, Y) = - \sum_{x \in \mathcal{A}} \sum_{y \in \mathcal{B}} \mathbb{P}(X = x, Y = y) \log \mathbb{P}(X = x, Y = y).$$

This generalises to any finite number of random variables.

Lemma 4.1. Let X, Y be random variables taking values in \mathcal{A} and \mathcal{B} respectively. Then

$$H(X, Y) \leq H(X) + H(Y),$$

with equality if and only if X and Y are independent.

Proof. Write $\mathcal{A} = \{x_1, \dots, x_m\}$, $\mathcal{B} = \{y_1, \dots, y_n\}$. Let

$$p_{ij} = \mathbb{P}(X = x_i, Y = y_j), \quad p_i = \mathbb{P}(X = x_i), \quad q_j = \mathbb{P}(Y = y_j).$$

Apply Gibbs' inequality to the probability distributions $\{p_{ij}\}$ and $\{p_i q_j\}$ to obtain

$$\begin{aligned} - \sum_{i,j} p_{ij} \log p_{ij} &\leq - \sum_{i,j} p_{ij} \log(p_i q_j) \\ &= - \sum_i \left(\sum_j p_{ij} \right) \log p_i - \sum_j \left(\sum_i p_{ij} \right) \log q_j \\ &= - \sum_i p_i \log p_i - \sum_j q_j \log q_j \\ &= H(X) + H(Y). \end{aligned}$$

With equality if and only if $p_{ij} = p_i q_j$ for all i, j .

□

Error-correcting codes

5 Noisy channels and Hamming's code

Definition. A *binary* $[m, n]$ -code is a subset C of $\{0, 1\}^n$ of size $m = |C|$. n is the length of the code and the elements of C are called codewords.

We use an $[n, m]$ -code to send one of m messages through a BSC (binary symmetric channel) making n uses of the channel. Clearly $1 \leq m \leq 2^n$, so $0 \leq \frac{1}{n} \log m \leq 1$.

Definition. For any $x, y \in \{0, 1\}^n$ the *Hamming distance* is

$$d(x, y) = |\{i : 1 \leq i \leq n, x_i \neq y_i\}|.$$

Definition.

- (i) The *ideal observer* decoding rule decodes $x \in \{0, 1\}^n$ as $c \in C$ maximising $\mathbb{P}(c \text{ sent} | x \text{ received})$.
- (ii) The *maximum likelihood* decoding rule decodes $x \in \{0, 1\}^n$ as $c \in C$ maximising $\mathbb{P}(x \text{ received} | c \text{ sent})$.
- (iii) The *minimum distance* decoding rule decodes $x \in \{0, 1\}^n$ as $c \in C$ minimising $d(x, C)$.

Lemma 5.1.

- (a) If all the messages are equally likely, then (i) and (ii) above are equivalent.
- (b) If $p < 1/2$ (error probability) then (ii) and (iii) are equivalent.

Remark. If $p = 1/2$ the code is called *useless*. If $p = 0$ the code is called *lossless*.

Proof.

- (a) We have

$$\mathbb{P}(c \text{ sent} | x \text{ received}) = \frac{\mathbb{P}(c \text{ sent}, x \text{ received})}{\mathbb{P}(x \text{ received})} = \frac{\mathbb{P}(c \text{ sent})}{\mathbb{P}(x \text{ received})} \mathbb{P}(x \text{ received} | c \text{ sent}).$$

So by hypothesis, $\mathbb{P}(c \text{ sent})$ is independent of $c \in C$. So for fixed x , maximising $\mathbb{P}(c \text{ sent} | x \text{ received})$ is the same as maximising $\mathbb{P}(x \text{ received} | c \text{ sent})$.

- (b) Let $r = d(x, c)$. Then $\mathbb{P}(x \text{ received} | c \text{ sent}) = p^r (1-p)^{n-r} = (1-p)^n \left(\frac{p}{1-p}\right)^r$. Since $p < 1/2$, $\frac{p}{1-p} < 1$. So maximising $\mathbb{P}(x \text{ received} | c \text{ sent})$ is the same as minimising r .

□

We choose to use minimum distance decoding from now on.

Example 5.1. Suppose 000, 111 are sent with probabilities $\alpha = 9/10$, $\beta = 1/10$ respectively through a BSC with error probability $p = 1/4$. Suppose 110 is received. Then

$$\mathbb{P}(000 \text{ sent} | 110 \text{ received}) = \frac{\alpha p^2(1-p)}{\alpha p^2(1-p) + (1-\alpha)p(1-p)^2} = \frac{3}{4},$$

$$\text{similarly } \mathbb{P}(111 \text{ sent} | 110 \text{ received}) = \frac{1}{4}.$$

So the ideal observer decodes as 000. But the maximum likelihood/minimum distance rules decode as 111.

Remarks.

- Minimum distance decoding may be expensive in terms of time and storage if $|C|$ is large.
- Need to specify a convention in case there is no unique maximiser (e.g make a random choice, or request the message is sent again).

We aim to detect, or even correct errors.

Definition. A code C is

- *d-error* detecting if changing up to d digits in each codeword can never produce another codeword. In other words, each codeword is of Hamming distance greater than d from every other codeword.
- *e-error* correcting if knowing that $x \in \{0,1\}^n$ differs from a codeword in at most e places we can deduce the codeword.

Examples.

- (a) A *repetition code* of length n has codewords $\underbrace{00 \dots 0}_{n \text{ times}}, \underbrace{11 \dots 1}_{n \text{ times}}$. This is a $[n, 2]$ -code. It is $(n-1)$ -error detecting and $\lfloor \frac{n-1}{2} \rfloor$ -error correcting. But the information rate is only $1/n$.
- (b) A *simple parity check code* or *paper tape code*: identify $\{0,1\}$ with \mathbb{F}_2 and let $C = \{(x_1, \dots, x_n) \in \{0,1\}^n : \sum_{i=1}^n x_i = 0\}$. This is a $[n, 2^{n-1}]$ -code, 1-error detecting but cannot correct errors. The information rate is $\frac{n-1}{n}$.
- (c) Hamming's original code (1950): a 1-error correcting binary $[7, 16]$ -code. Take $C \subseteq \mathbb{F}_2^7$ where

$$C = \{c \in \mathbb{F}_2^7 : c_1 + c_3 + c_5 + c_7 = 0, c_2 + c_3 + c_6 + c_7 = 0, c_4 + c_5 + c_6 + c_7 = 0\}.$$

The bits c_3, c_5, c_6, c_7 are arbitrary and c_1, c_2, c_4 are forced (called the check digits) so $|C| = 2^4$. To decode: suppose we receive $x \in \mathbb{F}_2^7$. We form the *syndrome*: $z = z_x = (z_1, z_2, z_4) \in \mathbb{F}_2^3$ where

$$z_1 = x_1 + x_3 + x_5 + x_7$$

$$z_2 = x_2 + x_3 + x_6 + x_7$$

$$z_4 = x_4 + x_5 + x_6 + x_7.$$

If $x \in C$, then $z_x = (0, 0, 0)$. If $d(x, c) = 1$ for some $c \in C$, then place where x and c differ is given by $z_1 + 2z_2 + 4z_4$ (not mod 2). Check: if $x = c + e_i$ where e_i has all 0's except a 1 in the i th position, then $z_x = z_{e_i}$, so check for each $1 \leq i \leq 7$.

Lemma 5.2. *The Hamming distance is a metric on \mathbb{F}_2^n .*

Proof. Trivial. □

Definition. The *minimum distance of a code* is the minimum of $d(c_1, c_2)$ for all codewords c_1, c_2 with $c_1 \neq c_2$.

Lemma 5.3. *Let C be a code with minimum distance $d > 0$. Then*

- (i) *C is $(d - 1)$ -error detecting, but cannot detect all sets of d errors.*
- (ii) *C is $\lfloor \frac{d-1}{2} \rfloor$ -error correcting, but cannot correct all sets of $\lfloor \frac{d-1}{2} \rfloor + 1$ errors.*

Proof.

- (i) If $x \in \mathbb{F}_2^n$ and $c \in C$ are such that $0 < d(x, c) \leq d - 1$, then we know that $x \notin C$ so this is $(d - 1)$ -error detecting. However there must exist $c_1, c_2 \in C$ such that $d(c_1, c_2) = d$, so we cannot say if there's an error if c_1 is 'corrupted' to c_2 in d errors.
- (ii) Take $e = \lfloor \frac{d-1}{2} \rfloor$. If $x \in \mathbb{F}_2^n$ and $c_1 \in C$ are such that $d(x, c_1) \leq e$ then for any $c_1 \neq c_2 \in C$ we have $d(x, c_2) \geq d(c_1, c_2) - d(c_1, x) \geq d - e > e$. So C is e -error correcting. Now take $c_1, c_2 \in C$ with $d(c_1, c_2) = d$. Then take $x \in \mathbb{F}_2^n$ such that x differs from c_1 is precisely $e + 1$ places where c_1 and c_2 differ. Then $d(c_1, x) = e + 1$ and $d(x, c_2) = d - (e + 1) \leq e + 1$. So C cannot be $(e + 1)$ -error correcting. □

Definition. A $[n, m]$ -code with minimum distance d is called a $[n, m, d]$ -code.

Notes.

- $m \leq 2^n$ with equality if and only if $C = \mathbb{F}_2^n$ (trivial code)
- $d \leq n$, with equality in case of the repetition code.

Example 5.2.

- (i) Repetition code of length n is a $[n, 2, n]$ -code, $(n - 1)$ -error detecting and $\lfloor \frac{n-1}{2} \rfloor$ -error correcting.
- (ii) Simple parity check code is a $[n, 2^{n-1}, 2]$ -code, 1-error detecting and 0-error correcting.
- (iii) Hamming's original code is 1-error correcting, implying $d \geq 3$. Also 0000000, 1110000 are distance 3 apart, so $d = 3$. So this is a $[7, 16, 3]$ -code and is 2-error detecting.

6 Covering estimates

Take $x \in \mathbb{F}_2^n$, $r \geq 0$. Then $\overline{B}(x, r) = \{y \in \mathbb{F}_2^n : d(x, y) \leq r\}$ is the *closed Hamming ball*. Denote $V(n, r) = |\overline{B}(x, r)| = \sum_{i=0}^r \binom{n}{i}$, the *volume*.

Lemma 6.1 (Hamming's bound). *An e -error correcting code C of length n has*

$$|C| \leq \frac{2^n}{V(n, e)}.$$

Proof. C is e -error correcting so $\{B(c, e)\}_{c \in C}$ are pairwise disjoint balls, so $\sum_{c \in C} |B(c, e)| = |C|V(n, e) \leq |\mathbb{F}_2^n| = 2^n$. \square

Lemma 6.2. *A code C of length n that can correct e errors is perfect if $|C| = \frac{2^n}{V(n, e)}$. Equivalently, for all $x \in \mathbb{F}_2^n$ there exists a unique $c \in C$ such that $d(x, c) \leq e$. In this case, any $e + 1$ errors will make you decode incorrectly.*

Example 6.1.

(a) Hamming $[7, 16, 3]$ -code is 1-error correcting and

$$\frac{2^n}{V(n, e)} = \frac{2^7}{V(7, 1)} = \frac{2^7}{1 + 7} = 2^4 = |C|.$$

(b) Binary repetition code of length n (for n odd) is perfect.

Remark. If $\frac{2^n}{V(n, e)} \notin \mathbb{Z}$ then there does not exist a perfect e -error correcting code of length n . Converse is also false ($n = 90, e = 2$ on Example Sheet 2).

Definition. Define $A(n, d) = \max\{m : \exists [n, m, d]\text{-code}\}$.

The $A(n, d)$ are unknown in general. But we have some special cases:

Examples.

- $A(n, 1) = 2^n$ (trivial code)
- $A(n, n) = 2$ (repetition code)
- $A(n, 2) = 2^{n-1}$ (simple parity check code)

Lemma 6.3. $A(n, d + 1) \leq A(n, d)$.

Proof. Let $m = A(n, d + 1)$ and take a $[n, m, d + 1]$ -code C . Let $c_1, c_2 \in C$ have $d(c_1, c_2) = d + 1$. Let c'_1 differ from c_1 in a single place where c_1 and c_2 differ. Hence $d(c'_1, c_2) = d$. If $c \in C \setminus \{c_1\}$, then $d(c, c_1) \leq d(c, c'_1) + d(c'_1, c_1)$ so $d + 1 \leq d(c, c'_1) + 1$. Hence $d(c, c'_1) \geq d$. So replacing c_1 with c'_1 , we get an $[n, m, d]$ -code. \square

Corollary 6.4. $A(n, d) = \max\{m : \exists [n, m, d']\text{-code for some } d' \geq d\}$.

Theorem 6.5.

$$\frac{2^n}{V(n, d - 1)} \underbrace{\leq}_{\text{GSV bound}} A(n, d) \underbrace{\leq}_{\text{Hamming bound}} \frac{2^n}{V(n, \lfloor \frac{d-1}{2} \rfloor)}.$$

Proof. We have already proved the Hamming bound. So let $m = A(n, d)$. Let C be a $[n, m, d]$ -code. Then there does not exist $x \in \mathbb{F}_2^n$ with $d(x, c) \geq d$ for all $c \in C$ (otherwise could replace C with $C \cup \{x\}$, contradicting maximality of m). Hence

$$\mathbb{F}_2^n \subseteq \bigcup_{c \in C} \overline{B}(c, d-1) \implies 2^n \leq \sum_{c \in C} |\overline{B}(c, d-1)| = mV(n, d-1).$$

□

Example 6.2. $n = 10, d = 3$, have $V(n, 1) = 11, V(n, 2) = 56$. The above theorem gives $19 \leq \frac{2^{10}}{56} \leq A(10, 3) \leq \frac{2^{10}}{11} \leq 93$. It was known that $72 \leq A(10, 3) \leq 93$, but the exact value of $A(10, 3)$ was only found in 1999.

Asymptotics of $V(n, r)$

We study $\frac{\log A(n, \lfloor n\delta \rfloor)}{n}$ as $n \rightarrow \infty$ to see how large the information rate can be for a given error rate.

Proposition 6.6. *Let $\delta \in (0, 1/2)$. Then*

$$(i) \log V(n, \lfloor n\delta \rfloor) \leq nH(\delta);$$

$$(ii) \frac{1}{n} \log A(n, \lfloor n\delta \rfloor) \geq 1 - H(\delta).$$

Where $H(\delta) = -\delta \log \delta - (1 - \delta) \log(1 - \delta)$.

Proof. First we show (i) \Rightarrow (ii): by the GSV bound,

$$A(n, \lfloor n\delta \rfloor) \geq \frac{2^n}{V(n, \lfloor n\delta \rfloor - 1)} \geq \frac{2^n}{V(n, \lfloor n\delta \rfloor)}$$

and so

$$\frac{\log A(n, \lfloor n\delta \rfloor)}{n} \geq 1 - \frac{\log V(n, \lfloor n\delta \rfloor)}{n} \geq 1 - H(\delta).$$

Now we prove (i): $H(\delta)$ is increasing for $\delta < 1/2$, so wlog we may assume $n\delta \in \mathbb{Z}$. Now

$$\begin{aligned} 1 &= (\delta + (1 - \delta))^n = \sum_{i=0}^n \binom{n}{i} \delta^i (1 - \delta)^{n-i} \geq \sum_{i=0}^{n\delta} \binom{n}{i} \delta^i (1 - \delta)^{n-i} \\ &= (1 - \delta)^n \sum_{i=0}^{n\delta} \binom{n}{i} \left(\frac{\delta}{1 - \delta} \right)^i \\ &\geq (1 - \delta)^n \sum_{i=0}^{n\delta} \binom{n}{i} \left(\frac{\delta}{1 - \delta} \right)^{n\delta} \\ &= \delta^{n\delta} (1 - \delta)^{n(1-\delta)} V(n, n\delta). \end{aligned}$$

Now taking logs:

$$0 \geq n(\delta \log \delta + (1 - \delta) \log(1 - \delta)) + \log V(n, n\delta).$$

□

The constant $H(\delta)$ in the above bound best possible:

Lemma 6.7. *We have*

$$\lim_{n \rightarrow \infty} \frac{\log V(n, \lfloor n\delta \rfloor)}{n} = H(\delta).$$

Proof. Exercise.

□

7 Constructing new codes from old

We're given C , a $[n, m, d]$ -code. Can check the details in the following:

Examples.

1. The *parity check extension* C^+ is

$$\left\{ \left(c_1, \dots, c_n, \sum_{i=1}^n c_i \right) : (c_1, \dots, c_n) \in C \right\}$$

Is a $[n+1, m, d']$ -code with $d \leq d' \leq d+1$, depending on whether d is odd or even.

2. Fix $1 \leq i \leq n$. Deleting the i th digit from each codeword gives the *punctured code* C^-

$$\{(c_1, \dots, c_{i-1}, c_{i+1}, \dots, c_n) : (c_1, \dots, c_n) \in C\}.$$

If $d \geq 2$, then it is a $[n-1, m, d']$ -code with $d-1 \leq d' \leq d$.

3. Fix $1 \leq i \leq n$, and $\alpha \in \mathbb{F}_2$. The *shortened code* C' is

$$\{(c_1, \dots, c_{i-1}, c_{i+1}, \dots, c_n) : (c_1, \dots, c_{i-1}, \alpha, c_{i+1}, \dots, c_n) \in C\}.$$

It has parameters $[n, m', d']$ with $d' \geq d$ and $m' \geq \frac{m}{2}$ for a suitable choice of α .

Shannon's Theorems

8 AEP and Shannon's first coding theorem

Definition. A *source* is a sequence of random variables X_1, X_2, \dots taking values in some alphabet \mathcal{A} . A source is *Bernoulli (memoryless)* if X_1, X_2, \dots are iid: write (X, X_n) . A source X_1, X_2, \dots is *reliably encodable at rate r* if there exists a sequence of subsets $(A_n)_{n \geq 1}$ with $A_n \subseteq \mathcal{A}^n$ such that:

1. $\lim_{n \rightarrow \infty} \frac{\log A_n}{n} = r;$
2. $\lim_{n \rightarrow \infty} \mathbb{P}((X_1, \dots, X_n) \in A_n) = 1.$

The *information rate* H of a source is the infimum of all reliable encoding rates.
Exercise: $0 \leq H \leq \log |\mathcal{A}|$ with both bounds attainable.

Shannon's first coding theorem computes the information rate of certain sources, including Bernoulli sources.

Reminders from IA Probability:

We have a probability space $(\Omega, \mathcal{F}, \mathbb{P})$. A discrete random variable X is a function $X : \Omega \rightarrow \mathcal{A}$. The probability mass function $p_x : \mathcal{A} \rightarrow [0, 1]$ is defined by $x \mapsto \mathbb{P}(X = x)$. Can consider $p \circ X = p(X) : \Omega \rightarrow [0, 1]$, a random variable taking values in $[0, 1]$.

Given a source X_1, X_2, \dots of random variables with values in \mathcal{A} , the probability mass function of $X^{(n)} = (X_1, \dots, X_n)$ is $p_{X^{(n)}}$ given by $(x_1, \dots, x_n) \mapsto \mathbb{P}((X_1, \dots, X_n) = (x_1, \dots, x_n))$. Since $p_{X^{(n)}} : \mathcal{A}^n \rightarrow [0, 1]$ and $X^{(n)} : \Omega \rightarrow \mathcal{A}^n$, you can form $p(X^{(n)}) = p_{X^{(n)}} \circ X^{(n)} : \Omega \rightarrow [0, 1]$.

Example 8.1. Let $\mathcal{A} = \{A, B, C\}$. Suppose

$$X^{(2)} = \begin{cases} AB & \text{with probability 0.3} \\ AC & \text{with probability 0.1} \\ BC & \text{with probability 0.1} \\ BA & \text{with probability 0.2} \\ CA & \text{with probability 0.25} \\ CB & \text{with probability 0.05} \end{cases}.$$

Then

$$p(X^{(2)}) = \begin{cases} 0.3 & \text{with probability 0.3} \\ 0.1 & \text{with probability 0.2} \\ 0.2 & \text{with probability 0.2} \\ 0.25 & \text{with probability 0.25} \\ 0.05 & \text{with probability 0.05} \end{cases}.$$

So some points are “lumped together”.

Given a source X_1, X_2, \dots *converges in probability* to a random variable L (possibly constant) if for all $\varepsilon > 0$, $\mathbb{P}(|X_n - L| > \varepsilon) \rightarrow 0$ as $n \rightarrow \infty$. We write $X_n \xrightarrow{\mathbb{P}} L$.

The *Weak Law of Large Numbers* (WLLN) says that if $(X; X_n)$ are iid real-valued random variables with finite expectation $\mathbb{E}X$, we have

$$\frac{1}{n} \sum_{i=1}^n X_i \xrightarrow{\mathbb{P}} \mathbb{E}X.$$

Example 8.2. If X_1, X_2, \dots are iid Bernoulli, then $p(X_1), p(X_2), \dots$ are iid random variables and $p(X_1, \dots, X_n) = p(X_1) \dots p(X_n)$. Note

$$-\frac{1}{n} \log p(X_1, \dots, X_n) = -\frac{1}{n} \sum_{i=1}^n \log p(X_i) \xrightarrow{\mathbb{P}} -\mathbb{E}(-\log p(X_1)) = H(X_1) \text{ as } n \rightarrow \infty.$$

Lemma 8.1. *The information rate of a Bernoulli source X_1, X_2, \dots is at most the expected word length of an optimal code $c : \mathcal{A} \rightarrow \{0, 1\}^*$ for X_1 .*

Proof. Let l_1, l_2, \dots be the lengths of codewords when we encode X_1, X_2, \dots using c . Let $\varepsilon > 0$. Set $A_n = \{x \in \mathcal{A}^n : c^*(x) \text{ has length at less than } n(\mathbb{E}l_1 + \varepsilon)\}$. Then

$$\begin{aligned} \mathbb{P}((X_1, \dots, X_n) \in A_n) &= \mathbb{P}\left(\sum_{i=1}^n l_i < n(\mathbb{E}l_1 + \varepsilon)\right) \\ &= \mathbb{P}\left(\left|\frac{1}{n} \sum_{i=1}^n l_i - \mathbb{E}l_1\right| < \varepsilon\right) \\ &\rightarrow 1 \text{ as } n \rightarrow \infty. \end{aligned}$$

Now, c is decipherable so c^* is injective. Hence $|A_n| \leq 2^{n(\mathbb{E}l_1 + \varepsilon)}$. Making A_n larger if necessary, $|A_n| = \lfloor e^{n(\mathbb{E}l_1 + \varepsilon)} \rfloor$ so

$$\frac{\log(A_n)}{n} \rightarrow \mathbb{E}l_1 + \varepsilon.$$

Hence X_1, X_2, \dots is reliably encodable at rate $r = \mathbb{E}l_1 + \varepsilon$ for all $\varepsilon > 0$. Hence the information rate is at most $\mathbb{E}l_1$. \square

Corollary 8.2. *A Bernouilli source has information rate less than $H(X_1) + 1$.*

Proof. Combine the above with the Noiseless Coding Theorem. \square

We encode X_1, X_2, \dots in blocks

$$\underbrace{X_1, \dots, X_N}_{Y_1}, \underbrace{X_{n+1}, \dots, X_{2N}, \dots}_{Y_2}, \dots$$

so Y_1, Y_2, \dots take values in \mathcal{A}^N . Exercise: show that if X_1, X_2, \dots has information rate H then Y_1, Y_2, \dots has information rate NH .

Proposition 8.3. *The information rate H of a Bernouilli source X_1, X_2, \dots is at most $H(X_1)$.*

Proof. Apply the previous corollary to Y_1, Y_2, \dots and obtain

$$NH < H(Y_1) + 1 = H(X_1, \dots, X_N) + 1 = \sum_{i=1}^N H(X_i) + 1 = NH(X_1, \dots, X_n) + 1.$$

Hence $H < H(X_1) + \frac{1}{N}$. Since N is arbitrary, $H \leq H(X_1)$. \square

Definition. A source X_1, X_2, \dots satisfies the *Asymptotic Equipartition Property* (AEP) for some constant $H \geq 0$ if

$$-\frac{1}{n} \log p(X_1, X_2, \dots) \xrightarrow{\mathbb{P}} H \text{ as } n \rightarrow \infty.$$

Example 8.3. Tossing a biased coin, $\mathbb{P}(H) = p$. Let $(X; X_n)$ be the results of independent coin tosses. After a large number N of tosses, expect on average pN heads and $(1-p)N$ tails. The probability of any particular sequence of pN heads and $(1-p)N$ tails is $p^{pN}(1-p)^{(1-p)N} = 2^{N(p \log p + (1-p) \log (1-p))} = 2^{-NH(X)}$. Not every sequence of tosses will be like this, but there is only a small probability of “atypical” sequences. With high probability we get a “typical” sequence and its probability will be close to $2^{-NH(X)}$.

Lemma 8.4. *The AEP for a source X_1, X_2, \dots is equivalent to the following property*

$\forall \varepsilon > 0 \exists n_0(\varepsilon)$ such that $\forall n \geq n_0(\varepsilon) \exists$ a “typical set” $T_n \subseteq \mathcal{A}^n$ such that

(i) $\mathbb{P}((X_1, \dots, X_n) \in T_n) > 1 - \varepsilon$;

(ii) $2^{-n(H+\varepsilon)} \leq p(x_1, \dots, x_n) \leq 2^{-n(H-\varepsilon)}$ for all $(x_1, \dots, x_n) \in T_n$.

Proof. Obvious and non-examinable. □

Theorem 8.5 (Shannon's First Coding Theorem (FCT)). *If a source X_1, X_2, \dots satisfies the AEP with constant H , then the source has information rate H .*

Proof. Let $\varepsilon > 0$ and let $T_n \subseteq \mathcal{A}^n$ be typical sets. Then for some $n_0(\varepsilon)$ and all $n \geq n_0(\varepsilon)$

$$p(x_1, \dots, x_n) \geq 2^{-n(H+\varepsilon)} \text{ for all } (x_1, \dots, x_n) \in T_n$$

$$\implies \mathbb{P}(T_n) \geq 2^{-n(H+\varepsilon)} |T_n| \implies 1 \geq 2^{-n(H+\varepsilon)} |T_n| \implies \frac{\log |T_n|}{n} \leq H + \varepsilon.$$

Taking $A_n = T_n$, shows the source is reliably encodable at rate $H + \varepsilon$. Conversely, if $H = 0$, we're done. Otherwise pick $0 < \varepsilon < H/2$ and suppose for contradiction that the source is reliably encodable at rate $H - 2\varepsilon$, say with sets $A_n \subseteq \mathcal{A}^n$. Let $T_n \subseteq \mathcal{A}^n$ be typical sets. Then for all $(x_1, \dots, x_n) \in T_n$

$$p(x_1, \dots, x_n) \leq 2^{-n(H-\varepsilon)}$$

$$\implies \mathbb{P}(A_n \cap T_n) \leq 2^{-n(H-\varepsilon)} |A_n|$$

$$\implies \frac{\log \mathbb{P}(A_n \cap T_n)}{n} \leq H - \varepsilon + \frac{\log |A_n|}{n} \rightarrow -(H - \varepsilon) + H - 2\varepsilon = -\varepsilon.$$

Hence $\log \mathbb{P}(A_n \cap T_n) \rightarrow -\infty$ and $\mathbb{P}(A_n \cap T_n) \rightarrow 0$. But $\mathbb{P}(T_n) \leq \mathbb{P}(A_n \cap T_n) + \mathbb{P}(\mathcal{A}^n \setminus A_n) \rightarrow 0$, a contradiction to $\mathbb{P}(T_n) \rightarrow 1$. Thus the information rate is exactly H . \square

Corollary 8.6. *A Bernoulli source X_1, X_2, \dots has information rate $H(X_1)$.*

Proof. We've already seen that $-\frac{1}{n} \log p(X_1, \dots, X_n) \xrightarrow{\mathbb{P}} H(X_1)$, so done by Shannon's First Coding Theorem. \square

Remarks.

- The AEP is useful for noiseless coding. We can
 - encode the typical sequences using a block code;
 - encode the atypical sequences arbitrarily.
- Many sources, which are not necessarily Bernoulli satisfy the AEP. Under suitable hypotheses the sequence $\frac{1}{n} H(X_1, \dots, X_n)$ is decreasing and the AEP is satisfied.

9 Capacity & Shannon's second coding theorem

Recall:

Definition. We model n uses of a channel by the n th extension, with input alphabet \mathcal{A}^n and output alphabet \mathcal{B}^n . A code C of length n is a function $\mathcal{M} \rightarrow \mathcal{A}^n$ where \mathcal{M} is the set of possible messages. Implicitly we also have a decoding rule $\mathcal{B}^n \rightarrow \mathcal{M}$. The size of C is $m = |\mathcal{M}|$. The information rate is $\rho(C) = \frac{1}{n} \log_2 m$. The error rate is $\hat{e}(C) = \max_{x \in \mathcal{M}} \mathbb{P}(\text{error} | x \text{ sent})$.

Definition. A channel can *transmit reliably at rate R* if there exists $(C_n)_{n=1}^\infty$ with each C_n a code of length n such that

$$\lim_{n \rightarrow \infty} \rho(C_n) = R \text{ \& } \lim_{n \rightarrow \infty} \hat{e}(C_n) = 0.$$

The *capacity* is the supremum of all reliable transmission rates.

Suppose we are given a source where

- it has information rate r bits per second;
- it emits symbols at s symbols per second.

Suppose we are also given a channel where

- it has capacity R bits per transmission;
- it transmits symbols at S transmissions per second.

Usually, information theorists take $S = s = 1$. If $rs \leq RS$ then you can encode and transmit reliably, and if $rs > RS$ you cannot.

We'll compute the capacity of a BSC with error probability p .

Proposition 9.1. *A binary symmetric channel with error probability $p < 1/4$ has non-zero capacity.*

Proof. We use the GSV bound. Pick $\delta \in (2p, 1/2)$. We claim reliable transmission at rate $R = 1 - H(\delta) > 0$. Let C_n be a code of length n , and suppose it has minimum distance $\lfloor n\delta \rfloor$ of maximal size. Then (by Proposition 6.6(ii))

$$|C_n| = A(n, \lfloor n\delta \rfloor) \geq 2^{n(1-H(\delta))}.$$

Replacing C_n by a subcode, we can assume $|C_n| = \lfloor 2^{nR} \rfloor$ and still minimum distance $\geq \lfloor n\delta \rfloor$. Using minimum distance decoding

$$\begin{aligned} \hat{e}(C_n) &\leq \mathbb{P} \left(\text{in } n \text{ uses, BSC makes } \geq \left\lfloor \frac{\lfloor n\delta \rfloor - 1}{2} \right\rfloor \text{ errors} \right) \\ &\leq \mathbb{P} \left(\text{in } n \text{ uses, BSC makes } \geq \left\lfloor \frac{n\delta - 1}{2} \right\rfloor \text{ errors} \right). \end{aligned}$$

Pick $\varepsilon > 0$ with $p + \varepsilon < \frac{\delta}{2}$. For n sufficiently large, $\frac{n\delta - 1}{2} = n \left(\frac{\delta}{2} - \frac{1}{2n} \right) > n(p + \varepsilon)$. Hence $\hat{e}(C_n) \leq \mathbb{P}(\text{BSC makes } \geq n(p + \varepsilon) \text{ errors}) \rightarrow 0$, using the next lemma. \square

Lemma 9.2. Let $\varepsilon > 0$. A BSC with error probability p is used to transmit n digits. Then

$$\lim_{n \rightarrow \infty} \mathbb{P}(\text{BSC makes } \geq n(p + \varepsilon) \text{ errors}) = 0.$$

Proof. Consider random variables U_i taking value 1 if the i th digit is mistransmitted, and value 0 otherwise. Then $\mathbb{E}U_i = p$, so $\mathbb{P}(\text{BSC makes } \geq n(p + \varepsilon) \text{ errors}) \leq \mathbb{P}(|\frac{1}{n} \sum_{i=1}^n U_i - p| \geq \varepsilon) \rightarrow 0$. \square

Definition (Conditional Entropy). Let X, Y be random variables taking values in alphabets \mathcal{A}, \mathcal{B} respectively. Then we define the *conditional entropy*

$$H(X|Y = y) = - \sum_{x \in \mathcal{A}} \mathbb{P}(X = x|Y = y) \log \mathbb{P}(X = x|Y = y)$$

$$H(X|Y) = \sum_{y \in \mathcal{B}} \mathbb{P}(Y = y) H(X|Y = y).$$

Clearly $H(X|Y) \geq 0$.

Lemma 9.3. We have

$$H(X, Y) = H(X|Y) + H(Y).$$

Proof.

$$\begin{aligned} H(X|Y) &= - \sum_{y \in \mathcal{B}} \sum_{x \in \mathcal{A}} \mathbb{P}(X = x|Y = y) \mathbb{P}(Y = y) \log \mathbb{P}(X = x|Y = y) \\ &= - \sum_{y \in \mathcal{B}} \sum_{x \in \mathcal{A}} \mathbb{P}(X = x, Y = y) \log \left(\frac{\mathbb{P}(X = x, Y = y)}{\mathbb{P}(Y = y)} \right) \\ &= - \sum_{y \in \mathcal{B}} \sum_{x \in \mathcal{A}} \mathbb{P}(X = x, Y = y) \log \mathbb{P}(X = x, Y = y) \\ &\quad + \sum_{y \in \mathcal{B}} \sum_{x \in \mathcal{A}} \mathbb{P}(X = x, Y = y) \log \mathbb{P}(Y = y) \\ &= H(X, Y) - \sum_{y \in \mathcal{B}} \mathbb{P}(Y = y) \log \mathbb{P}(Y = y) \\ &= H(X, Y) - H(Y). \end{aligned}$$

\square

Example 9.1. We roll fair die. Let X be the value defined by the roll, let Y be equal to 0 if X is even and 1 if X is odd. Then $H(X, Y) = H(X) = \log 6$, $H(Y) = \log 2 = 1$, $H(X|Y) = \log 3 = H(X, Y) - H(Y)$. Similarly $H(Y|X) = 0 = H(X, Y) - H(X)$.

Corollary 9.4. $H(X|Y) \leq H(X)$.

Proof. Combine the previous with Lemma 4.1. \square

Now replace X, Y with random vectors $X^{(r)} = (X_1, \dots, X_r), Y^{(s)} = (Y_1, \dots, Y_s)$. Similarly can define $H(X_1, \dots, X_r | Y_1, \dots, Y_s) = H(X^{(r)} | Y^{(s)})$.

Remark. $H(X, Y | Z)$ is the entropy of (X, Y) given Z , not the entropy of X and $Y | Z$.

Lemma 9.5. *Let X, Y, Z be random variables. Then $H(X, Y) \leq H(X | Y, Z) + H(Z)$.*

Proof. We expand $H(X, Y, Z)$ in two different ways.

$$H(X, Y, Z) = H(Z | X, Y) + H(X | Y) + H(Y),$$

$$H(X, Y, Z) = H(X | Y, Z) + H(Z | Y) + H(Y).$$

Since $H(Z | X, Y) \geq 0$, we have

$$\begin{aligned} H(X | Y) &\leq H(X | Y, Z) + H(Z | Y) \\ &\leq H(X | Y, Z) + H(Z). \end{aligned}$$

□

Proposition 9.6 (Fano's inequality). *Let X, Y be random variables taking values in \mathcal{A} , $|\mathcal{A}| = m$. Let $p = \mathbb{P}(X \neq Y)$. Then*

$$H(X | Y) \leq H(p) + p \log(m - 1).$$

Proof. Define $Z = 0$ if $X = Y$ and $Z = 1$ if $X \neq Y$. Then $\mathbb{P}(Z = 0) = 1 - p$, $\mathbb{P}(Z = 1) = p$. So $H(Z) = H(p)$. By the previous lemma,

$$H(X | Y) \leq H(p) + H(X | Y, Z). \quad (*)$$

Since $Z = 0$ implies $X = Y$, $H(X | Y = y, Z = 0) = 0$. Also there are $m - 1$ remaining possibilities for X , so $H(X | Y = y, Z = 1) \leq \log(m - 1)$. Therefore

$$\begin{aligned} H(X | Y) &= \sum_{y, z} \mathbb{P}(Y = y, Z = z) H(X | Y = y, Z = z) \\ &\leq \sum_{y \in \mathcal{A}} \mathbb{P}(Y = y, Z = 1) \log(m - 1) \\ &= \mathbb{P}(Z = 1) \log(m - 1) \\ &= p \log(m - 1). \end{aligned}$$

So by $(*)$, $H(X | Y) \leq H(p) + p \log(m - 1)$. □

Remark. $H(p)$ represents the information needed to decide whether or not there is an error, and $p \log(m - 1)$ represents the information needed to resolve the error assuming the worst possible case.

Definition. Let X, Y be random variables. Then the *mutual information* is $I(X; Y) := H(X) - H(X | Y)$. By Corollary 9.4, $I(X; Y) \geq 0$ with equality if and only if X, Y are independent. Clearly $I(X; Y) = I(Y; X)$.

Suppose we're given a discrete memoryless channel (DMC) with input alphabet \mathcal{A} , $|\mathcal{A}| = m$ and output alphabet \mathcal{B} . Let X be a random variable taking values in \mathcal{A} and used as input to the channel. Let Y be a random variable output, depending on X and the channel matrix.

Definition. The (*information*) *capacity* is $\max_X I(X; Y)$.

Remarks.

- The maximum is over all probability distributions (p_1, \dots, p_m) for X on \mathcal{A} .
- The maximum is attained since $(p_1, \dots, p_m) \mapsto I((p_1, \dots, p_m); Y)$ is continuous on the compact set

$$\{(p_1, \dots, p_m) \in [0, 1]^m : p_1 + \dots + p_m = 1\}.$$

- The information capacity depends only on the channel matrix.

Theorem 9.7 (Shannon's Second Coding Theorem). *For a DMC, the operational capacity is equal to the information capacity.*

Remark. We'll show one inequality in general and the other for the BSC only.

Assuming Shannon's Second Coding Theorem, let us compute the capacity of certain channels.

Example 9.2. BSC, error probability p . Input X : $\mathbb{P}(X = 0) = \alpha$, $\mathbb{P}(X = 1) = 1 - \alpha$. Output Y : $\mathbb{P}(Y = 0) = \alpha(1 - p) + (1 - \alpha)p$, $\mathbb{P}(Y = 1) = (1 - \alpha)(1 - p) + \alpha p$. Then C is

$$\begin{aligned} \max_{\alpha} I(X; Y) &= \max_{\alpha} (H(Y) - H(Y|X)) \\ &= \max_{\alpha} (H(\alpha(1 - p) + (1 - \alpha)p) - H(p)) \\ &= 1 - H(p). \end{aligned}$$

Where the maximum is attained for $\alpha = 1/2$. Hence $C = 1 + p \log p + (1 - p) \log(1 - p)$.

Remark. We can choose to calculate either $H(Y) - H(Y|X)$ or $H(X) - H(X|Y)$ depending on which is easier.

Example 9.3. BEC, erasure probability p . Input X : $\mathbb{P}(X = 0) = \alpha$, $\mathbb{P}(X = 1) = 1 - \alpha$. Output Y : $\mathbb{P}(Y = 0) = \alpha(1 - p)$, $\mathbb{P}(Y = *) = p$, $\mathbb{P}(Y = 1) = (1 - \alpha)(1 - p)$. Can calculate $H(X|Y = 0) = 0$, $H(X|Y = 1) = 0$, $H(X|Y = *) = H(\alpha)$, which gives $H(X|Y) = pH(\alpha)$. So

$$\begin{aligned} C = \max_{\alpha} I(X; Y) &= \max_{\alpha} (H(X) - H(X|Y)) = \max_{\alpha} (H(\alpha) - pH(\alpha)) \\ &= (1 - p) \max_{\alpha} H(\alpha) \\ &= 1 - p. \end{aligned}$$

Where the max is attained for $\alpha = 1/2$.

Now model using a channel n times as the n th extension, i.e replace alphabets \mathcal{A}, \mathcal{B} with $\mathcal{A}^n, \mathcal{B}^n$,

$$\mathbb{P}(y_1, \dots, y_n \text{ received} | x_1, \dots, x_n \text{ sent}) = \prod_{i=1}^n \mathbb{P}(y_i | x_i).$$

Lemma 9.8. The n th extension of a DMC with information capacity C has information capacity nC .

Proof. Take random variable input X_1, \dots, X_n producing output Y_1, \dots, Y_n . Since the channel is memoryless,

$$H(Y_1, \dots, Y_n | X_1, \dots, X_n) = \sum_{i=1}^n H(Y_i | X_1, \dots, X_n) = \sum_{i=1}^n H(Y_i | X_i).$$

Hence

$$\begin{aligned}
 I(X_1, \dots, X_n; Y_1, \dots, Y_n) &= H(Y_1, \dots, Y_n) - H(Y_1, \dots, Y_n | X_1, \dots, X_n) \\
 &= H(Y_1, \dots, Y_n) - \sum_{i=1}^n H(Y_i | X_i) \\
 &\leq \sum_{i=1}^n (H(Y_i) - H(Y_i | X_i)) \\
 &= \sum_{i=1}^n I(X_i; Y_i) \leq nC.
 \end{aligned}$$

To finish, we find a distribution for X_1, \dots, X_n giving equality. Equality is attained by taking X_1, \dots, X_n independent, each of the same distribution such that $I(X_i; Y_i) = C$. Indeed, if X_1, \dots, X_n are independent, so are Y_1, \dots, Y_n . \square

Proposition 9.9. *For a DMC the (operational) capacity is at most the information capacity. Let C be the information capacity. Suppose reliable transmission is possible at some rate $R > C$, i.e there exist a sequence of codes $(C_n)_{n \geq 1}$ with C_n having length n and size $\lceil 2^{nR} \rceil$ for all n such that*

$$\lim_{n \rightarrow \infty} \rho(C_n) = R \text{ and } \lim_{n \rightarrow \infty} \hat{e}(C_n) = 0.$$

(Recall $\hat{e}(C_n) = \max_{c \in C_n} \mathbb{P}(\text{error} | c \text{ sent})$.)

Proof. Define the average error rate $e(C_n) = \frac{1}{|C_n|} \sum_{c \in C_n} \mathbb{P}(\text{error} | c \text{ sent})$. Note $e(C_n) \leq \hat{e}(C_n)$ and so $e(C_n) \rightarrow 0$ as $n \rightarrow \infty$. Take input random variable X equidistributed over C_n (i.e takes all values with equal probability). Let Y be the random variable output when X is transmitted and decoded. So $e(C_n) = \mathbb{P}(X \neq Y)$, define $p := e(C_n)$. Now we have

$$H(X) = \log |C_n| = \log \lceil 2^{nR} \rceil \geq nR - 1 \text{ for sufficiently large } n.$$

And

$$\begin{aligned}
 H(X|Y) &\leq H(p) + p \log(|C_n| - 1) && \text{(Fano's inequality)} \\
 &\leq 1 + pnR.
 \end{aligned}$$

Recall $I(X; Y) = H(X) - H(X|Y)$. By the previous lemma $nC \geq I(X; Y)$ so

$$\begin{aligned}
 nC &\geq nR - 1 - (1 + pnR) \\
 \implies pnR &\geq n(R - C) - 2 \implies p \geq \frac{n(R - C) - 2}{nR} \rightarrow \frac{R - C}{R} \neq 0.
 \end{aligned}$$

Since $R > C$, which contradicts $p \rightarrow 0$. Hence we conclude that we cannot transmit reliably at any rate exceeding C . \square

To complete the proof of Shannon's SCT for a BSC with error probability p , we need to show that the operational capacity is at least $1 - H(p)$ (i.e the information capacity).

Proposition 9.10. *Consider BSC with error probability p . Suppose $R < 1 - H(p)$. Then there is a sequence of codes $(C_n)_{n \geq 1}$ with C_n of length n and size $\lceil 2^{nR} \rceil$ for all n such that*

$$\lim_{n \rightarrow \infty} \rho(C_n) = R \text{ and } \lim_{n \rightarrow \infty} e(C_n) = 0.$$

Remark. Note that the above proposition deals with the average error rate e , not the error rate \hat{e} .

Proof. We'll use the method of 'random coding'. Wlog $p < 1/2$. Take $\varepsilon > 0$ such that

$$p + \varepsilon < 1/2 \text{ and } R < 1 - H(p + \varepsilon).$$

(This is possible by continuity of H). We use minimum distance decoding (in case of a tie, make arbitrary choice). Let $m = \lceil 2^{nR} \rceil$. Pick $C = \{c_1, \dots, c_m\}$ at random from $\mathcal{C} = \{[n, m]\text{-codes}\}$, a set of size $\binom{2^n}{m}$. Choose $1 \leq i \leq m$ at random (i.e each with probability $1/m$). We sent c_i through the channel and get output Y . Then

$$\begin{aligned} \mathbb{P}(Y \text{ not decoded as } c_i) &= \text{average value of } e(C) \text{ as } C \text{ runs through } \mathcal{C} \\ &= \frac{1}{|\mathcal{C}|} \sum_{C \in \mathcal{C}} e(C). \end{aligned}$$

We can pick C_n , an $[n, m]$ -code with $e(C_n)$ at most this average. So it is enough to show

$$\mathbb{P}(Y \text{ is not decoded as } c_i) \rightarrow 0 \text{ as } n \rightarrow \infty.$$

Let $r = \lfloor n(p + \varepsilon) \rfloor$. Then if $B(Y, r) \cap C = \{c_i\}$, Y is correctly decoded as c_i . So

$$\mathbb{P}(Y \text{ is not decoded as } c_i) \leq \mathbb{P}(c_i \notin B(Y, r)) + \mathbb{P}(B(Y, r) \cap C \supsetneq \{c_i\}).$$

Consider the two cases separately, according as $d(c_i, Y) > r$ (case (i)) or $d(c_i, Y) \leq r$ (case (ii)). For (i) we have

$$\begin{aligned} \mathbb{P}(d(c_i, Y) > r) &= \mathbb{P}(\text{BSC makes } > r \text{ errors}) \\ &= \mathbb{P}(\text{BSC makes } > n(p + \varepsilon) \text{ errors}) \\ &\rightarrow 0 \text{ as } n \rightarrow \infty \text{ by a previous theorem.} \end{aligned}$$

For (ii), if $j \neq i$, $\mathbb{P}(c_j \in B(Y, r) | c_i \in B(Y, r)) = \frac{V(n, r) - 1}{2^n - 1} \leq \frac{V(n, r)}{2^n}$ (since

$V(n, r) \leq 2^n$. Hence

$$\begin{aligned}
 \mathbb{P}(B(Y, r) \cap C \supsetneq \{c_i\}) &\leq \sum_{j \neq i} \mathbb{P}(c_j \in B(Y, r) \text{ and } c_i \in B(Y, r)) \\
 &\leq \sum_{j \neq i} \mathbb{P}(c_j \in B(Y, r) | c_i \in B(Y, r)) \\
 &\leq (m-1) \frac{V(n, r)}{2^n} \\
 &\leq \frac{mV(n, r)}{2^n} \\
 &\leq 2^{nR} 2^{nH(p+\varepsilon)2^{-n}} \\
 &= 2^{n(R-(1-H(p+\varepsilon)))} \xrightarrow{n \rightarrow \infty} 0
 \end{aligned}$$

since $R < 1 - H(p + \varepsilon)$. \square

Proposition 9.11. *We can replace e by \hat{e} in the previous result.*

Proof. Pick R' with $R < R' < 1 - H(p)$. The previous result constructs a sequence $(C'_n)_{n \geq 1}$ of codes, C'_n of length n and size $\lfloor 2^{nR'} \rfloor$ for all n , and $e(C'_n) \rightarrow 0$ as $n \rightarrow \infty$.

For each n , order the codewords of C'_n by $\mathbb{P}(\text{error} | c \text{ sent})$ and delete the worst half, to get a code C_n with $\hat{e}(C_n) \leq 2e(C'_n)$. So $\hat{e}(C_n) \rightarrow 0$ as $n \rightarrow \infty$. Since C_n has length n and size $\lfloor 2^{nR'} \rfloor / 2 = \text{floor } 2^{nR'-1}$. But $2^{nR'-1} = 2^{n(R'-1/n)} \geq 2^{nR}$ for all n sufficiently large. So can replace C'_n by a code of smaller size $\lfloor 2^{nR} \rfloor$ and still get $\hat{e}(C_n) \rightarrow 0$ and $\rho(C_n) \rightarrow R$. \square

Remarks.

1. A BSC with error probability p has operational capacity $1 - H(p)$. Indeed the above proposition says we can transmit reliably at any rate $R < 1 - H(p)$, whereas 9.9 says the capacity is at most $1 - H(p)$.
2. Shannon's SCT shows that good codes exist, but the proof does not say how to construct them.

10 The Kelly criterion

Let $1 > p > 0$, $u > 0$, $1 > w \geq 0$. Suppose a coin is tossed n times in succession, $\mathbb{P}(H) = p$. If I pay k ahead of a particular throw, then get back ku if the throw is heads, or get back 0 if throw is tails.

Say we have initial bankroll $X_0 = 1$. If we have bankroll of X_n after n th toss, bet wX_n , retaining $(1 - w)X_n$. Then bankroll X_{n+1} after $(n + 1)$ th throw is

$$X_{n+1} = \begin{cases} X_n(wu + (1 - w)) & \text{if } n\text{th toss is heads} \\ X_n(1 - w) & \text{if } n\text{th toss is tails} \end{cases}$$

Let $Y_{n+1} = X_{n+1}/X_n$. Then Y_1, Y_2, \dots , is a sequence of iid random variables. Hence $\log Y_1, \log Y_2, \dots$ is a sequence of iid random variables. Note $\log X_n = \sum_{i=1}^n \log Y_i$.

Lemma 10.1. Let $\mathbb{E} \log Y_1 = \mu$, $\text{Var}(\log Y_1) = \sigma^2$. Then if $a > 0$

$$(i) \quad \mathbb{P} \left(\left| \frac{1}{n} \sum_{i=1}^n \log Y_i - \mu \right| \geq a \right) \leq \frac{\sigma^2}{na^2}$$

$$(ii) \quad \mathbb{P} \left(\left| \frac{\log X_n}{n} - \mu \right| \geq a \right) \leq \frac{\sigma^2}{na^2}$$

$$(iii) \quad \text{Given } \varepsilon > 0 \text{ and } \delta > 0 \text{ there exists } N \in \mathbb{N} \text{ such that } \mathbb{P} \left(\left| \frac{\log X_n}{n} - \mu \right| \geq \delta \right) \leq \varepsilon \text{ for all } n \geq N.$$

Proof. Obvious by Chebyshev. \square

Lemma 10.2. Consider a single toss of a coin, $\mathbb{P}(H) = p < 1$. Suppose a bet of k on H has payout ratio $u > 0$. Suppose we have initial bankroll 1, and we bet w on H , retaining $1 - w$ ($0 \leq w < 1$). If $Y =$ fortune after toss then $\mathbb{E} \log Y = p \log(1 + (u - 1)w) + (1 - p) \log(1 - w)$. The value of $\mathbb{E} \log Y$ is maximised by taking $w = 0$ if $up \leq 1$, and otherwise taking $w = \frac{up-1}{u-1}$.

Proof. Not given (see Körner's book). \square

Remark. Put $q = 1 - p$. If $up > 1$, at the optimum w have

$$\mathbb{E} \log Y = p \log p + q \log q + \log u - q \log(u - 1)$$

and $-(p \log p + q \log q)$ is the entropy of a simple probabilistic system.

The suggestion that those making a long sequence of bets should maximise the expectation of the log is called Kelly's criterion.

Exercise: show that if you bet less than the optimum, your bankroll will tend to increase, but more slowly; and if you bet more than some critical proportion, your bankroll will tend to decrease.

11 Algebraic Coding Theory

11.1 Linear codes

Definition. A code $C \subseteq \mathbb{F}_2^n$ is *linear* if

- $0 \in C$;
- whenever $x, y \in C$, we have $x + y \in C$.

i.e C is a \mathbb{F}_2 -vector subspace of \mathbb{F}_2^n .

Definition. The *rank* of a linear code C is its dimension as a \mathbb{F}_2 -vector space. A linear code of length n , rank k is called a (n, k) -code. If the minimum distance of the code is d , it is a (n, k, d) -code.

Let v_1, \dots, v_k be a basis of C . Then $C = \{\sum_{i=1}^k \lambda_i v_i : \lambda_1, \dots, \lambda_k \in \mathbb{F}_2\}$. So $|C| = 2^k$. So an (n, k) -code is a $[n, 2^k]$ -code. The information rate is k/n .

Definition. The *weight* of $x \in \mathbb{F}_2^n$ is $w(x) = d(x, 0)$.

Lemma 11.1. *The minimum distance of a linear code is the minimum weight of a non-zero code word.*

Proof. Obvious. □

Definition. For $x, y \in \mathbb{F}_2^n$, let $x \cdot y = \sum_{i=1}^n x_i y_i \in \mathbb{F}_2$. This operation is symmetric and bilinear, but not positive definite.

Definition. Take some $P \subseteq \mathbb{F}_2^n$. The *parity check code* defined by P is

$$C = \{x \in \mathbb{F}_2^n : p \cdot x = 0, \forall p \in P\}.$$

Examples.

- $P = \{(1, \dots, 1)\}$ gives the simple parity check code;
- $P = \{(1, 0, 1, 0, 1, 0, 1), (0, 1, 1, 0, 0, 1, 1), (0, 0, 0, 1, 1, 1, 1)\}$ gives Hamming's $[7, 4, 3]$ -code;
- Show that C^+ and C^- are linear if C is linear. When is C' linear?

Lemma 11.2. *Every parity check code is linear.*

Proof. Obvious. □

Definition. Take $C \subseteq \mathbb{F}_2^n$ a linear code. The *dual code*

$$C^\perp = \{x \in \mathbb{F}_2^n : x \cdot y = 0, \forall y \in C\}.$$

This is a parity-check code, so is linear. Note $C \cap C^\perp$ is not necessarily $\{0\}$.

Lemma 11.3. $\dim(C) + \dim(C^\perp) = n$.

Proof. See later. □

Corollary 11.4. *Let C be linear. Then $(C^\perp)^\perp = C$. In particular, C is a parity check code.*

Proof. Clearly $C \subseteq (C^\perp)^\perp$. Equality follows from the previous lemma. □

Definition. Let C be a (n, k) -code.

- (i) A *generator matrix* G for C is a $k \times n$ matrix with rows a basis of C ;
- (ii) A *parity check matrix* H for C is a generator matrix for C^\perp . It is a $(n - k) \times n$ matrix.

The codewords of C can be viewed either as:

- Linear combinations of rows of G ;
- Linear dependence relations between the columns of H , i.e $C = \{x \in \mathbb{F}_2^n : Hx = 0\}$.

Syndrome decoding

Definition. Let C be a (n, k) -code. The *syndrome* of $x \in \mathbb{F}_2^n$ is Hx . If we receive $x = c + z$, where c is a codeword and z is the ‘error pattern’, the syndrome $Hx = H(c + z) = Hc + Hz = Hz$. If C is e -error correcting, we precompute a list of all possible Hx for all z with $w(z) \leq e$. On receiving $x \in \mathbb{F}_2^n$ we look for Hx in our list. If successful, $Hx = Hz$ so $H(x - z) = 0$ and we decode x as $c = x - z \in C$ with $d(x, c) = w(z) \leq e$.

Definition. We say codes $C_1, C_2 \subseteq \mathbb{F}_2^n$ are *equivalent* if reordering each codeword of C_1 using the same permutation gives the codewords of C_2 . i.e $\sigma C_1 = C_2$ for some $\sigma \in S_n$. Usually we only consider codes up to equivalence.

Lemma 11.5. Every (n, k) -linear code is equivalent to one with generator $G = (I_k | B)$ for some $k \times (n - k)$ matrix B .

Proof. Given a $k \times n$ generator matrix G for C , using row operations, we can transform G into row-echelon form, i.e

$$G_{ij} = \begin{cases} 0 & j < l(i) \\ 1 & j = l(i) \end{cases}$$

for some $l(1) < l(2) < \dots < l(k)$. Permuting the columns replaces code by an equivalent code, so wlog we can assume $l(i) = i$ for all $1 \leq i \leq k$ which gives the required form upon applying further row operations. \square

Remark. A message $y \in \mathbb{F}_2^k$, viewed as a row vector is encoded as yG . So if $G = (I_k | B)$, then $yG = (y | yB)$, where y is the message and yB are check digits.

Now we will prove

Lemma 11.6. $\dim(C) + \dim(C^\perp) = n$.

Proof. Wlog C has generator matrix $G = (I_k | B)$. G has k linearly independent columns so the linear map $\gamma : \mathbb{F}_2^n \rightarrow \mathbb{F}_2^k$ defined by $x \mapsto Gx$ is surjective and $\text{Ker}(\gamma) = C^\perp$. By the rank-nullity theorem, $\dim \mathbb{F}_2^n = \dim \text{Ker}(\gamma) + \dim \text{Im}(\gamma) = n$, i.e $n = \dim(C) + \dim(C^\perp)$. \square

Lemma 11.7. An (n, k) -linear code with generator matrix $G = (I_k, B)$ has parity check matrix $H = (B^T | I_{n-k})$.

Proof. $GH^T = (I_k | B) \begin{pmatrix} B \\ I_{n-k} \end{pmatrix} = B + B = 0$. So the rows of H generate a subcode of C^\perp . But $\text{rank}(H) = n - k$ and $n - k = \dim C^\perp$. Hence C^\perp has generator matrix H , and H is a parity check matrix for C . \square

Lemma 11.8. Let C be a linear code with parity check matrix H . Then $d(C) = d$ if and only if

- (i) any $(d - 1)$ columns of H are linearly independent;
- (ii) some set of d columns of H are linearly dependent.

Proof. Exercise. \square

12 Examples of codes

Hamming codes

Definition. For $d \geq 1$, let $n = 2^d - 1$. Let H be the $d \times n$ matrix whose columns are the non-zero elements of \mathbb{F}_2^d . The *Hamming* $(n, n-d)$ -linear code is the code with parity matrix H (original code is $d = 3$).

Lemma 12.1. *The Hamming $(n, n-d)$ -code C has minimum distance $d(C) = 3$, is perfect and is 1-error correcting.*

Proof. Any two columns of H are linearly independent, but there are 3 columns that are linearly dependent, hence $d(C) = 3$ by the previous lemma. By a previous result, C is $\lfloor \frac{3-1}{2} \rfloor = 1$ -error correcting.

Note that $n = 2^d - 1$, $e = 1$ so $V(n, e) = 1 + 2^d - 1 = 2^d$ and $|C| = 2^{n-d} = \frac{2^n}{2^d} = \frac{2^n}{V(n, e)}$ and C is perfect. \square

Reed-Muller codes

Take $X = \{p_1, \dots, p_n\}$, set of size n . There is a correspondence between $\mathcal{P}(X)$ and \mathbb{F}_2^n via $\mathcal{P}(X) \rightarrow \{f : X \rightarrow \mathbb{F}_2\} \rightarrow \mathbb{F}_2^n$ the composition of $A \mapsto \mathbb{1}_A$ and $f \mapsto (f(p_1), \dots, f(p_n))$. The symmetric difference then corresponds to vector addition, and the intersection $A \cap B$ to the wedge product $x \wedge y = (x_1 y_1, \dots, x_n y_n)$.

Take $X = \mathbb{F}_2^d$, so $n = 2^d = |X|$. Let $v_0 = (1, 1, 1, \dots, 1)$. Let $v_i = \mathbb{1}_{H_i}$ where $H_i = \{p \in X : p_i = 0\}$ for $1 \leq i \leq d$.

Definition. Let $0 \leq r \leq d$. The *Reed-Muller code* $\text{RM}(d, r)$ of order r and length 2^d is the linear code spanned by v_0 and all wedge products of r or fewer of the v_i . By convention, the empty wedge product is v_0 .

Example 12.1. $d = 3$, $X = \mathbb{F}_2^3 = \{p_1, \dots, p_8\}$ in binary order.

X	000	001	010	011	100	101	110	111
v_0	1	1	1	1	1	1	1	1
v_1	1	1	1	1	0	0	0	0
v_2	1	1	0	0	1	1	0	0
v_3	1	0	1	0	1	0	1	0
$v_1 \wedge v_2$	1	1	0	0	0	0	0	0
$v_2 \wedge v_3$	1	0	0	0	1	0	0	0
$v_1 \wedge v_3$	1	0	1	0	0	0	0	0
$v_1 \wedge v_2 \wedge v_3$	1	0	0	0	0	0	0	0

Examples.

- $\text{RM}(3, 0)$ - span of v_0 , a repetition code of length 8;
- $\text{RM}(3, 1)$ - span of v_0, v_1, v_2, v_3 , equivalent to a parity check extension of Hamming's $(7, 4)$ -code;
- $\text{RM}(3, 2)$ - an $(8, 7)$ -code, actually simple parity check code of length 8;
- $\text{RM}(3, 3) = \mathbb{F}_2^8$, trivial code.

Theorem 12.2.

(i) The vectors $v_{i_1} \wedge \dots \wedge v_{i_s}$ where $1 \leq i_1 < i_2 < \dots, i_s \leq d$ and $0 \leq s \leq d$ form a basis for \mathbb{F}_2^n ;

(ii) $\text{rank}(\text{RM}(d, r)) = \sum_{s=0}^r \binom{d}{s}$.

Proof. In (i) we've listed $\sum_{s=0}^d \binom{d}{s} = (1+1)^d = 2^d = n$ vectors, so it is enough to show spanning, i.e. $\text{RM}(d, d) = \mathbb{F}_2^n$. Let $p \in X$. Put

$$y_i = \begin{cases} v_i & p_i = 0 \\ v_0 + v_i & p_i = 1 \end{cases}.$$

Then $\mathbb{1}_{\{p\}} = y_1 \wedge \dots \wedge y_d$. Expand this using the distributive law to see that $\mathbb{1}_{\{p\}} \in \text{RM}(d, d)$. Since $\{\mathbb{1}_{\{p\}}\}_{p \in X}$ clearly spans, we are done.

For (ii), $\text{RM}(d, r)$ is spanned by $v_{i_1} \wedge \dots \wedge v_{i_s}$ for $1 \leq i_1 < \dots < i_s \leq d$ with $0 \leq s \leq r$. These vectors are linearly independent by (i), so a basis. Hence $\text{rank}(\text{RM}(d, r)) = \sum_{s=0}^r \binom{d}{s}$. \square

Definition. Let C_1, C_2 be linear codes of length n with $C_2 \subseteq C_1$. The *bar product* is $C_1 | C_2 = \{(x | x+y) : x \in C_1, y \in C_2\}$. It is a linear code of length $2n$.

Lemma 12.3.

(i) $\text{rank}(C_1 | C_2) = \text{rank}(C_1) + \text{rank}(C_2)$;

(ii) $d(C_1 | C_2) = \min\{2d(C_1), d(C_2)\}$.

Proof. (i) If C_1 has basis x_1, \dots, x_k , C_2 has basis y_1, \dots, y_l , then $C_1 | C_2$ has basis $\{(x_i | x_i) : 1 \leq i \leq k\} \cup \{(0 | y_i) : 1 \leq i \leq l\}$. So $\text{rank}(C_1 | C_2) = k + l = \text{rank}(C_1) + \text{rank}(C_2)$;

(ii) Let $0 \neq (x | x+y) \in C_1 | C_2$. If $y \neq 0$ then $w(x | x+y) = w(x) + w(x+y) \geq w(y) \geq d(C_2)$. If $y = 0$ then $w(x | x+y) = 2w(x) \geq 2d(C_1)$. So $d(C_1 | C_2) \geq \min\{2d(C_1), d(C_2)\}$. But there exists $0 \neq x \in C_1$ with $w(x) = d(C_1)$, so $d(C_1 | C_2) \leq w(x | x) = 2d(C_1)$. There exists $0 \neq y \in C_2$ with $w(y) = d(C_2)$, so $d(C_1 | C_2) \leq w(0 | y) = w(y) = d(C_2)$. \square

Theorem 12.4.

- (i) $\text{RM}(d, r) = \text{RM}(d-1, r) \mid \text{RM}(d-1, r-1)$;
- (ii) $\text{RM}(d, r)$ has minimum distance 2^{d-r} .

Proof.

- (i) Exercise, noting $\text{RM}(d-1, r-1) \subseteq \text{RM}(d-1, r)$ so the bar product is defined;
- (ii) If $r = 0$ then $\text{RM}(d, 0)$ is a repetition code of length $n = 2^d$, having minimum distance 2^{d-0} . If $r = d$, then $\text{RM}(d, d) = \mathbb{F}_2^n$ having minimum distance $1 = 2^{d-d}$. We prove the $0 < r < d$ case by induction on d . Recall from (i) that $\text{RM}(d, r) = \text{RM}(d-1, r) \mid \text{RM}(d-1, r-1)$. By induction, the minimum distance of $\text{RM}(d-1, r)$ is 2^{d-1-r} and the minimum distance of $\text{RM}(d-1, r-1)$ is 2^{d-1-r} . Hence the minimum distance of $\text{RM}(d, r)$ is $\min\{2 \cdot 2^{d-1-r}, 2^{d-r}\} = 2^{d-r}$.

□

13 Cyclic codes

GRM Recap

Fact. Let F be a field. The rings \mathbb{Z} and $F[X]$ both have a division algorithm: if $a, b \in \mathbb{Z}$, $b \neq 0$, there exist $q, r \in \mathbb{Z}$ such that $a = qb + r$ and $0 \leq r < |b|$. If $f, g \in F[X]$, $g \neq 0$, there exist $q, r \in F[X]$ such that $f = qg + r$ with $\deg r < \deg g$.

Recall ideals I : subsets of a ring R such that whenever $r \in R$, $x \in I$ we have $rx \in I$ (e.g. $2\mathbb{Z} \trianglelefteq \mathbb{Z}$). Principle ideals generated by $r \in R$ are $(x) = \{rx : r \in R\}$. By the division algorithm, every ideal in \mathbb{Z} or $F[X]$ is principle.

If $I \trianglelefteq R$ then the set of cosets $R/I = \{x + I : x \in R\}$ is a ring, called the *quotient ring*. In particular, we identify $\mathbb{Z}/n\mathbb{Z}$ and $\{0, 1, \dots, n-1\}$ and reduce modulo n after each operation.

Similarly, $F[X]/(f(X)) \leftrightarrow \left\{ \sum_{i=0}^{n-1} a_i X^i : a_i \in F \right\} \leftrightarrow F^n$, reducing ‘modulo f ’ after each operation using the division algorithm.

Definition. A linear code $C \subseteq \mathbb{F}_2^n$ is *cyclic* if

$$(a_0, a_1, \dots, a_{n-1}) \in C \implies (a_{n-1}, a_0, a_1, \dots, a_{n-2}) \in C.$$

Identify $\mathbb{F}_2[X]/(X^n - 1)$ with \mathbb{F}_2^n via

$$a_0 + a_1X + \dots + a_{n-1}X^{n-1} + (X^n - 1) \leftrightarrow^\pi (a_0, \dots, a_{n-1}).$$

Lemma 13.1. A code $C \subseteq \mathbb{F}_2^n$ is cyclic if and only if $\pi(C) = \mathcal{C}$ satisfies

- (i) $0 \in \mathcal{C}$;
- (ii) $f, g \in \mathcal{C} \Rightarrow f + g \in \mathcal{C}$;
- (iii) $f \in \mathbb{F}_2[X], g \in \mathcal{C} \Rightarrow fg \in \mathcal{C}$.

In other words, \mathcal{C} is an ideal of $\mathbb{F}_2[X]/(X^n - 1)$.

Proof. If $g(X) = a_0 + a_1X + \dots + a_{n-1}X^{n-1} + (X^n - 1)$ then $Xg(X) = a_{n-1} + a_0X + \dots + a_{n-2}X^{n-1} + (X^n - 1)$. So \mathcal{C} is cyclic iff (i) and (ii) hold (linearity) and $g(X) \in \mathcal{C} \Rightarrow Xg(X) \in \mathcal{C}$. Note that closure under multiplication by X implies closure under multiplication by any $f(X)$, so we have the result. \square

From now on, we identify $\pi(C)$ and \mathcal{C} .

Basic problem:

$$\begin{aligned} \text{cyclic codes of length } n &\leftrightarrow \text{ideals in } \mathbb{F}_2[X]/(X^n - 1) \\ &\leftrightarrow \text{ideals in } \mathbb{F}_2[X] \text{ containing } (X^n - 1) \\ &\leftrightarrow \text{polynomials } g(X) \in \mathbb{F}_2[X] \text{ with } g(X) \mid (X^n - 1) \end{aligned}$$

Theorem 13.2. Let $C \leq \mathbb{F}_2[X]/(X^n - 1)$ be a cyclic code. Then there exists a unique polynomial $g(X)$ such that

- (i) $C = \{f(X)g(X) + (X^n - 1) : f(X) \in \mathbb{F}_2[X]\} = (g)$;
- (ii) $g(X) \mid X^n - 1$.

In particular, $p(X) \in \mathbb{F}_2[X]$ represents a codeword $g(X) \mid p(X)$. We say $g(X)$ is the generator polynomial of C .

Proof.

- (i) Let $g(X) \in \mathbb{F}_2[X]$ be of least degree representing a non-zero codeword. Now $\deg g \leq n - 1$. Since C is cyclic, we have $(g) \subseteq C$.

Let $p(X) \in \mathbb{F}_2[X]$ represent a codeword. By the division algorithm, $p(X) = q(X)g(X) + r(X)$ where $\deg r < \deg g$. So $r(X) = p(X) - q(X)g(X) \in C$, contradicting minimality of $\deg g$ unless $r = 0$. Hence $g(X) \mid p(X)$ and so $C \subseteq (g)$, i.e. $C = (g)$.

(ii) Take $p(X) = X^n - 1$, giving (ii).

Uniqueness: suppose $g_1(X), g_2(X)$ both satisfy (i) and (ii). Then $g_1(X) \mid g_2(X)$ and $g_2(X) \mid g_1(X)$, i.e. $g_1(X), g_2(X)$ are equal up to a unit. But the only unit in $\mathbb{F}_2[X]$ is 1, so $g_1(X) = g_2(X)$. \square

Lemma 13.3. *Let C be a cyclic code of length n with generator polynomial $g(X) = a_0 + a_1X + \dots + a_kX^k$ ($a_k \neq 0$). Then C has a basis given by $g(X), Xg(X), \dots, X^{n-k-1}g(X)$. In particular, $\text{rank}(C) = n - k$.*

Proof. Trivial check. \square

Corollary 13.4. *Given generator polynomial of degree k as above, a generator matrix for C is given by*

$$G = \begin{pmatrix} a_0 & a_1 & \dots & a_k & 0 & \dots & 0 \\ 0 & a_0 & \dots & a_{k-1} & a_k & \dots & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ 0 & 0 & \dots & a_0 & a_1 & \dots & a_k \end{pmatrix}$$

a $(n - k) \times n$ matrix.

Corollary 13.5. *If $g(X)h(X) = X^n - 1$ where $g(X)$ is the generator polynomial, then writing $h(X) = b_0 + b_1X + \dots + b_{n-k}X^{n-k}$ ($b_{n-k} \neq 0$), h is called the parity check polynomial and the parity check matrix is*

$$H = \begin{pmatrix} b_{n-k} & b_{n-k-1} & \dots & b_0 & 0 & \dots & 0 \\ 0 & b_{n-k} & \dots & b_1 & b_0 & \dots & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ 0 & 0 & \dots & b_{n-k} & b_{n-k-1} & \dots & b_0 \end{pmatrix}$$

a $k \times n$ matrix.

Proof. Check that the “dot product” of the i th row of G and the j th row of H is the coefficient of $X^{(n-k-i)+j}$ in $g(X)h(X) = X^n - 1$. But $1 \leq i \leq n - k$ and $1 \leq j \leq k$ so $0 < (n - k - i) + j < n$ so all these coefficients are 0. Hence the row of G and H are “orthogonal”. Also as $b_{n-k} \neq 0$, $\text{rank}(H) = k = \text{rank}(C^\perp)$, so H is the parity check matrix. \square

Remarks. The cyclic code given by $h(X)$ is the dual code C^\perp with coefficients ‘reversed’. (Given a polynomial $f(X) = \sum_{i=0}^m f_iX^i$ of degree m in $\mathbb{F}_2[X]$, the reverse polynomial is $\check{f}(X) = \sum_{i=0}^m f_{m-i}X^i = X^{\deg f} f(1/X)$.)

Lemma 13.6. *If n is odd, then $X^n - 1 = f_1(X) \dots f_t(X)$ with $f_i(X)$ distinct irreducibles in $\mathbb{F}_2[X]$ (false if n is even, e.g. $X^2 - 1 = (X - 1)^2$). Thus there exist 2^t cyclic codes of length n .*

Proof. Not given (see II Galois Theory). \square

14 BCH codes

These are a certain type of cyclic code.

Fact (A). Suppose p prime, $\mathbb{F}_p = \mathbb{Z}/p\mathbb{Z}$ is a field. If $f(X) \in \mathbb{F}_p[X]$ is irreducible, then $K = \mathbb{F}_p[X]/(f)$ is a field of order $p^{\deg f}$.

Fact (B). Let $q = p^\alpha$ for p prime, $\alpha \in \mathbb{N}$. Then there exists a field \mathbb{F}_q of order q and it is unique up to isomorphism.

Fact (C). The multiplicative group $\mathbb{F}_q^\times = \mathbb{F}_q \setminus \{0\}$ is cyclic. A generator of this group is called a *primitive element*.

Notation: n an odd integer. Pick $r \geq 1$ such that $2^r \equiv 1 \pmod{n}$ (exists since $(n, 2) = 1$). Let $K = \mathbb{F}_{2^r}$ and let $\mu_n(K) = \{x \in K : x^n = 1\} \leq K^\times$. Then $\mu_n(K)$ is a cyclic group of order n (as K^\times is cyclic and $n \mid (2^r - 1)$). So $\mu_n(K)$ is generated by some $\alpha \in K$, called a *primitive n th root of unity*.

Definition. The *cyclic code of length n with defining set $A \subseteq \mu_n(K)$* is

$$C = \{f(x) + (X^n - 1) : f(x) \in \mathbb{F}_2[X], f(a) = 0 \forall a \in A\}.$$

The *generator polynomial* $g(x)$ is the non-zero polynomial of least degree such that $g(a) = 0$ for all $a \in A$. Equivalently, the lcm of the minimal polynomials of the elements $a \in A$.

Definition. The cyclic code with defining set $\{\alpha, \alpha^2, \dots, \alpha^{\delta-1}\}$ for $\alpha \in \mu_n(K)$ is called a BCH (Bose, Ray-Chandhari, Hocquenghem) with *design distance* δ .

Theorem 14.1. A BCH code C with design distance δ has $d(C) \geq \delta$.

We first prove

Lemma 14.2 (Vandermonde).

$$\det \begin{pmatrix} 1 & 1 & 1 & \dots & 1 \\ x_1 & x_2 & x_3 & \dots & x_n \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ x_1^{n-1} & x_2^{n-1} & x_3^{n-1} & \dots & x_n^{n-1} \end{pmatrix} = \prod_{1 \leq j < i \leq n} (x_i - x_j).$$

Proof. Exercise. □

Proof of theorem. Consider

$$\begin{pmatrix} 1 & \alpha & \alpha^2 & \dots & \alpha^{n-1} \\ 1 & \alpha^2 & \alpha^4 & \dots & \alpha^{2(n-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \alpha^{\delta-1} & \alpha^{2(\delta-1)} & \dots & \alpha^{(\delta-1)(n-1)} \end{pmatrix}.$$

By the previous lemma, any $\delta - 1$ columns of H are linearly independent. But any codeword of C is a dependence relation between the columns of H . Hence any non-zero codeword has weight $\geq \delta$, and $d(C) \geq \delta$. □

Remark. H in the above proof is *not* a parity-check matrix: its entries are not in \mathbb{F}_2 .

Decoding BCH codes

Let C be a cyclic code with defining set $\{\alpha, \alpha^2, \dots, \alpha^{\delta-1}\}$ where $\alpha \in K$ is a primitive n th root of unity. By the above theorem, we should be able to correct $t := \lfloor \frac{\delta-1}{2} \rfloor$ errors.

We send $c \in C$ and receive $r = c + e$, where e is the error pattern with $\leq t$ non-zero errors. Note

$$\mathbb{F}_2^n \leftrightarrow \mathbb{F}_2[X]/(X^n - 1)$$

$$r, c, e \leftrightarrow r(X), c(X), e(X).$$

Note also $c(\alpha^j) = 0$ for each $1 \leq j \leq \delta - 1$. Hence $r(\alpha^j) = e(\alpha^j)$.

Definition. The *error locator polynomial* is $\sigma(X) = \prod_{i \in \mathcal{E}} (1 - \alpha^i X) \in K[X]$ where $\mathcal{E} = \{i : e_i = 1, 0 \leq i \leq n - 1\}$.

Problem: assuming $\deg \sigma = |\mathcal{E}|$, where $2t + 1 \leq \delta$, recover $\sigma(X)$ from $r(X)$.

Theorem 14.3. Assume $\deg \sigma = |\mathcal{E}| \leq t$. Then $\sigma(X)$ is the unique polynomial in $K[X]$ of least degree such that

$$(i) \quad \sigma(0) = 1;$$

$$(ii) \quad \sigma(X) \sum_{j=1}^{2t} r(\alpha^j) X^j \equiv \omega(X) \pmod{X^{2t+1}} \text{ for some } \omega(X) \in K[X] \text{ of degree at most } t.$$

Proof. Define $\omega(X) = -X\sigma'(X)$, the error co-locator. Then

$$\omega(X) = \sum_{i \in \mathcal{E}} \alpha^i X \prod_{j \in \mathcal{E} \setminus \{i\}} (1 - \alpha^j X).$$

Work in the ring of formal power series $K[[X]] = \{\sum_{i=0}^{\infty} \beta_i X^i : \beta_i \in K\}$. Then

$$\begin{aligned} \frac{\omega(X)}{\sigma(X)} &= \sum_{i \in \mathcal{E}} \frac{\alpha^i X}{1 - \alpha^i X} = \sum_{i \in \mathcal{E}} \sum_{j=1}^{\infty} (\alpha^i X)^j \\ &= \sum_{j=1}^{\infty} \left(\sum_{i \in \mathcal{E}} (\alpha^i)^j \right) X^j \\ &= \sum_{j=1}^{\infty} e(\alpha^j) X^j \end{aligned}$$

so $\sigma(X) \sum_{j=1}^{\infty} e(\alpha^j) X^j = \omega(X)$. By the definition of c , $c(\alpha^j) = 0$ for all $1 \leq j \leq \delta - 1$ and so for all $1 \leq j \leq 2t$. As $r = c + e$, $r(\alpha^j) = e(\alpha^j)$ for all $1 \leq j \leq 2t$. Hence

$$\sigma(X) \sum_{j=1}^{2t} r(\alpha^j) X^j \equiv \omega(X) \pmod{X^{2t+1}}.$$

This verifies (i),(ii) with $\omega(X) = -X\sigma'(X)$, so $\deg \omega = \deg \sigma = |\mathcal{E}| \leq 2t$.

Now we show uniqueness: suppose we have $\tilde{\sigma}, \tilde{\omega} \in K[X]$ which also satisfy (i) and (ii). Without loss of generality $\deg \tilde{\sigma} \leq \deg \sigma$. Now, $\sigma(X)$ has distinct non-zero roots, so $\sigma(X)$ and $\omega(X) = -X\sigma'(X)$ are coprime. By (ii) $\tilde{\sigma}\omega(X) = \sigma(X)\tilde{\omega}(X) \pmod{X^{2t+1}}$, so $\sigma(X)\tilde{\omega}(X) = \tilde{\sigma}(X)\omega(X)$ since $\sigma, \tilde{\sigma}, \omega, \tilde{\omega}$ all have degree $\leq t$. But since σ, ω are coprime we must have $\sigma(X) \mid \tilde{\sigma}(X)$. Since $\deg \tilde{\sigma} \leq \deg \sigma$, we have $\tilde{\sigma}(X) = \lambda\sigma(X)$ for some $\lambda \in K$. But (i) says $\lambda = 1$ so $\sigma = \tilde{\sigma}$. \square

Decoding BCH codes

We receive $r(X)$

- Compute $\sum_{j=1}^{2t} r(\alpha^j)X^k$;
- Set $\sigma(X) = 1 + \sigma_1X + \dots + \sigma_tX^t$ and compute coefficients of X^i for $t+1 \leq i \leq 2t$ to obtain linear equations for $\sigma_1, \dots, \sigma_t$ (i.e. $\sum_{j=0}^t \sigma_j r(\alpha^{i-j}) = 0$);
- Solve these over K , keeping solutions of least degree;
- Compute $\mathcal{E} = \{0 \leq i \leq n-1 : \sigma(\alpha^{-i}) = 0\}$ and check $|\mathcal{E}| = \deg \sigma$;
- Set $e(X) = \sum_{i \in \mathcal{E}} X^i$. Then $c(X) = r(X) + e(X)$.

Examples.

- (i) $n = 7$, $X^7 - 1 = (X + 1)(X^3 + X + 1)(X^3 + X^2 + 1)$ in $\mathbb{F}_2[X]$. Let $g(X) = X^3 + X + 1$ so $h(X) = (X + 1)(X^3 + X^2 + 1) = X^4 + X^2 + X + 1$. Parity check matrix

$$\begin{pmatrix} 1 & 0 & 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 0 & 1 & 1 & 1 \end{pmatrix}$$

with columns running over $\mathbb{F}_2^3 \setminus \{0\}$, the Hamming (7,4)-code.

- (ii) Let K be a splitting field for $X^7 - 1 \in \mathbb{F}_2[X]$, e.g $K = \mathbb{F}_{2^3}$. Let $\beta \in K$ be a root of $g(X) = X^3 + X + 1$. Note $\beta^3 = \beta + 1$ so $\beta^6 = \beta^2 + 1$ and $g(\beta^2) = 0$. So BCH code C defined by $\{\beta, \beta^2\}$ has generator polynomial $g(X)$, so it is Hamming's (7,4)-code. By theorem 14.1, $d(C) \geq 3$.
- (iii) CD players use BCH codes.

15 Shift registers

Definition. A (general) feedback shift register (FSR) is a map $f : \mathbb{F}_2^d \rightarrow \mathbb{F}_2^d$ given by $f(x_0, \dots, x_{d-1}) = (x_1, x_2, \dots, x_{d-1}, C(x_0, \dots, x_{d-1}))$ where $C : \mathbb{F}_2^d \rightarrow \mathbb{F}_2$. We say the register has length d . The stream associated to an initial fill $(y_0, y_1, \dots, y_{d-1})$ is the infinite sequence $(y_0, y_1, \dots, y_j, \dots)$ with $y_n = C(y_{n-d}, y_{n-d+1}, \dots, y_{n-1})$ for all $n \geq d$.

Definition. The function f as above is a linear feedback shift register (LFSR) if

$$C(x_0, x_1, \dots, x_{d-1}) = \sum_{i=0}^{d-1} a_i x_i$$

for some $a_i \in \mathbb{F}_2$ (usually $a_0 = 1$). The stream produced by a LFSR is a recurrence relation $y_n = \sum_{i=0}^{d-1} a_i y_{n-d+i}$ for $n \geq d$, over \mathbb{F}_2 with auxiliary polynomial $P(X) = X^d + a_{d-1}X^{d-1} + \dots + a_1X + a_0$.

Definition. The feedback polynomial is $\check{P}(X) = a_0X^d + \dots + a_{d-1}X + 1$. A sequence y_0, y_1, \dots of elements of \mathbb{F}_2 is said to have generating function $G(X) = \sum_{j=0}^{\infty} y_j X^j \in \mathbb{F}_2[[X]]$.

Theorem 15.1. The stream $(y_n)_{n \geq 0}$ comes from a LFSR with auxiliary polynomial $P(X)$ if and only if its generating function $G(X)$ is (formally) of the form $A(X)/\check{P}(X)$ with $A \in \mathbb{F}_2[X]$ such that $\deg A < \deg \check{P}$ [note $\check{P}(X) = X^{\deg P} P(X^{-1})$].

Proof. Take $P(X), \check{P}(X)$ as above. We require $\left(\sum_{j=0}^{\infty} y_j X^j\right) \left(\sum_{i=0}^d a_{d-i} X^i\right)$ to be a degree strictly less than d . This holds precisely when the coefficient of

X^n in $G(X)\tilde{P}(X)$ is 0 for all $n \geq d$, i.e

$$\begin{aligned} \sum_{i=0}^d a_{d-i} y_{n-i} &= 0 \quad \forall n \geq d \\ \iff y_n &= \sum_{i=0}^{d-1} a_i y_{n-d+i} \quad \forall n \geq d \\ \iff (y_n)_{n \geq 1} &\text{ comes from a LFSR with auxiliary polynomial } P. \end{aligned}$$

□

Remark. The two problems

- to recover the LFSR from its stream;
- decoding BCH codes

both involve writing a power series as a ratio of two polynomials.

The Berlekamp-Massey Method

Let $(x_n)_{n \geq 1}$ be the output of a binary LFSR. Can we find the (unknown) d and a_0, \dots, a_{d-1} such that

$$x_n + \sum_{i=1}^d a_{d-i} x_{n-i} = 0 \quad \forall n \geq d?$$

We have

$$\begin{pmatrix} x_d & x_{d-1} & \cdots & x_1 & x_0 \\ x_{d+1} & x_d & \cdots & x_2 & x_1 \\ \cdots & \cdots & \cdots & \cdots & \cdots \\ x_{2d} & x_{2d-1} & \cdots & x_{d+1} & x_d \end{pmatrix} \begin{pmatrix} a_d \\ a_{d-1} \\ \vdots \\ a_0 \end{pmatrix}. \quad (*)$$

Look successively at matrices

$$A_0 = (a_0), A_1 = \begin{pmatrix} x_1 & x_0 \\ x_2 & x_1 \end{pmatrix}, A_2 = \begin{pmatrix} x_2 & x_1 & x_0 \\ x_3 & x_2 & x_1 \\ x_4 & x_3 & x_2 \end{pmatrix}, \dots$$

starting from A_r if we happen to know that $d \geq r$. For each A_i compute $\det A_i$.

- If $\det A_i \neq 0$, then $d \neq i$ so replace i with $i + 1$ and repeat;
- if $\det A_i = 0$, solve $(*)$ on the assumption $d = i$. Check this candidate for a_0, \dots, a_{d-1} over as many terms of the stream as you wish. If this fails, we know $d > i$ so replace i with $i + 1$ and repeat.

Cryptography

16 Introduction and examples

We want to modify a message such that it becomes unintelligible to an ‘enemy’, E. Certain (secret) information is shared by A&B, called the *key*, from \mathcal{K} . The unencrypted message is called the *plaintext* from \mathcal{M} . The encrypted message is called *ciphertext* \mathcal{C} . A *cryptosystem* consists of $(\mathcal{K}, \mathcal{M}, \mathcal{C})$ together with functions

$$e : \mathcal{M} \times \mathcal{K} \rightarrow \mathcal{C} \quad (\text{encryption})$$

$$d : \mathcal{C} \times \mathcal{K} \rightarrow \mathcal{M} \quad (\text{decryption})$$

such that $d(e(m, k), k) = m$ for all $m \in \mathcal{M}$, $k \in \mathcal{K}$.

Examples. Suppose $\mathcal{M} = \mathcal{C} = \{A, B, \dots, Z\}^* = \Sigma^*$

- (i) Simple substitution: $\mathcal{K} = \{\text{permutations of } \Sigma\}$. Each letter of plaintext is replaced by the image under the permutation.
- (ii) Vigenère cipher: $\mathcal{K} = \Sigma^d$ for some $d \in \mathbb{N}$. Identify Σ and $\mathbb{Z}/26\mathbb{Z}$. Write out the key repeatedly below the message, then add modulo 26. For example:

ATTACKATDAWN
LEMONLEMONLE
LXFOPVEFRNHR.

- (iii) Caesar cipher: take $d = 1$ in the above (key is a single letter).

Breaking a cryptosystem

E might know:

- the functions e, d ;
- the probability distributions on \mathcal{M}, \mathcal{K}

but not the key. E seeks to recover plaintexts from ciphertexts. 3 possible attacks:

- Level 1: *ciphertext-only*. E only knows some piece of ciphertext;
- Level 2: *known-plaintext*: E possesses a considerable length of plaintext and matching ciphertext and seeks the key (knows $m, e(m, k)$ but not k);
- Level 3: *chosen plaintext*: E may acquire the ciphertext for any message she chooses (can generate $e(m, k)$ for any m).

Remarks.

- The previous examples fail at Level 1 if the messages are in English (highly predictable) and sufficiently long. Even for random plaintext, both can fail at Level 2.
- For modern applications, Level 3 is desirable.

Consider the cryptosystem $(\mathcal{M}, \mathcal{K}, \mathcal{C})$. Model the keys and messages as independent random variables K, M taking values in \mathcal{K}, \mathcal{M} respectively. Then $C := e(K, M) \in \mathcal{C}$ is also random.

Definition. Say $(\mathcal{M}, \mathcal{K}, \mathcal{C})$ has *perfect secrecy* if $H(M|C) = H(M)$, i.e M and C are independent.

Exercise: show that perfect secrecy implies $|\mathcal{K}| \geq |\mathcal{M}|$.

Definition.

1. The *message equivocation* is $H(M|C)$;
2. The *key equivocation* is $H(K|C)$.

Lemma 16.1. $H(M|C) \leq H(K|C)$.

Proof. Note $M = d(C, K)$, hence $H(M|C, K) = 0$, $H(C, K) = H(M, C, K)$. Hence

$$\begin{aligned} H(K|C) &= H(K, C) - H(C) = H(M, C, K) - H(M|K, C) - H(C) \\ &= H(M, K, C) - H(C) \\ &= H(K|M, C) + H(M, C) - H(C) \\ &= H(K|M, C) + H(M|C) \leq H(M|C). \end{aligned}$$

□

Let $\mathcal{M} = \mathcal{C} = \mathcal{A}$. Send n messages $M^{(n)} = (M_1, \dots, M_n)$ encrypted as $C^{(n)} = (C_1, \dots, C_n)$ using the same key.

Definition. The *unicity distance* is the least n such that $H(K|C^{(n)}) = 0$, i.e the smallest number of encrypted messages required to uniquely determine the key.

Now,

$$\begin{aligned} H(K|C^{(n)}) &= H(K, C^{(n)}) - H(C^{(n)}) = H(K, M^{(n)}, C^{(n)}) - H(C^{(n)}) \\ &= H(K, M^{(n)}) - H(C^{(n)}) \\ &= H(K) + H(M^{(n)}) - H(C^{(n)}) \end{aligned}$$

since $K, M^{(n)}$ are independent.

We assume:

- All keys are equally likely ($H(K) = \log |\mathcal{K}|$);
- $H(M^{(n)}) \sim nH$ for some H constant and all n sufficiently large;
- All sequences of ciphertext are equally likely ($H(C^{(n)}) = n \log |\mathcal{A}|$).

Therefore $H(K|C^{(n)}) = \log |\mathcal{K}| + nH - n \log |\mathcal{A}| \geq 0$, or equivalently

$$n \leq U := \frac{\log |\mathcal{K}|}{\log |\mathcal{A}| - H}$$

where U is the unicity distance. Letting $R = 1 - \frac{H}{\log |\mathcal{A}|}$ we have $U = \frac{\log |\mathcal{K}|}{R}$ and R is called the *redundancy*.

Recall $0 \leq H \leq \log |\mathcal{A}|$. To make the unicity distance large we can make $|\mathcal{K}|$ large or use a message source with small redundancy.

17 The one-time pad and stream ciphers

Consider streams in \mathbb{F}_2 :

- A *plaintext stream* is p_0, p_1, p_2, \dots ;
- A *key stream* is k_0, k_1, k_2, \dots ;
- A *ciphertext stream* is z_0, z_1, z_2, \dots where $z_n = p_n + k_n$ for all $n \geq 0$.

Definition. A *one-time pad* is where k is generated uniformly at random, i.e $k_i = 0$ or 1 with probability $1/2$ and the k_i are iid. Hence $z = p + k$ gives a stream of iid random variables with $z_i = 0$ or 1 with probability $1/2$. So without the key stream, deciphering is impossible.

Exercise: a one time pad has perfect secrecy.

Problems:

- How to generate the k_i 's? This is hard in practice.
- How to share the key stream? Same as the initial problem!

In most applications, the one-time pad is not practical. Instead, share k_0, k_1, \dots, k_{d-1} using a shared feedback shift register (FSR) of length d to generate k . Then apply the following:

Lemma 17.1. *Let x_0, x_1, x_2, \dots be a stream produced by a FSR of length d . Then there exist $M, N \leq 2^d$ such that $x_{N+r} = x_r$ for all $r \geq M$, i.e x is eventually periodic.*

Proof. Let the register be $f : \mathbb{F}_2^d \rightarrow \mathbb{F}_2^d$. Let $v_i = (x_i, x_{i+1}, \dots, x_{i+d-1})$, then $v_{i+1} = f(v_i)$. Since $|\mathbb{F}_2^d| = 2^d$, the vectors v_0, \dots, v_{2^d} cannot all be distinct, so there exist $0 \leq a < b \leq 2^d$ such that $v_a = v_b$. Let $M = a$ and $N = b - a$ so $v_m = v_{M+N}$ and $v_r = v_{r+N}$ for all $r \geq M$ (induction), so $x_r = x_{r+N}$ for all $r \geq M$. \square

Remarks.

1. The maximum period of a FSR of length d is 2^d ;
2. The maximum period of a LFSR of length d is 2^{d-1} (Example Sheet 4 Q12);
3. Stream ciphers using LFSR fail at Level 2 due to the Berlekamp-Massey method;
4. Why should this cryptosystem be used?
 - It's cheap, fast and easy to use;
 - Messages are encrypted & decrypted 'on the fly' (codeword-by-codeword);
 - It's error tolerant.

Recall the stream produced by LFSR

$$a_d x_n = x_n = \sum_{i=0}^d a_{d-i} x_{n-i} \quad \forall n \geq d$$

has auxilliary polynomial $P(X) = X^d + a_{d-1}X^{d-1} + \dots + a_0$. Solutions of recursions are linear combinations of powers of roots of P . Over \mathbb{C} , the general solution is a linear combination of $\alpha^n, n\alpha^n, \dots, n^{t-1}\alpha^n$ for α a root of $P(X)$ with multiplicity t . Beware that $n^2 = n$ in \mathbb{F}_2 . Over \mathbb{F}_2 we need two modifications:

- Work in a splitting field K for $P(X) \in \mathbb{F}_2[X]$;
- Replace $n^i \alpha^n$ by $\binom{n}{i} \alpha^n$.

Can also generate new key streams from old ones:

Lemma 17.2. *Let $(x_n), (y_n)$ be outputs from LFSR's of length M, N respectively.*

- (i) *The sequence $(x_n + y_n)$ is the output of a LFSR of length $M + N$;*
- (ii) *The sequence $(x_n y_n)$ is the output of a LFSR of length MN .*

The following proof is ***non-examinable***:

Proof. Assume the auxilliary polynomials $P(X), Q(X)$ each have distinct roots, say $\alpha_1, \dots, \alpha_M$ and β_1, \dots, β_N in some extension $K \supseteq \mathbb{F}_2$. Then $x_n = \sum_{i=1}^M \lambda_i \alpha_i^n$, $y_n = \sum_{j=1}^N \mu_j \beta_j^n$, $\lambda_i, \mu_j \in K$.

- (i) $x_n + y_n = \sum_{i=1}^M \lambda_i \alpha_i^n + \sum_{j=1}^N \mu_j \beta_j^n$. This is produced by a LFSR with auxilliary polynomial $P(X)Q(X)$;
- (ii) $x_n y_n = \sum_{i=1}^M \sum_{j=1}^N \lambda_i \mu_j (\alpha_i \beta_j)^n$ is the output of a LFSR with auxilliary polynomial $\prod_{i=1}^M \prod_{j=1}^N (X - \alpha_i \beta_j)$.

□

Conclusions:

- Adding the output of two LFSR's is no more economical than producing the same string with a single LFSR;
- Multiplying streams looks promising until we realise $x_n y_n = 0$ 75% of the time;
- Non-linear FSR's look appealing but in general they are hard to analyse. In particular, E may understand them better than you do.

Remark. Stream ciphers are examples of *symmetric* cryptosystems - decryption is the same, or easily deduced from, the encryption algorithm. Now we consider *asymmetric* cryptosystems.

18 Assymmetric ciphers

The key is split into two parts:

- The *private key* for decryption;
- The *public key* for encryption.

Knowing the encryption & decryption algorithms and the public key, it should still be hard to find the private key or to decrypt messages. This aim implies secrecy at Level 3. Also no key exchange problem.

Base the system on mathematical problems that are believed to be ‘hard’. Two examples:

- Factoring*. Let $N = pq$ for p, q large primes. Given N , find p and q ;
- Discrete logarithms*. Let p be a large prime and g a primitive root modulo p (i.e a generator of \mathbb{F}_p^\times). Given x , find a such that $x \equiv g^a \pmod{p}$.

Definition. An algorithm *runs in polynomial time* if

$$\# \text{ operations} \subseteq c \cdot (\text{input size})^d$$

for some constants c, d .

Examples.

- An algorithm for factoring N has input size $\log_2 N$ (# binary digits).

Polynomial time algorithms include:

- Arithmetic of integers (+, −, ×, division algorithm);
- Computation of GCD (Euclidean algorithm);
- Modular exponentiation: i.e compute $x^a \pmod{N}$ using repeated squaring;
- Primality testing (Agrawal, Kayal, Saxena; 2002)

Elementary methods for computing factoring and discrete logarithms take much longer than polynomial time:

- Dividing by successive primes up to \sqrt{N} : takes time $\mathcal{O}(\sqrt{N}) = \mathcal{O}(2^{B/2})$ where $B = \log_2 N$;
- ‘Baby-step, giant-step’ (Shanks): set $m = \lceil \sqrt{p} \rceil$, write $a = qm + r$, $0 \leq q, r < m$. Then $x \equiv g^a \equiv g^{qm+r} \pmod{p}$ iff $g^{qm} \equiv g^{-r}x \pmod{p}$. So list $g^{qm} \pmod{p}$ for $q = 0, \dots, m-1$ and $g^{-r}x \pmod{p}$ for $r = 0, \dots, m-1$. Sort these two lists and look for a match. This takes $\mathcal{O}(\sqrt{p} \log p)$.

The best known method for solving these problems uses a factor base called the number field sieve. It has running time $\mathcal{O}(e^{c(\log N)^{1/3}(\log \log N)^{2/3}})$, where c is a known constant.

Modular arithmetic

Recall $\varphi(n) = |\{1 \leq a \leq n : (a, n) = 1\}| = |(\mathbb{Z}/n\mathbb{Z})^\times|$, the *Euler totient function*. Lagrange's theorem gives $a^{\varphi(N)} \equiv 1 \pmod{N}$ for each a with $(a, N) = 1$, the Euler-Fermat theorem. With $N = p$, this says $a^{p-1} \equiv 1 \pmod{p}$ for all a with $(a, p) = 1$, Fermat's little theorem.

Lemma 18.1. *Let $p = 4k - 1$ be prime for some $k \in \mathbb{Z}$. Let $d \in \mathbb{Z}$. If the equation $x^2 \equiv d \pmod{p}$ has a solution then $x \equiv d^k \pmod{p}$ is a solution.*

Proof. Let x_0 be a solution. Without loss of generality we may assume $x_0 \not\equiv 0 \pmod{p}$. Then $d^{2k-1} \equiv x_0^{2(2k-1)} \equiv x_0^{p-1} \equiv 1 \pmod{p}$ by Fermat's little theorem. Then $(d^k)^2 \equiv d \pmod{p}$. \square

19 Examples of publickey cryptography

Rabin cryptosystem

Private key: 2 large distinct primes $p, q \equiv 3 \pmod{4}$.

Public key: $N = pq$.

$\mathcal{M} = \mathcal{C} = \{1, \dots, N-1\} = \mathbb{Z}_N^\times$. Encrypt $m \in \mathcal{M}$ as $c = m^2 \pmod{N}$. The ciphertext is c . Usually we restrict so that $(m, N) = 1$ and $m > \sqrt{N}$.

Recieve c . Use the above lemma to solve for x_1, x_2 such that $x_1^2 \equiv c \pmod{p}$ and $x_2^2 \equiv c \pmod{q}$. Then use the Chinese remainder theorem (CRT) to find x such that $x = x_1 \pmod{p}$, $x = x_2 \pmod{q}$ so $x^2 \equiv c \pmod{N}$. Indeed, running the Euclidean algorithm on p, q gives r, s such that $rp + sq = 1$. Take $x = (sq)x_1 + (rp)x_2$.

Lemma 19.1.

- (i) *Let p be an odd prime, $(d, p) = 1$. Then $x^2 \equiv d \pmod{p}$ has either zero or two solutions.*
 - (ii) *Let $N = pq$ for p, q distinct odd primes, $(d, N) = 1$. Then $x^2 \equiv d \pmod{N}$ has zero or four solutions.*
1. We have $x^2 \equiv y^2 \iff p \mid (x+y)(x-y) \iff p \mid (x+y) \text{ or } p \mid x-y \iff x \equiv \pm y \pmod{p}$.
 2. If x_0 is a solution then by the CRT there exist solutions x with $x \equiv \pm x_0 \pmod{p}$, $x \equiv \pm x_0 \pmod{q}$ for any of the four choices of sign. By (i) these are the only possibilities.

To decrypt Rabin, find all 4 solutions to $x^2 \equiv c \pmod{N}$. Messages should include enough redundancy that only one of these possibilities makes sense.

Theorem 19.2. *Breaking the Rabin cryptosystem is essentially as difficult as factoring N .*

Proof. We've already seen that factoring N allows us to decrypt messages. Conversely, suppose we have an algorithm for computing square roots modulo N . Pick $x \pmod{N}$ at random. Use the algorithm to find a y for which $y^2 \equiv x^2 \pmod{N}$. With probability $1/2$, $x \equiv y \pmod{N}$. Then $(N, x - y)$ is a non-trivial factor of N . If this fails, start again with another x . After r trials, $\mathbb{P}(\text{failure}) = 2^{-r} \rightarrow 0$. \square

Suppose $N = pq$ where p, q are distinct odd primes. We claim that if we know some multiple m of $\varphi(N) = (p-1)(q-1)$, then factoring N is 'easy'.

Write $o_p(x) = \text{ord}(x)$ in $(\mathbb{Z}/p\mathbb{Z})^\times$. Write $m = 2^a b$ where $a \geq 1$, b odd. Let

$$X = \{x \in (\mathbb{Z}/N\mathbb{Z})^\times : o_p(x^b) \neq o_q(x^b)\}.$$

Theorem 19.3.

(i) *If $x \in X$ then there exists t with $0 \leq t < a$ such that $(x^{2^t b} - 1, N)$ is a non-trivial factor of N (i.e. p or q).*

(ii) $|X| \geq \frac{1}{2} |(\mathbb{Z}/N\mathbb{Z})^\times| = \frac{1}{2} \varphi(N)$.

Proof.

(i) By Euler-Fermat $x^{\varphi(N)} \equiv 1 \pmod{N} \Rightarrow x^m \equiv 1 \pmod{N}$. But $m = 2^a b$ so putting $y \equiv x^b \pmod{N}$ we get $y^{2^a} \equiv 1 \pmod{N}$, so $o_p(y)$ and $o_q(y)$ are powers of 2. Since $x \in X$, $o_p(y) \neq o_q(y)$ and wlog $o_p(y) < o_q(y)$. Say $o_p(y) = 2^t$, so $0 \leq t < a$. Then $y^{2^t} \equiv 1 \pmod{p}$ but $y^{2^t} \not\equiv 1 \pmod{q}$. Hence $(y^{2^t} - 1, N) = p$ as required. \square

The RSA cryptosystem

Named for Rivest, Shamir, Adleman (1977) (also Clifford Cocks at GCHQ in 1973).

- Choose large distinct primes p, q at random;
- p, q kept secret. Set $N = pq$;
- Choose e randomly such that $(e, \varphi(N)) = 1$: e is the *encrypting exponent*;
- Euclid says there exist d, k with $de - k\varphi(N) = 1$. d is the *decrypting exponent*.

Public key: (N, e) (ordered pair not gcd), encrypt $m \in \mathcal{M}$ as $c \equiv m^e \pmod{N}$.

Private key: (N, d) (again an ordered pair), decrypt $c \in \mathcal{C}$ as $x \equiv c^d \pmod{N}$.

By Euler-Fermat, $x \equiv (m^e)^d \equiv m^{de} \equiv m^{k\varphi(N)+1} \equiv m \pmod{N}$ (noting the probability $(m, N) \neq 1$ is small enough to ignore). So the decrypting function is inverse to the encrypting function.

If E intercepts the ciphertext she needs to find the plaintext knowing only the public key. This is tantamount to finding factors p, q of N so that $\varphi(N)$ can be computed. Formally

Corollary 19.4. *Finding the RSA private key (N, d) from the public key (N, e) is essentially as difficult as factoring N .*

Proof. We've already seen that factoring N allows us to find d . Conversely, assume we have an algorithm that gives d in terms of N, e . Then $de \equiv 1 \pmod{\varphi(N)}$ and taking $m = de - 1$ in the proof of theorem 19.3 (i) we can factor N . Choose $x \in (\mathbb{Z}/N\mathbb{Z})^\times$, then $\mathbb{P}(x \in X) \geq 1/2$. When $x \in X$ we can factor N . If $x \notin X$ can choose another random x . After r random choices we've found a factor of N with probability at least $1 - 2^{-r}$. \square

We now prove (ii) of

Theorem 19.5.

(i) If $x \in X$ then there exists t with $0 \leq t < a$ such that $(x^{2^t b} - 1, N)$ is a non-trivial factor of N (i.e. p or q).

(ii) $|X| \geq \frac{1}{2} |(\mathbb{Z}/N\mathbb{Z})^\times| = \frac{1}{2} \varphi(N)$.

Proof.

(ii) The Chinese Remainder Theorem provides a multiplicative group isomorphism $(\mathbb{Z}/N\mathbb{Z})^\times \rightarrow (\mathbb{Z}/p\mathbb{Z})^\times \times (\mathbb{Z}/q\mathbb{Z})^\times$, $x \mapsto (x \pmod{p}, x \pmod{q})$. We claim that if we partition $(\mathbb{Z}/p\mathbb{Z})^\times$ according to the value of $o_p(x^b)$ then each subset has size at most $\frac{1}{2} |(\mathbb{Z}/p\mathbb{Z})^\times| = \frac{1}{2} (p-1)$.

We show that some subset has size exactly $\frac{1}{2} |(\mathbb{Z}/p\mathbb{Z})^\times|$. Let g be a primitive root modulo p , i.e. $(\mathbb{Z}/p\mathbb{Z})^\times = \langle g \rangle$. By Fermat's Little Theorem $g^{p-1} \equiv 1 \pmod{p}$ so $g^{2^t b} \equiv 1 \pmod{p}$ and so $o_p(g^b)$ is a power of two. Suppose $o_p(g^b) = 2^t$ (some $0 \leq t \leq a$). Let $x = g^k$ for some $0 \leq k \leq p-2$, then $x^b = (g^b)^k$ so $o_p(x) = \frac{2^t}{\gcd(2^t, k)}$, so $o_p(x^t) = 2^t$ if and only if k is odd, i.e.

$$o_p(x^b) = o_p(g^{bk}) \begin{cases} = o_p(g^b) = 2^t & k \text{ odd} \\ < 2^t & k \text{ even} \end{cases}$$

Thus $\{g^k \pmod{p} : k \text{ odd}\}$ is the required set. This proves the claim. To finish, for each $y \in (\mathbb{Z}/q\mathbb{Z})^\times$, $\{x \in (\mathbb{Z}/p\mathbb{Z})^\times : o_p(x^b) \neq o_q(y^b)\}$ has at least $\frac{p-1}{2}$ elements. Therefore

$$\begin{aligned} |X| &= |\{(x, y) \in (\mathbb{Z}/p\mathbb{Z})^\times \times (\mathbb{Z}/q\mathbb{Z})^\times : o_p(x^b) \neq o_q(y^b)\}| \\ &\geq \frac{1}{2} (p-1)(q-1) = \frac{1}{2} \varphi(N). \end{aligned}$$

□

Remark. We've shown that finding (N, d) from (N, e) is as hard as factoring N . It is unknown whether decrypting messages sent via RSA is as hard as factoring.

RSA avoids the issue of key sharing, but it is slow. Symmetric ciphers (like the one-time pad) are often faster.

Example 19.1 (Sharmir's padlock). $\mathcal{A} = \mathbb{Z}_p$.

- A chooses $a \in \mathbb{Z}_{p-1}$, computes g^a . She finds a' such that $aa' \equiv 1 \pmod{p-1}$;
- B chooses $b \in \mathbb{Z}_{p-1}$, computes g^b . He finds b' such that $bb' \equiv 1 \pmod{p-1}$.

Message $m \in \mathbb{Z}_p$

$$m \xrightarrow{A} \underbrace{m^a}_{=c} \pmod{p} \xrightarrow{B} \underbrace{c^b}_{=d} \pmod{p} \xrightarrow{A} \underbrace{d^{a'}}_{=e} \pmod{p} \xrightarrow{B} e^{b'} \pmod{p}.$$

Then by FLT: $e^{b'} \equiv d^{a'b'} \equiv c^{ba'b'} \equiv c^{a'} \equiv m^{aa'} \equiv m \pmod{p}$.

This suggests:

Diffie-Hellman key exchange

A,B wish to agree on a secret key k . Take p a large prime, g a primitive root modulo p .

- A chooses α and sends $g^\alpha \pmod{p}$ to B;
- B chooses β and sends $g^\beta \pmod{p}$ to A;
- Both parties compute $k = g^{\alpha\beta}$ and use this as the secret key;
- E must find $g^{\alpha\beta}$ knowing g, g^α, g^β .

D-H conjectured that this was as difficult as solving the discrete logarithm problem.

21 Signatures and authentication

Consider a message m sent by A to B. Possible aims include:

- *Secrecy*: A,B can be sure that no third party can read the message;
- *Integrity*: A,B can be sure that no third party can alter the message;
- *Authenticity*: B can be sure that A sent m ;
- Non-repudiation: B can prove to a third party that A sent the message.

Example 21.1 (Authenticity using RSA). Alice uses private key (N, d) to encrypt m . Anyone can decrypt m using the public key (N, e) (as $(m^d)^e \equiv (m^e)^d \equiv m \pmod{p}$), but they cannot forge messages sent by Alice [B picks random message m , sends to A. If B then receives a message back from A, which after decryption ends in m , then he can be sure it comes from A].

Signature schemes preserve integrity and non-repudiation. They also prevent tampering:

Example 21.2 (Homomorphism attack). Bank sends messages of form (M_1, M_2) where M_1 is the client's name, M_2 is the amount to be transferred to M_1 's account. Messages are encoded using RSA as

$$(Z_1, Z_2) = (M_1^e \pmod{N}, M_2^e \pmod{N}).$$

Say you transfer £100 to your account, observe the encrypted message (Z_1, Z_2) , then send (Z_1, Z_2^3) . You then become a millionaire without breaking RSA.

Example 21.3 (Copying). Keep sending (Z_1, Z_2) as above. This attack is defeated by time-stamping.

Definition. A message m is signed as (m, s) where s is a function of m and the private key k , i.e $s = s(m, k)$. B then checks signature using A's public key. The *signature* (or *trapdoor*) function $s : \mathcal{M} \times \mathcal{K} \rightarrow \mathcal{S}$ (where \mathcal{S} is the set of all possible signatures) is designed so that no one without knowledge of the private key can sign messages, yet anyone can check the signature is valid.

Note. We are concerned about the 'signature of the message' not the 'signature of the sender'.

The signature is designed so no one can sign messages without knowledge of the private key, yet anyone can check the signature is valid.

Signatures using RSA

Private key (N, d) , public key (N, e) . A signs m as $(m, m^d \pmod{N})$. Signature s is verified by checking $s^e \equiv m \pmod{N}$.

Problems:

- (i) Homomorphism attack still works;
- (ii) Existential forgery: anyone can produce valid signed messages of form $(s^e \pmod{N}, s)$ after choosing s first. Hopefully these messages are not meaningful.

Solutions:

- (i) Use a better signature scheme;
- (ii) Rather than signing message m , we sign $h(m)$, where h is a *hash function* $h : \mathcal{M} \rightarrow \{1, \dots, N-1\}$, a publicly known function for which it's very difficult to find pairs $x_1 \neq x_2$ for which $h(x_1) = h(x_2)$ (called a *collision*);
- (iii) Example of md5sum files in Carne.

Elgamal signature scheme

A chooses a large prime p , some random integer u ($1 < u < p$), g a primitive root modulo p .

Public key: $p, g, y := g^u \pmod{p}$. Private key: u . Let $h : \mathcal{M} \rightarrow \{1, \dots, p-1\}$ be a hash function. To send $m \in \{0, \dots, p-1\}$, A randomly chooses $k \in \{1, \dots, p-2\}$ with $(k, p-1) = 1$. A then computes r, s as $1 \leq r \leq p-1$, $1 \leq s \leq p-2$ satisfying

$$r \equiv g^k \pmod{p} \tag{1}$$

$$h(m) \equiv ur + ks \pmod{p-1} \tag{2}$$

[since $(k, p-1) = 1$, we can solve for s]. Then A signs m with signature (r, s) . Now

$$g^{h(m)} \equiv g^{ur+ks} \equiv (g^u)^r (g^k)^s \equiv y^r r^s \pmod{p}.$$

Then B accepts signature if $g^{h(m)} \equiv y^r r^s \pmod{p}$.

How to forge a signature? Obvious attacks involve the Discrete Logarithm Problem (find u from $y \equiv g^u \pmod{p}$).

Choice of k

Lemma 21.1. *Given a, b, m , the congruence*

$$ax \equiv b \pmod{m} \quad (3)$$

has either 0 or $\gcd(a, m)$ solutions for x (modulo m).

Proof. Let $d = \gcd(a, m)$. If $d \nmid b$ then there are no solutions. Otherwise we re-write (3) as

$$\frac{a}{d}x \equiv \frac{b}{d} \pmod{\frac{m}{d}}. \quad (4)$$

Now $\gcd(\frac{a}{d}, \frac{m}{d}) = 1$, so (4) has a unique solution $\pmod{\frac{m}{d}}$, hence (3) has d solutions modulo m . \square

It is vital that A chooses a new k to sign each message. Otherwise suppose m_1, m_2 have signatures $(r, s_1), (r, s_2)$:

$$h(m_1) \equiv ur + ks_1 \pmod{p-1}, \quad h(m_2) \equiv ur + ks_2 \pmod{p-1}. \quad (5)$$

$$\implies h(m_1) - h(m_2) \equiv k(s_1 - s_2) \pmod{p-1}.$$

So let $d = (p-1, s_1 - s_2)$, then by the previous lemma there are d solutions for k modulo $p-1$. Choose the one that gives the correct value for $r \equiv g^k \pmod{p}$. Then

$$ur \equiv \frac{h(m_1) - s_1}{k} \pmod{p-1}.$$

This gives $(p-1, r)$ solutions for u (again by the previous lemma). So choose the one that gives $y \equiv g^u \pmod{p}$. In this way you've found Alice's private key as well as the exponent k that she uses for signatures.

21.1 The Digital Signatures Algorithm

The following is a variant of Elgamal developed by the NSA.

Setup: public key (p, q, g) constructed as follows:

- (a) p prime of exactly N bits, where N is a multiple of 64, such that $512 \leq N \leq 1024$ (so $2^{N-1} < p < 2^N$);
- (b) q prime of 160 bits which divides $p-1$;
- (c) $g \equiv h^{(p-1)/q} \pmod{p}$ where h is a primitive root modulo p . Thus g is an element of order q in \mathbb{Z}_p^\times ;
- (d) A chooses private key x , $1 < x < q$ and publishes her public key $y \equiv g^x$.

Signing: for A to sign message m , $0 \leq m < q$, she chooses random k , $1 < k < q$ and computes

$$s_1 \equiv (g^k \pmod{p}) \pmod{q} \text{ and } s_2 \equiv k^{-1}(m + xs_1) \pmod{q}.$$

Her signature for a message m is then the pair (s_1, s_2) , which she sends to B together with m .

Verification: B verifies A's signed message as follows:

- (a) He computes $w = s_2^{-1} \pmod{q}$, $u_1 \equiv mw \pmod{q}$, $u_2 \equiv s_1w \pmod{q}$ and $v = (g^{u_1}y^{u_2} \pmod{p}) \pmod{q}$;
- (b) B accepts the signed message if and only if $v = s_1$.

Proposition 21.2. *The Digital Signatures Algorithm (DSA) works, i.e B accepts correctly signed messages.*

Proof. First note

$$(m + xs_1)w \equiv ks_2s_2^{-1} \pmod{q}. \quad (6)$$

Now

$$v = (g^{u_1}y^{u_2} \pmod{p}) \pmod{q} = (g^{mw}g^{xs_1w} \pmod{p}) \pmod{q}$$

recalling $g^q \equiv 1 \pmod{p}$. Thus by (6)

$$v = (g^{(m+xs_1)w} \pmod{p}) \pmod{q} = (g^k \pmod{p}) \pmod{q} = s_1.$$

So for a correctly signed message, the verification procedure works. \square

Exercise: Suppose A sends messages m_1, m_2 to B,C respectively and provides signatures for each message using DSA. Show that if A is lazy and instead of choosing two different random values of k , she uses the same value for both transmissions, then it's possible for E to recover the private key x from the signed messages.

22 Commitment schemes

Alice wants to send a bit m ($= 0$ or 1) to B, in such a way that

- B cannot determine the value of the bit without A's help;
- A cannot change the bit once it's been sent.

Applications:

- Coin tossing when A and B are in different rooms;
- Selling racing tips (tipster needs to ensure payment before revealing the tip);
- Online polling (results revealed to all parties only after everyone has voted).

We give two examples of this *commit & reveal strategy*, known as bit commitment.

Example 22.1 (Using a public key cryptosystem).

- e_A, d_A encryption/decryption functions for a public key cryptosystem used by A where e_A is published but d_A is kept secret by A;
- A makes her choice $m \in \mathbb{F}_2$ and commits $c = e_A(m)$ to B. With a secure cipher, B cannot decipher this;

- To reveal her choice, A sends to B her private key so he can find d_A . He computes $d_A(c) = d_A(e_A(m)) = m$ to find A's choice. He can also check d_A, e_A are indeed inverse functions, so he can be confident that A has not send the wrong private key (to pretend her choice was different).

Example 22.2 (Using coding theory). A,B have two ways to communicate: via a noisy channel (a BSC(p)), or a clear channel (transmission with no errors). Assume $0 < p < 1/2$ and it corrupts bits independent of any action by A or B, so neither can affect its behaviour.

- B chooses a binary linear code $C \subseteq \mathbb{F}_2^N$ with minimum distance d ;
- A chooses a random, non-trivial linear map $\theta : C \rightarrow \mathbb{F}_2$;
- Both publish their choices;
- To send a bit $m \in \mathbb{F}_2$, A chooses a random codeword $c \in C$ such that $\theta(c) = m$ and sends c to B via the noisy channel. B receives $r = c + e \in \mathbb{F}_2^N$ (e the error);
- The expected value for $d(r, c) = d(e, 0)$ is Np , with N chosen so that $Np \gg d$ so B can't tell what the original c was, hence can't find $\theta(c) = m$ and determine A's choice;
- To reveal: A sends c to B via the clear channel. B can check $d(c, r)$ is close to the expected value Np . If it is, he accepts that $\theta(c)$ was A's choice. If not, he rejects it;
- It's possible that many more or fewer bits of c were corrupted than expected, in which case he rejects even though she has not cheated. Choose N, d such that the probability of this occurring is very small. If it does occur, A,B repeat the process.

We've seen B can't read A's guess until she has revealed it. Why is it that A can't cheat by changing her guess and sending a different codeword $c' \neq c$ though the clear channel? Alice knows the codeword c that she originally chose, but she doesn't know how it will be corrupted via the noisy channel. All she knows is that the received $r \in \mathbb{F}_2^N$ is at most distance $\approx Np$ from c . If she sends c' then she must ensure that $d(r, c') \approx Np$. The probability that this happens is small unless she chooses c' very close to c , which can't happen as the code has minimum distance d .

20 Secret sharing schemes

If CMS is attacked by MIO, the Faculty will retreat to a bunker, known as MR2. Entry to MR2 involves tapping out a positive integer S (the 'secret') known only to the Leader, L. Each of the n members of the Faculty knows a certain pair of numbers (their *shadow/share*) and it is required that, in the absence of L, any k members of the Faculty can reconstruct S from their shadows, but no $k - 1$ can do so. How can this be done?