Introduction

We model communication:

$$\underbrace{\text{SOURCE}}_{\text{message}} \to \underbrace{\text{ENCODER}}_{\text{codewords}} \xrightarrow[\text{errors,noise}]{\text{CHANNEL}} \underbrace{\text{DECODER}}_{\text{received word error correction}} \to \underbrace{\text{RECIEVER}}_{\text{message}}.$$

Examples: optical signals, electrical telegraph, SMS (compression), postcodes, CDs (error correction), zip/gz files (compression).
Given a source and a channel, modelled probabilistically, the basic problem is to design an encoder and decoder to transmit messages economically (noiseless coding; compression) and reliably (noisy coding).

Examples:

- Noiseless coding: Morse code: common letters are assigned shorter codewords, e.g $A \mapsto \bullet-$, $E \mapsto \bullet$, $Q \mapsto --\bullet-$, $S \mapsto \bullet\bullet\bullet$, $O \mapsto ---$, $Z \mapsto --\bullet\bullet$. Noiseless coding is adapted to source.
- Noisy coding: Every book has an ISBN $a_1, a_2, \ldots, a_9, a_{10}$, $a_i \in \{0, 1, \ldots, 9\}$ for $1 \leq i \leq 9$ and $a_{10} \in \{0, 1, \ldots 9, X\}$ with $\sum_{j=1}^{10} j a_j \equiv 0 \pmod{11}$. This detects common errors - e.g one incorrect digit, transposition of two digits. Noisy coding is adapted to the channel.

Plan:

(I) Noiseless coding - entropy

(II) Error correcting codes - noisy channels

(III) Information theory - Shannon's theorems

(IV) Examples of codes

(V) Cryptography

Books: [GP], [W], [CT], [TW], Buchmann, Körner. Online notes: Carne, Körner.

Basic Definitions

Definition (Communication channel)**.** A *communication channel* accepts symbols from a alphabet $\mathcal{A} = \{a_1, \ldots, a_r\}$ and it outputs symbols from alphabet $\mathcal{B} = \{b_1, \ldots, b_s\}$. Channel modelled by the probabilities $\mathbb{P}(y_1 \ldots y_n \text{ received} | x_1 \ldots x_n \text{sent})$.
A *discrete memoryless channel* (DMC) is a channel with

$$p_{ij} = \mathbb{P}(b_j \text{ received} | a_i \text{ sent})$$

the same for each channel use and independent of all past and future uses. The channel matrix is $P = (b_{ij})$, a $r \times s$ stochastic matrix.

Definition (Binary symmetric channel)**.** The *binary symmetric channel* (BSC) with error probability $p \in [0, 1)$ from $\mathcal{A} = \mathcal{B} = \{0, 1\}$. The channel matrix is

$$\begin{pmatrix} 1-p & p \\ p & 1-p \end{pmatrix}.$$

A symbol is transmitted correctly with probability $1 - p$. Usually assume $p < 1/2$.

The *binary erasure channel* (BEC) has $\mathcal{A} = \{0, 1\}$, $\mathcal{B} = \{0, 1, *\}$. The channel matrix is

$$\begin{pmatrix} 1-p & 0 & p \\ 0 & 1-p & p \end{pmatrix}.$$

So $p = \mathbb{P}$(symbol can't be read).

Definition. We model n uses of a channel by the nth extension, with input alphabet $\mathcal{A}^n$ and output alphabet $\mathcal{B}^n$. A *code C of length n* is a function $\mathcal{M} \to \mathcal{A}^n$ where $\mathcal{M}$ is the set of possible messages. Implicitly we also have a decoding rule $\mathcal{B}^n \to \mathcal{M}$. The *size* of C is $m = |\mathcal{M}|$. The *information rate* is $\rho(C) = \frac{1}{n} \log_2 m$. The *error rate* is $\hat{e}(C) = \max_{x \in \mathcal{M}} \mathbb{P}(\text{error}|x \text{ sent})$.

Remark. For the remainder of the course we write log instead of $\log_2$.

Definition. A channel can *transmit reliably at rate R* if there exists $(C_n)_{n=1}^{\infty}$ with each C_n a code of length n such that

$$\lim_{n\to\infty} \rho(C_n) = R \;\&\; \lim_{n\to\infty} \hat{e}(C_n) = 0.$$

The *capacity* is the supremum of all reliable transmission rates. We'll see in Chapter 9 that a BSC with error probability $p < 1/2$ has non-zero capacity.

1 Noiseless coding

1.1 Prefix-free codes

For an alphabet $\mathcal{A}$, $|\mathcal{A}| < \infty$, let $\mathcal{A}^* = \bigcup_{n \geq 0} \mathcal{A}^n$, the set of all finite strings from $\mathcal{A}$. The *concatenation* of strings $x = x_1 \dots x_r$ and $y = y_1 \dots y_s$ is $xy = x_1 \dots x_r y_1 \dots y_s$.

Definition. Let $\mathcal{A}, \mathcal{B}$ be alphabets. A code is a function $c : \mathcal{A} \to \mathcal{B}^*$. The strings $c(a)$ for $a \in \mathcal{A}$ are called *codewords* or *words* (CWS).

Example 1.1 (Greek fire code)**.** $\mathcal{A} = \{\alpha, \beta, \ldots, \omega\}$ (greek alphabet), $\mathcal{B} = \{1, 2, 3, 4, 5\}$., $c : \alpha \mapsto 11, \beta \mapsto 12, \ldots, \psi \mapsto 53$, $\omega \mapsto 54$. xy means hold up x torches and another y torches nearby.

Example 1.2. $\mathcal{A} =$ words in a dictionary, $\mathcal{B} = \{A, B, \ldots, Z, \omega\}$. $c : \mathcal{A} \to \mathcal{B}$ splits the word and follows with a space. Send message $x_1 \ldots x_n \in \mathcal{A}^*$ as $c(x_1) \ldots c(x_n) \in \mathcal{B}^*$. So c extends to a function $c^* : \mathcal{A}^* \to \mathcal{B}^*$.

Definition. c is said to be *decipherable* if the induced map c^* (as in the previous example) is injective. In other words, each string from $\mathcal{B}$ corresponds to at most one message.

Clearly if c is decipherable, it is necessary for c to be injective. However it is not sufficient:

Example 1.3. $\mathcal{A} = \{1, 2, 3, 4\}$, $\mathcal{B} = \{0, 1\}$. Define $c : 1 \mapsto 0$, $2 \mapsto 1$, $3 \mapsto 00$, $4 \mapsto 01$. Then $c^*(114) = 0001 = c^*(312) = c^*(144)$ yet c is injective.

Notation: $|\mathcal{A}| = m$, $|\mathcal{B}| = a$, call c am a-ary code of size m. For example a 2-ary code is a binary one, and a 3-ary code is a ternary code.

Our aim is to construct decipherable codes with short word lengths. Assuming c is injective, the following codes are always decipherable:

(i) A block code has all codewords of the same length (e.g Greek fire code);

(ii) A comma code reserves a letter from $\mathcal{B}$ to signal the end of a word (e.g Example 1.2);

(iii) A prefix-free code is a code where no codeword is a prefix of any other distinct word (if $x, y \in \mathcal{B}^*$ then x is a prefix of y if $y = xz$ for some string $z \in \mathcal{B}^*$).

(i) and (ii) are special cases of (iii). As we can decode the message as it is received, prefix-free codes are sometimes called *instantaneous*.

Exercise: find a decipherable code which is not prefix-free.

Definition (Kraft's inequality)**.** $|\mathcal{A}| = m$, $|\mathcal{B}| = a$, $c : \mathcal{A} \to \mathcal{B}^*$ has word lengths $l_1, \ldots, l_m$. Then Kraft's inequality is

$$\sum_{i=1}^{m} a^{-l_i} \leq 1. \qquad (*)$$

Theorem 1.1. *A prefix-free code exists if and only if Kraft's inequality* $(*)$ *holds.*

Proof. Rewrite $(*)$ as

$$\sum_{l=1}^{s} n_l a^{-l} \leq 1, \qquad (**)$$

where n_l is the number of codewords with length l, and $s = \max_{1 \leq i \leq m} l_i$.

Now if $c : \mathcal{A} \to \mathcal{B}^*$ is prefix-free,

$$n_1 a^{s-1} + n_2 a^{s-2} + \ldots + n_{s-1} a + n_a \leq a^s.$$

Indeed the LHS is the number of strings of length s in B with some codeword of c as a prefix, and the RHS is the total number of strings of length S. Dividing through by a^s we get $(**)$.

Now given $n_1, \ldots, n_s$ satisfying $(**)$, we try to construct a prefix-free code c with n_l codewords of length l, $\forall l \leq s$. Proceed by induction on s, $s = 1$ is clear (since $(**)$ gives $n_1 \leq a$ so can construct code).

By the induction hypothesis, there exists a prefix-code $\hat{c}$ with n_l codewords of length l for all $l \leq s - 1$. Then $(**)$ implies

$$n_1 a^{s-1} + n_2 a^{s-2} + \ldots n_{s-1} a + n_s \leq a^s.$$

The first $s - 1$ terms on the LHS sum to the number of strings of length s with a codeword of $\hat{c}$ as a prefix and the RHS is the number of strings of length s. Hence we can add at least n_s new codewords of length s to $\hat{c}$ and maintain the prefix-free property.

□

Remark. This proof is constructive: just choose codewords in order of increasing length, ensuring that no previous codeword is a prefix.

Theorem 1.2 (McMillan)**.** *Any decipherable code satisfies Kraft's inequality.*

Proof (Karush, 1961). Let $c : \mathcal{A} \to \mathcal{B}^*$ be a decipherable code with word lengths $l_1, \ldots, l_m$. Set $s = \max_{1 \leq i \leq m} l_i$. For $R \in \mathbb{N}$

$$\left(\sum_{i=1}^{m} a^{-l_i} \right)^R = \sum_{l=1}^{Rs} b_l a^{-l}, \tag{$\dagger$}$$

where b_l is the number of ways of choosing R codewords of total length l. Since c is decipherable, any string of length l formed from codewords must correspond to at most one sequence of codewords, i.e $b_l \leq |\mathcal{B}^l| = a^l$. Subbing this into $(\dagger)$

$$\left(\sum_{i=1}^{m} a^{-l_i} \right)^R \leq \sum_{i=1}^{Rs} a^l a^{-l} = Rs,$$

so

$$\sum_{i=1}^{m} a^{-l_i} \leq (Rs)^{1/R} \to 1 \text{ as } R \to \infty.$$

Hence $\sum_{i=1}^{m} a^{-l_i} \leq 1$. □

Corollary 1.3. *A decipherable code with prescribed word lengths exists if and only if a prefix-free code with the same word lengths exists.*

Proof. Combine previous two theorems. □

Therefore we can restrict our attention to prefix-free codes.

2 Shannon's Noiseless Coding Theorem

Entropy is a measure of 'randomness' or 'uncertainty'. Suppose we have a random variable X taking a finite set of values $x_1, \ldots, x_n$ with probabilities $p_1, \ldots, p_n$ respectively. The *entropy* $H(X)$ of X is the expected number of fair coin tosses needed to simulate X (roughly speaking).

Example 2.1. Suppose $p_1 = p_2 = p_3 = p_4 = 1/4$. Identify (x_1, x_2, x_3, x_4) with (HH, HT, TH, TT). Then the entropy is 2.

Example 2.2. Suppose $(p_1, p_2, p_3, p_4) = (1/2, 1/4, 1/8, 1/8)$. Identify (x_1, x_2, x_3, x_4) with (H, TH, TTH, TTT). Then the entropy is

$$\frac{1}{2} \times 1 + \frac{1}{4} \times 2 + \frac{1}{8} \times 3 + \frac{1}{8} \times 3 = \frac{7}{4}.$$

In a sense, the previous example (2.1) was 'more random' than this.

Definition (Entropy)**.** The *entropy* of X is

$$H(X) = -\sum_{i=1}^{b} p_i \log p_i.$$

(Recall that log =: $\log_2$ here.) Note $H(X) \geq 0$. It is measured in *bits* (binary digits). Conventionally, we take $0 \log 0 = 0$.

Example 2.3. Take a biased coin $\mathbb{P}(H) = p$, $\mathbb{P}(T) = 1 - p$. Write $H(p, 1-p) := H(p)$. Then

$$H(p) = -p \log p - (1-p) \log(1-p).$$

Note that $H'(p) = \log \frac{1-p}{p}$. Hence the entropy is maximised for $p = 1/2$ (giving entropy 1).

Proposition 2.1 (Gibbs' inequality)**.** *Let $(p_1, \ldots, p_n), (q_1, \ldots, q_n)$ be probability distributions. Then*

$$-\sum_{i=1}^{n} p_i \log p_i \leq -\sum_{i=1}^{n} p_i \log q_i.$$

(The RHS is sometimes called the cross entropy or mixed entropy) Furthermore we have equality iff $p_i = q_i$ for all i.

Proof. Since $\log x = \frac{\ln x}{\ln 2}$, we may replace log with ln. Put $I = \{1 \leq i \leq n : p_i \neq 0\}$. Now $\ln x = x - 1$ for all $x > 0$ with equality iff $x = 1$. Hence $\ln \frac{q_i}{p_i} \leq \frac{q_i}{p_i} - 1$ for all $i \in I$. So

$$\sum_{i \in I} p_i \ln \frac{q_i}{p_i} \leq \underbrace{\sum_{i \in I} q_i}_{\leq 1} - \underbrace{\sum_{i \in I} p_i}_{=1} \leq 0$$

$$\implies -\sum_{i \in I} p_i \ln p_i \leq -\sum_{i \in I} p_i \ln q_i$$

$$\implies -\sum_{i=1}^{n} p_i \ln p_i \leq -\sum_{i=1}^{n} p_i \ln q_i.$$

If equality holds, then $\sum_{i\in I} q_i = 1$ and $\frac{p_i}{q_i} = 1$ for all $i \in I$. So $q_i = p_i$ for all $1 \leq i \leq n$. □

Corollary 2.2. *$H(p_1, p_2, \ldots, p_n) \leq \log n$ with equality iff $p_1 = p_2 = \ldots = p_n = 1/n$.*

Proof. Take $q_1 = q_2 = \ldots = q_n = 1/n$ in Gibbs' inequality. □

Let $\mathcal{A} = \{\mu_1, \ldots, \mu_m\}$, $|\mathcal{B}| = a$ $(m, n \geq 2)$. The random variable X takes values $\mu_1, \ldots, \mu_m$ with probabilities $p_1, \ldots, p_m$.

Definition. If $c : \mathcal{A} \to \mathcal{B}^*$ is a code, we say it is *optimal* if has the smallest possible expected word length. i.e $\mathbb{E}S := \sum_{i=1}^{n} p_i l_i$ is minimal amongst all decipherable codes.

Theorem 2.3 (Shannon's Noiseless Coding Theorem)**.** *The expected word length $\mathbb{E}S$ of an optimal code satisfies*

$$\frac{H(X)}{\log a} \leq \mathbb{E}S < \frac{H(X)}{\log a} + 1.$$

Remark. The lower bound is actually true for any decipherable code.

Proof. We first get the lower bound. Let $c : \mathcal{A} \to \mathcal{B}^*$ be decipherable with word lengths $l_1, \ldots, l_m$. Let $q_i = \frac{a^{-l_i}}{D}$ where $D = \sum_{i=1}^{m} a^{-l_i}$. Note $\sum_{i=1}^{m} q_i = 1$. By Gibbs' inequality

$$\begin{aligned} H(X) &\leq -\sum_{i=1}^{m} p_i \log q_i \\ &= -\sum_{i=1}^{m} p_i(-l_i \log a - \log D) \\ &= \left(\sum_{i=1}^{m} p_i l_i\right) \log a + \log D. \end{aligned}$$

By McMillan, $D \leq 1$ so $\log D \leq 0$. Hence

$$H(X) \leq \left(\sum_{i=1}^{m} p_i l_i\right) \log a \implies \frac{H(X)}{\log a} \leq \mathbb{E}S.$$

And we have equality iff $p_i = a^{-l_i}$ for some integers $l_1, \ldots, l_m$. Note we have only used decipherability so far.

Now we get the upper bound. Take $l_i = \lceil -\log_a p_i \rceil$. Then

$$-\log_a p_i \leq l_i < -\log_a p_i + 1.$$

Hence $\log_a p_i \geq -l_i$, so $p_i \geq a^{-l_i}$. Therefore $\sum_{i=1}^m a^{-l_i} \leq \sum_{i=1}^m p_i = 1$. By Kraft's inequality, there exists a prefix-free code c with word lengths $l_1, \ldots, l_m$. c has expected word length

$$\mathbb{E}S = \sum_{i=1}^m p_i l_i < \sum_{i=1}^m p_i(-\log_a p_i + 1) = \frac{H(X)}{\log a} + 1.$$

□

Example 2.4 (Shannon-Fano Coding)**.** We mimic the above proof: given $p_1, \ldots, p_m$, set $l_i = \lceil -\log_a p_i \rceil$. Construct a prefix-free code with word lengths l_i by choosing codewords in order of increasing length, ensuring any new codeword has no previous codeword as a prefix (Kraft's inequality ensures we can do this).

Example 2.5. Take $a = 2, m = 5$.

i	p_i	$\lceil -\log_2 p_i \rceil$	code
1	0.4	2	00
2	0.2	3	010
3	0.2	3	011
4	0.1	4	1000
5	0.1	4	1001

Then $\mathbb{E}S = \sum_{i=1}^m p_i l_i = 2.8$, $H = H/\log a = 2.12$. [See also Carne p13.]

3 Huffman Coding

How to construct an optimal code? Take $\mathcal{A} = \{\mu_1, \ldots, \mu_m\}$, $p_i = \mathbb{P}(X = \mu_i)$. For simplicitly take $|\mathcal{B}| = a = 2$. Without loss of generality $p_1 \geq p_2 \geq \ldots \geq p_m$. Huffman gave an inductive definition of codes that we can prove are optimal. If $m = 2$, we take codewords $0, 1$. If $m > 2$, first take the Huffman code for messages $\mu_1, \ldots, \mu_{m-2}, \nu$ with probabilities $p_1, \ldots, p_{m-2}, p_{m-1} + p_m$. Then append 0 (respectively 1) to the codeword for ν to give a codeword for μ_{m-1} (respectively μ_m).

Notes.

- Huffman codes are prefix-free;
- Huffman codes are not unique: choice is needed if some of the p_i are equal.

Example 3.1. Revisit Example 2.5. We have

i	p_i	$c^{(1)}$	$p_i^{(2)}$	$c^{(2)}$	$p_i^{(3)}$	$c^{(3)}$	$p_i^{(4)}$	$c^{(4)}$
1	0.4	1	0.4	1	0.4	1	0.6	0
2	0.2	01	0.2	01	0.4	00	0.4	1
3	0.2	000	0.2	000	0.2	01		
4	0.1	0010	0.2	001				
5	0.1	0011						

Theorem 3.1. *Huffman codes are optimal (Huffman, 1952).*

Proof. We show by induction on m that Huffman codes of size $m = |\mathcal{A}|$ are optimal.

$\underline{m = 2}$: codewords are $0, 1$ - clearly optimal.

$\underline{m > 2}$: let c_m be a Huffman code for X_m, which takes values $\mu_1, \ldots, \mu_m$ with probabilities $p_1 \geq p_2 \geq \ldots \geq p_m$; each c_m is constructed from Huffman code c_{m-1} for X_{m-1} which takes values $\mu_1, \ldots, \mu_{m-2}, \nu$ with probabilities $p_1, \ldots, p_{m-2}, p_{m-1} + p_m$. Then the expected word length is

$$\mathbb{E}S_m = \mathbb{E}S_{m-1} + p_{m-1} + p_m. \qquad (*)$$

Let c'_m be an optimal code for X_m. Wlog c'_m is still prefix-free. Wlog the last two codewords of c'_m have maximal length and differ only in the final position (see next lemma). Say

$$c'_m(\mu_{m-1}) = y0, \ c'_m(\mu_m) = y1 \text{ for some } y \in \{0,1\}^*.$$

Let c'_{m-1} be some prefix-free code for X_{m-1}, given by

$$c'_{m-1}(\mu_i) = \begin{cases} c'_m(\mu_i) & 1 \leq i \leq m-2 \\ c'_{m-1}(\nu) = y \end{cases}.$$

Then the expected word length satisfies

$$\mathbb{E}S'_m = \mathbb{E}S'_{m-1} + p_{m-1} + p_m. \qquad (**)$$

By the inductive hypothesis, c_{m-1} is optimal, so $\mathbb{E}S_{m-1} \leq \mathbb{E}S'_{m-1}$. By $(*)$ and $(**)$ this implies $\mathbb{E}S_m \leq \mathbb{E}S'_m$.

□

Lemma 3.2. *Suppose letters $\mu_1, \ldots, \mu_m$ in $\mathcal{A}$ are sent with probabilities $p_1, p_2, \ldots, p_m$. Let c be an optimal (prefix-free) code with word lengths $l_1, \ldots, l_m$. Then*

(i) If $p_i > p + j$, then $l_i \leq l_j$;

(ii) Amongst all codewords of maximal length there exist two that differ only in the final digit.

Proof. (i) is obvious. For (ii), could otherwise just delete the final digit of the codeword of maximal length (since prefix-free). □

Remark. Note not all optimal codes are Huffman (look at the case $m = 4$).

Our main result says that if we have a prefix-free optimal code with word lengths $l_1, \ldots, l_m$ and associated probabilities $p_1, \ldots, p_m$, then there is a Huffman code with these word lengths.

4 Joint Entropy

If X, Y are random variables with values in $\mathcal{A}$ and $\mathcal{B}$ respectively, then (X, Y) is a random variable with values in $\mathcal{A} \times \mathcal{B}$, and the *entropy* $H(X, Y)$ is called the joint entropy, given by

$$H(X, Y) = -\sum_{x \in \mathcal{A}} \sum_{y \in \mathcal{B}} \mathbb{P}(X = x, Y = y) \log \mathbb{P}(X = x, Y = y).$$

This generalises to any finite number of random variables.

Lemma 4.1. *Let X, Y be random variables taking values in $\mathcal{A}$ and $\mathcal{B}$ respectively. Then*

$$H(X, Y) \leq H(X) + H(Y),$$

with equality if and only if X and Y are independent.

Proof. Write $\mathcal{A} = \{x_1, \ldots, x_m\}$, $\mathcal{B} = \{y_1, \ldots, y_n\}$. Let

$$p_{ij} = \mathbb{P}(X = x_i, Y = Y_j),\ p_i = \mathbb{P}(X = x_i),\ q_j = \mathbb{P}(Y = y_j).$$

Apply Gibbs' inequality to the probability distributions $\{p_{ij}\}$ and $\{p_i q_j\}$ to obtain

$$\begin{aligned}
-\sum_{i,j} p_{ij} \log p_{ij} &\leq -\sum_{i,j} p_{ij} \log(p_i q_j) \\
&= -\sum_i \left(\sum_j p_{ij} \right) \log p_i - \sum_j \left(\sum_i p_{ij} \right) \log q_j \\
&= -\sum_i p_i \log p_i - \sum_j q_j \log q_j \\
&= H(X) + H(Y).
\end{aligned}$$

With equality if and only if $p_{ij} = p_i q_j$ for all i, j. □

Error-correcting codes

5 Noisy channels and Hamming's code

Definition. A *binary* $[m, n]$*-code* is a subset C of $\{0,1\}^n$ of size $m = |C|$. n is the length of the code and the elements of C are called codewords.

We use an $[n, m]$-code to send one of m messages through a BSC (binary symmetric channel) making n uses of the channel. Clearly $1 \leq m \leq 2^n$, so $0 \leq \frac{1}{n} \log m \leq 1$.

Definition. For any $x, y \in \{0,1\}^n$ the *Hamming distance* is

$$d(x, y) = |\{i : 1 \leq i \leq n, x_i \neq y_i\}|.$$

Definition.

(i) The *ideal observer* decoding rule decodes $x \in \{0,1\}^n$ as $c \in C$ maximising $\mathbb{P}(c \text{ sent}|x \text{ received})$.

(ii) The *maximum likelihood* decoding rule decodes $x \in \{0,1\}^n$ as $c \in C$ maximising $\mathbb{P}(x \text{ received}|c \text{ sent})$

(iii) The *minimum distance* decoding rule decodes $x \in \{0,1\}$ as $c \in C$ minimising $d(x, C)$.

Lemma 5.1.

(a) If all the messages are equally likely, then (i) and (ii) above are equivalent.

(b) If $p < 1/2$ (error probability) then (ii) and (iii) are equivalent.

Remark. If $p = 1/2$ the code is called *useless*. If $p = 0$ the code is called *lossless*.

Proof.

(a) We have

$$\mathbb{P}(c \text{ sent}|x \text{ received}) = \frac{\mathbb{P}(c \text{ sent}, x \text{ received})}{\mathbb{P}(x \text{ received})} = \frac{\mathbb{P}(c \text{ sent})}{\mathbb{P}(x \text{ received})} \mathbb{P}(x \text{received}|c \text{sent}).$$

So by hypothesis, $\mathbb{P}(c \text{ sent})$ is independent of $c \in C$. So for fixed x, maximising $\mathbb{P}(c \text{ sent}|x \text{ received})$ is the same as maximising $\mathbb{P}(x \text{ received}|c \text{ sent})$.

(b) Let $r = d(x, c)$. Then $\mathbb{P}(x \text{ received}|c \text{ sent}) = p^r(1-p)^{n-r} = (1-p)^n \left(\frac{p}{1-p}\right)^r$. Since $p < 1/2$. $\frac{p}{1-p} < 1$. So maximising $\mathbb{P}(x \text{ received}|c \text{ sent})$ is the same as minimising r.

□

We choose to use minimum distance decoding from now on.

Example 5.1. Suppose 000, 111 are sent with probabilities $\alpha = 9/10$, $\beta = 1/10$ respectively through a BSC with error probability $p = 1/4$. Suppose 110 is received. Then

$$\mathbb{P}(000 \text{ sent}|110 \text{ received}) = \frac{\alpha p^2(1-p)}{\alpha p^2(1-p) + (1-\alpha)p(1-p)^2} = \frac{3}{4},$$

$$\text{similarly } \mathbb{P}(111 \text{ sent}|110 \text{ received}) = \frac{1}{4}.$$

So the ideal observer decodes as 000. But the maximum likelihood/minimum distance rules decode as 111.

Remarks.

- Minimum distance decoding may be expensive in terms of time and storage if $|C|$ is large.
- Need to specify a convention in case there is no unique maximiser (e.g make a random choice, or request the message is sent again).

We aim to detect, or even correct errors.

Definition. A code C is

- *d-error* detecting if changing up to d digits in each codeword can never produce another codeword. In other words, each codeword is of Hamming distance greater than d from every other codeword.
- *e-error* correcting if knowing that $x \in \{0,1\}^n$ differs fom a codeword in at most e places we can deduce the codeword.

Examples.

(a) A *repitition code* of length n has codewords $\underbrace{00\ldots0}_{n \text{ times}}, \underbrace{11\ldots1}_{n \text{ times}}$. This is a $[n, 2]$-code. It is $(n-1)$-error detecting and $\lfloor \frac{n-1}{2} \rfloor$-error correcting. But the information rate is only $1/n$.

(b) A *simple parity check code* or *paper tape code*: identify $\{0,1\}$ with $\mathbb{F}_2$ and let $C = \{(x_1, \ldots, x_n) \in \{0,1\}^n : \sum_{i=1}^n x_i = 0\}$. This is a $[n, 2^{n-1}]$-code, 1-error detecting but cannot correct errors. The information rate is $\frac{n-1}{n}$.

(c) Hamming's original code (1950): a 1-error correcting binary $[7, 16]$-code. Take $C \subseteq \mathbb{F}_2^7$ where

$$C = \{c \in \mathbb{F}_2^7 : c_1 + c_3 + c_5 + c_7 = 0, c_2 + c_3 + c_6 + c_7 = 0, c_4 + c_5 + c_6 + c_7 = 0\}.$$

The bits c_3, c_5, c_6, c_7 are arbitrary and c_1, c_2, c_4 are forced (called the check digits) so $|C| = 2^4$. To decode: suppose we recieve $x \in \mathbb{F}_2^7$. We form the *syndrome*: $z = z_x = (z_1, z_2, z_4) \in \mathbb{F}_2^7$ where

$$\begin{aligned} z_1 &= x_1 + x_3 + x_5 + x_7 \\ z_2 &= x_2 + x_3 + x_6 + x_7 \\ z_4 &= x_4 + x_5 + x_6 + x_7. \end{aligned}$$

If $x \in C$, then $z_x = (0, 0, 0)$. If $d(x, c) = 1$ for some $c \in C$, then place where x and c differ is given by $z_1 + 2z_2 + 4z_4$ (not mod 2). Check: if $x = c + e_i$ where e_i has all 0's except a 1 in the ith position, then $z_x = z_{e_i}$, so check for each $1 \leq i \leq 7$.

Lemma 5.2. *The Hamming distance is a metric on $\mathbb{F}_2^n$.*

Proof. Trivial. □

Definition. The *minimum distance of a code* is the minimum of $d(c_1, c_2)$ for all codewords c_1, c_2 with $c_1 \neq c_2$.

Lemma 5.3. *Let C be a code with minimum distance $d > 0$. Then*

(i) C is $(d-1)$-error detecting, but cannot detect all sets of d errors.

(ii) C is $\lfloor \frac{d-1}{2} \rfloor$-error correting, but cannot correct all sets of $\lfloor \frac{d-1}{2} \rfloor + 1$ errors.

Proof.

(i) If $x \in \mathbb{F}_2^n$ and $c \in C$ are such that $0 < d(x, c) \leq d - 1$, then we know that $x \notin C$ so this is $(d-1)$-error detecting. However there must exist $c_1, c_2 \in C$ such that $d(c_1, c_2) = d$, so we cannot say if there's an error if c_1 is 'corrupted' to c_2 in d errors.

(ii) Take $e = \lfloor \frac{d-1}{2} \rfloor$. If $x \in \mathbb{F}_2^n$ and $c_1 \in C$ are such that $d(x, c_1) \leq \lfloor e$ then for any $c_1 \neq c_2 \in C$ we have $d(x, c_2) \geq d(c_1, c_2) - d(c_1, x) \geq d - e > e$. So C is e-error correcting. Now take $c_1, c_2 \in C$ with $d(c_1, c_2) = d$. Then take $x \in \mathbb{F}_2^n$ such that x differs from c_1 is prescisely $e + 1$ places where c_1 and c_2 differ. Then $d(c_1, x) = e + 1$ and $d(x, c_2) = d - (e + 1) \leq e + 1$. So C cannot be $(e+1)$-error correcting.

□

Definition. A $[n, m]$-code with minimum distance is called a $[n, m, d]$-*code*.

Notes.

- $m \leq 2^n$ with equality if and only if $C = \mathbb{F}_2^n$ (trivial code)
- $d \leq n$, with equality in case of the repitition code.

Example 5.2.

(i) Repitition code of length n is a $[n, 2, n]$-code, $(n-1)$-error detecting and $\lfloor \frac{n-1}{2} \rfloor$-error correcting.

(ii) Simple parity check code is a $[n, 2^{n-1}, 2]$-code, 1-error detecting and 0-error correcting.

(iii) Hamming's original code is 1-error correcting, implying $d \geq 3$. Also $0000000, 1110000$ are distance 3 apart, so $d = 3$. So this is a $[7, 16, 3]$-code and is 2-error detecting.

6 Covering estimates

Take $x \in \mathbb{F}_2^n$, $r \geq 0$. Then $\overline{B}(x,r) = \{y \in \mathbb{F}_2^n : d(x,y) \leq r\}$ is the *closed Hamming ball*. Denote $V(n,r) = |\overline{B}(x,r)| = \sum_{i=0}^{r} \binom{n}{i}$, the *volume*.

Lemma 6.1 (Hamming's bound)**.** *An e-error correcting code C of length n has*

$$|C| \leq \frac{2^n}{V(n,e)}.$$

Proof. C is e-error correcting so $\{B(c,e)\}_{c\in C}$ are pairwise disjoint balls, so $\sum_{c\in C} |B(c,e)| = |C|V(n,e) \leq |\mathbb{F}_2^n| = 2^n$. □

Lemma 6.2. *A code C of length n that can correct e errors is perfect if $|C| = \frac{2^n}{V(n,e)}$. Equivalently, for all $x \in \mathbb{F}_2^n$ there exists a unique $c \in C$ such that $d(x,c) \leq e$. In this case, any $e+1$ errors will make you decode incorrectly.*

Example 6.1.

(a) Hamming $[7,16,3]$-code is 1-error correcting and

$$\frac{2^n}{V(n,e)} = \frac{2^7}{V(7,1)} = \frac{2^7}{1+7} = 2^4 = |C|.$$

(b) Binary repitition code of length n (for n odd) is perfect.

Remark. If $\frac{2^n}{V(n,e)} \notin \mathbb{Z}$ then there does not exist a perfect e-error correcting code of length n. Converse is also false ($n = 90, e = 2$ on Example Sheet 2).

Definition. Define $A(n,d) = \max\{m : \exists [n,m,d]\text{-code}\}$.

The $A(n,d)$ are unknown in general. But we have some special cases:

Examples.

- $A(n,1) = 2^n$ (trivial code)
- $A(n,n) = 2$ (repitition code)
- $A(n,2) = 2^{n-1}$ (simple parity check code)

Lemma 6.3. $A(n,d+1) \leq A(n,d)$.

Proof. Let $m = A(n,d+1)$ and take a $[n,m,d+1]$-code C. Let $c_1, c_2 \in C$ have $d(c_1,c_2) = d+1$. Let c_1' differ from c_1 in a single place where c_1 and c_2 differ. Hence $d(c_1',c_2) = d$. If $c \in C \setminus \{c_1\}$, then $d(c,c_1) \leq d(c,c_1') + d(c_1',c_1)$ so $d+1 \leq d(c,c_1')+1$. Hence $d(c,c_1') \geq d$. So replacing c_1 with c_1', we get an $[n,m,d]$-code. □

Corollary 6.4. $A(n,d) = \max\{m : \exists [n,m,d']\text{-code for some } d' \geq d\}$.

Theorem 6.5.

$$\frac{2^n}{V(n,d-1)} \underbrace{\leq}_{GSV\ bound} A(n,d) \underbrace{\leq}_{Hamming\ bound} \frac{2^n}{V\left(n, \lfloor \frac{d-1}{2} \rfloor\right)}.$$

Proof. We have already proved the Hamming bound. So let $m = A(n, d)$. Let C be a $[n, m, d]$-code. Then there does not exist $x \in \mathbb{F}_2^n$ with $d(x, c) \geq d$ for all $c \in C$ (otherwise could replace C with $C \cup \{x\}$, contradicting maximality of m). Hence

$$\mathbb{F}_2^n \subseteq \bigcup_{c \in C} \overline{B}(c, d-1) \implies 2^n \leq \sum_{c \in C} |\overline{B}(c, d-1)| = mV(n, d-1).$$

□

Example 6.2. $n = 10$, $d = 3$, have $V(n, 1) = 11$, $V(n, 2) = 56$. The above theorem gives $19 \leq \frac{2^{10}}{56} \leq A(10, 3) \leq \frac{2^{10}}{11} \leq 93$. It was known that $72 \leq A(10, 3) \leq 93$, but the exact value of $A(10, 3)$ was only found in 1999.

Asymptotics of $V(n, r)$

We study $\frac{\log A(n, \lfloor n\delta \rfloor)}{n}$ as $n \to \infty$ to see how large the information rate can be for a given error rate.

Proposition 6.6. *Let $\delta \in (0, 1/2)$. Then*

(i) $\log V(n, \lfloor n\delta \rfloor) \leq nH(\delta)$;

(ii) $\frac{1}{n} \log A(n, \lfloor n\delta \rfloor) \geq 1 - H(\delta)$.

Where $H(\delta) = -\delta \log \delta - (1 - \delta) \log(1 - \delta)$.

Proof. First we show (i) ⇒ (ii): by the GSV bound,

$$A(n, \lfloor n\delta \rfloor) \geq \frac{2^n}{V(n, \lfloor n\delta \rfloor - 1)} \geq \frac{2^n}{V(n, \lfloor n\delta \rfloor)}$$

and so

$$\frac{\log A(n, \lfloor n\delta \rfloor)}{n} \geq 1 - \frac{\log V(n, \lfloor n\delta \rfloor)}{n} \geq 1 - H(\delta).$$

Now we prove (i): $H(\delta)$ is increasing for $\delta < 1/2$, so wlog we may assume $n\delta \in \mathbb{Z}$. Now

$$\begin{aligned} 1 = (\delta + (1 - \delta))^n = \sum_{i=0}^{n} \binom{n}{i} \delta^i (1 - \delta)^{n-i} &\geq \sum_{i=0}^{n\delta} \binom{n}{i} \delta^i (1 - \delta)^{n-i} \\ &= (1 - \delta)^n \sum_{i=0}^{n\delta} \binom{n}{i} \left(\frac{\delta}{1 - \delta} \right)^i \\ &\geq (1 - \delta)^n \sum_{i=0}^{n\delta} \binom{n}{i} \left(\frac{\delta}{1 - \delta} \right)^{n\delta} \\ &= \delta^{n\delta} (1 - \delta)^{n(1-\delta)} V(n, n\delta). \end{aligned}$$

Now taking logs:

$$0 \geq n(\delta \log \delta + (1-\delta)\log(1-\delta)) + \log V(n, n\delta).$$

□

The constant $H(\delta)$ in the above bound best possible:

Lemma 6.7. *We have*

$$\lim_{n\to\infty} \frac{\log V(n, \lfloor n\delta \rfloor)}{n} = H(\delta).$$

Proof. Exercise. □

7 Constructing new codes from old

We're given C, a $[n, m, d]$-code. Can check the details in the following:

Examples.

1. The *parity check extension* C^+ is

$$\left\{ \left(c_1, \ldots, c_n, \sum_{i=1}^{n} c_i \right) : (c_1, \ldots, c_n) \in C \right\}$$

   Is a $[n+1, m, d']$-code with $d \leq d' \leq d+1$, depending on whether d is odd or even.

2. Fix $1 \leq i \leq n$. Deleting the ith digit from each codeword gives the *punctured code* C^-

$$\{(c_1, \ldots, c_{i-1}, c_{i+1}, \ldots, c_n) : (c_1, \ldots, c_n) \in C\}.$$

   If $d \geq 2$, then it is a $[n-1, m, d']$-code with $d - 1 \leq d' \leq d$.

3. Fix $1 \leq i \leq n$, and $\alpha \in \mathbb{F}_2$. The *shortened code* C' is

$$\{(c_1, \ldots, c_{i-1}, c_{i+1}, \ldots, c_n) : (c_1, \ldots, c_{i-1}, \alpha, c_{i+1}, \ldots, c_n) \in C\}.$$

   It has parameters $[n, m', d']$ with $d' \geq d$ and $m' \geq \frac{m}{2}$ for a suitable choice of α.

Shannon's Theorems

8 AEP and Shannon's first coding theorem

Definition. A *source* is a sequence of random variables $X_1, X_2, \ldots$ taking values in some alphabet $\mathcal{A}$. A source is *Bernouilli* (*memoryless*) if $X_1, X_2, \ldots$ are iid: write (X, X_n). A source $X_1, X_2, \ldots$ is *reliably encodable at rate* r if there exists a sequence of subsets $(A_n)_{n\geq 1}$ with $A_n \subseteq \mathcal{A}^n$ such that:

1. $\lim_{n\to\infty} \frac{\log A_n}{n} = r$;
2. $\lim_{n\to\infty} \mathbb{P}((X_1, \ldots, X_n) \in A_n) = 1$.

The *information rate* H of a source is the infimum of all reliable encoding rates. Exercise: $0 \leq H \leq \log |\mathcal{A}|$ with both bounds attainable.

Shannon's first coding theorem computes the information rate of certain sources, including Bernouilli sources.

Reminders from IA Probability:

We have a probability space $(\Omega, \mathcal{F}, \mathbb{P})$. A discrete random variable X is a function $X : \Omega \to \mathcal{A}$. The probability mass function $p_x : \mathcal{A} \to [0,1]$ is defined by $x \mapsto \mathbb{P}(X = x)$. Can consider $p \circ X = p(X) : \Omega \to [0,1]$, a random varible taking values in $[0,1]$.

Given a source $X_1, X_2, \ldots$ of random variables with values in $\mathcal{A}$, the probability mass function of $X^{(n)} = (X_1, \ldots, X_n)$ is $p_{X^{(n)}}$ given by $(x_1, \ldots, x_n) \mapsto \mathbb{P}((X_1, \ldots, X_n) = (x_1, \ldots, x_n))$. Since $p_{X^{(n)}} : \mathcal{A}^n \to [0,1]$ and $X^{(n)} : \Omega \to \mathcal{A}^n$, you can form $p(X^{(n)}) = p_{X^{(n)}} \circ X^{(n)} : \Omega \to [0,1]$.

Example 8.1. Let $\mathcal{A} = \{\mathrm{A}, \mathrm{B}, \mathrm{C}\}$. Suppose

$$X^{(2)} = \begin{cases} \mathrm{AB} & \text{with probability } 0.3 \\ \mathrm{AC} & \text{with probability } 0.1 \\ \mathrm{BC} & \text{with probability } 0.1 \\ \mathrm{BA} & \text{with probability } 0.2 \\ \mathrm{CA} & \text{with probability } 0.25 \\ \mathrm{CB} & \text{with probability } 0.05 \end{cases}.$$

Then

$$p(X^{(2)}) = \begin{cases} 0.3 & \text{with probability } 0.3 \\ 0.1 & \text{with probability } 0.2 \\ 0.2 & \text{with probability } 0.2 \\ 0.25 & \text{with probability } 0.25 \\ 0.05 & \text{with probability } 0.05 \end{cases}.$$

So some points are "lumped together".

Given a source $X_1, X_2, \ldots$ *converges in probability* to a random variable L (possibly constant) if for all $\varepsilon > 0$, $\mathbb{P}(|X_n - L| > \varepsilon) \to 0$ as $n \to \infty$. We write $X_n \xrightarrow{\mathbb{P}} L$.

The *Weak Law of Large Numbers* (WLLN) says that if $(X; X_n)$ are iid real-valued random variables with finite expectation $\mathbb{E}X$, we have

$$\frac{1}{n} \sum_{i=1}^{n} X_i \xrightarrow{\mathbb{P}} \mathbb{E}X.$$

Example 8.2. If $X_1, X_2, \ldots$ are iid Bernouilli, then $p(X_1), p(X_2), \ldots$ are iid random variables and $p(X_1, \ldots, X_n) = p(X_1) \ldots p(X_n)$. Note

$$-\frac{1}{n} \log p(X_1, \ldots, X_n) = -\frac{1}{n} \sum_{i=1}^{n} \log p(X_i) \xrightarrow{\mathbb{P}} -\mathbb{E}(-\log p(X_1)) = H(X_1) \text{ as } n \to \infty.$$

Lemma 8.1. *The information rate of a Bernouilli source $X_1, X_2, \ldots$ is at most the expected word length of an optimal code $c : \mathcal{A} \to \{0,1\}^*$ for X_1.*

Proof. Let $l_1, l_2, \ldots$ be the lengths of codewords when we encode $X_1, X_2, \ldots$ using c. Let $\varepsilon > 0$. Set $A_n = \{x \in \mathcal{A}^n : c^*(x) \text{ has length at less than } n(\mathbb{E}l_1 + \varepsilon)\}$. Then

$$\begin{aligned}\mathbb{P}((X_1, \ldots, X_n) \in A_n) &= \mathbb{P}\left(\sum_{i=1}^{n} l_i < n(\mathbb{E}l_1 + \varepsilon)\right) \\ &= \mathbb{P}\left(\left|\frac{1}{n}\sum_{i=1}^{n} l_i - \mathbb{E}l_1\right| < \varepsilon\right) \\ &\to 1 \text{ as } n \to \infty.\end{aligned}$$

Now, c is decipherable so c^* is injective. Hence $|A_n| \leq 2^{n(\mathbb{E}l_1+\varepsilon)}$. Making A_n larger if necessary, $|A_n| = \lfloor e^{n(\mathbb{E}l_1+\varepsilon)} \rfloor$ so

$$\frac{\log(A_n)}{n} \to \mathbb{E}l_1 + \varepsilon.$$

Hence $X_1, X_2, \ldots$ is reliably encodable at rate $r = \mathbb{E}l_1 + \varepsilon$ for all $\varepsilon > 0$. Hence the information rate is at most $\mathbb{E}l_1$. □

Corollary 8.2. *A Bernouilli source has information rate less than $H(X_1) + 1$.*

Proof. Combine the above with the Noiseless Coding Theorem. □

We encode $X_1, X_2, \ldots$ in blocks

$$\underbrace{X_1, \ldots, X_N}_{Y_1}, \underbrace{X_{n+1}, \ldots, X_{2N}}_{Y_2}, \cdots$$

so $Y_1, Y_2, \ldots$ take values in $\mathcal{A}^N$. Exercise: show that if $X_1, X_2, \ldots$ has information rate H then $Y_1, Y_2, \ldots$ has information rate NH.

Proposition 8.3. *The information rate H of a Bernouilli source $X_1, X_2, \ldots$ is at most $H(X_1)$.*

Proof. Apply the previous corollary to $Y_1, Y_2, \ldots$ and obtain

$$NH < H(Y_1)+1 = H(X_1, \ldots, X_N)+1 = \sum_{i=1}^{N} H(X_i)+1 = NH(X_1, \ldots, X_n)+1.$$

Hence $H < H(X_1) + \frac{1}{N}$. Since N is arbitrary, $H \leq H(X_1)$. □

Definition. A source $X_1, X_2, \ldots$ satisfies the *Asympotitic Equipartition Property* (AEP) for some constant $H \geq 0$ if

$$-\frac{1}{n}\log p(X_1, X_2, \ldots) \xrightarrow{\mathbb{P}} H \text{ as } n \to \infty.$$

Example 8.3. Tossing a biased coin, $\mathbb{P}(\mathrm{H}) = p$. Let $(X; X_n)$ be the results of independent coin tosses. After a large number N of tosses, expect on average pN heads and $(1-p)N$ tails. The probability of any particular sequence of pN heads and $(1-p)N$ tails is $p^{pN}(1-p)^{(1-p)N} = 2^{N(p\log p+(1-p)\log(1-p))} = 2^{-NH(X)}$. Not every sequence of tosses will be like this, but there is only a small probability of "atypical" sequences. With high probability we get a "typical" sequence and its probability will be close to $2^{-NH(X)}$.

Lemma 8.4. *The AEP for a source $X_1, X_2, \ldots$ is equivalent to the following property*

$$\forall \varepsilon > 0\ \exists n_0(\varepsilon) \text{ such that } \forall n \geq n_0(\varepsilon)\ \exists \text{ a "typical set" } T_n \subseteq \mathcal{A}^n \text{ such that}$$

(i) $\mathbb{P}((X_1, \ldots, X_n) \in T_n) > 1 - \varepsilon$;

(ii) $2^{-n(H+\varepsilon)} \leq p(x_1, \ldots, x_n) \leq 2^{-n(H-\varepsilon)}$ *for all* $(x_1, \ldots, x_n) \in T_n$.

Proof. Obvious and non-examinable. □

Theorem 8.5 (Shannon's First Coding Theorem (FCT))**.** *If a source $X_1, X_2, \ldots$ satisfies the AEP with constant H, then the source has information rate H.*

Proof. Let $\varepsilon > 0$ and let $T_n \subseteq \mathcal{A}^n$ be typical sets. Then for some $n_0(\varepsilon)$ and all $n \geq n_0(\varepsilon)$

$$p(x_1, \ldots, x_n) \geq 2^{-n(H+\varepsilon)} \text{ for all } (x_1, \ldots, x_n) \in T_n$$

$$\implies \mathbb{P}(T_n) \geq 2^{-n(H+\varepsilon)}|T_n| \implies 1 \geq 2^{-n(H+\varepsilon)}|T_n| \implies \frac{\log |T_n|}{n} \leq H + \varepsilon.$$

Taking $A_n = T_n$, shows the source is reliably encodable at rate $H + \varepsilon$. Conversely, if $H = 0$, we're done. Otherwise pick $0 < \varepsilon < H/2$ and suppose for contradiction that the source is reliably encodable at rate $H - 2\varepsilon$, say with sets $A_n \subseteq \mathcal{A}^n$. Let $T_n \subseteq \mathcal{A}^n$ be typical sets. Then for all $(x_1, \ldots, x_n) \in T_n$

$$p(x_1, \ldots, x_n) \leq 2^{-n(H-\varepsilon)}$$

$$\implies \mathbb{P}(A_n \cap T_n) \leq 2^{-n(H-\varepsilon)}|A_n|$$

$$\implies \frac{\log \mathbb{P}(A_n \cap T_n)}{n} \leq H - \varepsilon + \frac{\log |A_n|}{n} \to -(H - \varepsilon) + H - 2\varepsilon = -\varepsilon.$$

Hence $\log \mathbb{P}(A_n \cap T_n) \to -\infty$ and $\mathbb{P}(A_n \cap T_n) \to 0$. But $\mathbb{P}(T_n) \leq \mathbb{P}(A_n \cap T_n) + \mathbb{P}(\mathcal{A}^n \setminus A_n) \to 0$, a contradiction to $\mathbb{P}(T_n) \to 1$. Thus the information is rate is exactly H. □

Corollary 8.6. *A Bernouilli source $X_1, X_2, \ldots$ has information rate $H(X_1)$.*

Proof. We've already seen that $-\frac{1}{n} \log p(X_1, \ldots, X_n) \xrightarrow{\mathbb{P}} H(X_1)$, so done by Shannon's First Coding Theorem. □

Remarks.

- The AEP is useful for noiseless coding. We can
  - encode the typical sequences using a block code;
  - encode the atypical sequences arbitrarily.
- Many sources, which are not necessarily Bernouilli satisfy the AEP. Under suitable hypotheses the sequence $\frac{1}{n} H(X_1, \ldots, X_n)$ is decreasing and the AEP is satisfied.

9 Capacity & Shannon's second coding theorem

Recall:

Definition. We model n uses of a channel by the nth extension, with input alphabet $\mathcal{A}^n$ and output alphabet $\mathcal{B}^n$. A *code C of length n* is a function $\mathcal{M} \to \mathcal{A}^n$ where $\mathcal{M}$ is the set of possible messages. Implicitly we also have a decoding rule $\mathcal{B}^n \to \mathcal{M}$. The *size* of C is $m = |\mathcal{M}|$. The *information rate* is $\rho(C) = \frac{1}{n} \log_2 m$. The *error rate* is $\hat{e}(C) = \max_{x \in \mathcal{M}} \mathbb{P}(\text{error}|x \text{ sent})$.

Definition. A channel can *transmit reliably at rate* R if there exists $(C_n)_{n=1}^{\infty}$ with each C_n a code of length n such that

$$\lim_{n\to\infty} \rho(C_n) = R \;\&\; \lim_{n\to\infty} \hat{e}(C_n) = 0.$$

The *capacity* is the supremum of all reliable transmission rates.

Suppose we are given a souce where

- it has information rate r bits per second;
- it emits symbols at s symbols per second.

Suppose we are also given a channel where

- it has capacity R bits per transmission;
- it transmits symbols at S transmissions per second.

Usually, information theorists take $S = s = 1$. If $rs \leq RS$ then you can encode and transmit reliably, and if $rs > RS$ you cannot.

We'll compute the capacity of a BSC with error probability p.

Proposition 9.1. *A binary symmetric channel with error probability $p < 1/4$ has non-zero capacity.*

Proof. We use the GSV bound. Pick $\delta \in (2p, 1/2)$. We claim reliable transmission at rate $R = 1 - H(\delta) > 0$. Let C_n be a code of length n, and suppose it has minimum distance $\lfloor n\delta \rfloor$ of maximal size. Then (by Proposition 6.6(ii))

$$|C_n| = A(n, \lfloor n\delta \rfloor) \geq 2^{n(1-H(\delta))}.$$

Replacing C_n by a subcode, we can assume $|C_n| = \lfloor 2^{nR} \rfloor$ and still minimum distance $\geq \lfloor n\delta \rfloor$. Using minimum distance decoding

$$\begin{aligned}\hat{e}(C_n) &\leq \mathbb{P}\left(\text{in } n \text{ uses, BSC makes} \geq \left\lfloor \frac{\lfloor n\delta - 1 \rfloor}{2} \right\rfloor \text{ errors}\right) \\ &\leq \mathbb{P}\left(\text{in } n \text{ uses, BSC makes} \geq \left\lfloor \frac{n\delta - 1}{2} \right\rfloor \text{ errors}\right).\end{aligned}$$

Pick $\varepsilon > 0$ with $p + \varepsilon < \frac{\delta}{2}$. For n sufficiently large, $\frac{n\delta - 1}{2} = n\left(\frac{\delta}{2} - \frac{1}{2n}\right) > n(p+\varepsilon)$. Hence $\hat{e}(C_n) \leq \mathbb{P}(\text{BSC makes} \geq n(p+\varepsilon) \text{ errors}) \to 0$, using the next lemma. □

Lemma 9.2. *Let $\varepsilon > 0$/ A BSC with error probability p is used to transmit n digits. Then*

$$\lim_{n\to\infty} \mathbb{P}(BSC \text{ makes} \geq n(p+\varepsilon) \text{ errors}) = 0.$$

Proof. Consider random variables U_i taking value 1 if the ith digit is mistransmitted, and value 0 otherwise. Then $\mathbb{E}U_i = p$, so $\mathbb{P}(BSC$ makes $\geq n(p+\varepsilon)$ errors$) \leq \mathbb{P}\left(\left|\frac{1}{n}\sum_{i=1}^{n} - p\right| \geq \varepsilon\right) \to 0$. □

Definition (Conditional Entropy)**.** Let X, Y be random variables taking values in alphabets $\mathcal{A}, \mathcal{B}$ respectively. Then we define the *conditional entropy*

$$H(X|Y=y) = -\sum_{x\in\mathcal{A}} \mathbb{P}(X=x|Y=y)\log\mathbb{P}(X=x|Y=y)$$

$$H(X|Y) = \sum_{y\in\mathcal{B}} \mathbb{P}(Y=y)H(X|Y=y).$$

Clearly $H(X|Y) \geq 0$.

Lemma 9.3. *We have*

$$H(X,Y) = H(X|Y) - H(Y).$$

Proof.

$$\begin{aligned}
H(X|Y) &= -\sum_{y\in\mathcal{B}}\sum_{x\in\mathcal{A}} \mathbb{P}(X=x|Y=y)\mathbb{P}(Y=y)\log\mathbb{P}(X=x|Y=y) \\
&= -\sum_{y\in\mathcal{B}}\sum_{x\in\mathcal{A}} \mathbb{P}(X=x,Y=y)\log\left(\frac{\mathbb{P}(x=x,Y=y)}{\mathbb{P}(Y=y)}\right) \\
&= -\sum_{y\in\mathcal{B}}\sum_{x\in\mathcal{A}} \mathbb{P}(X=x,Y=y)\log\mathbb{P}(X=x,Y=y) \\
&\quad + \sum_{y\in\mathcal{B}}\sum_{x\in\mathcal{A}} \mathbb{P}(X=x,Y=y)\log\mathbb{P}(Y=y) \\
&= H(X,Y) - \sum_{y\in\mathcal{B}} \mathbb{P}(Y=y)\log\mathbb{P}(Y=y) \\
&= H(X,Y) - H(Y).
\end{aligned}$$

□

Example 9.1. We roll fair die. Let X be the value defined by the roll, let Y be equal to 0 if X is even and 1 if X is odd. Then $H(X,Y) = H(X) = \log 6$, $H(Y) = \log 2 = 1$, $H(X|Y) = \log 3 = H(X,Y) - H(Y)$. Similarly $H(Y|X) = 0 = H(X,Y) - H(X)$.

Corollary 9.4. $H(X|Y) \leq H(X)$.

Proof. Combine the previous with Lemma 4.1. □

Now replace X, Y with random vectors $X^{(r)} = (X_1, \ldots, X_r)$, $Y^{(s)} = (Y_1, \ldots, Y_s)$. Similarly can define $H(X_1, \ldots, X_r | Y_1, \ldots, Y_s) = H(X^{(r)} | Y^{(s)})$.

Remark. $H(X, Y|Z)$ is the entropy of (X, Y) given Z, not the entropy of X and $Y|Z$.

Lemma 9.5. *Let X, Y, Z be random variables. Then $H(X, Y) \leq H(X|Y, Z) + H(Z)$.*

Proof. We expand $H(X, Y, Z)$ in two different ways.

$$H(X, Y, Z) = H(Z|X, Y) + H(X|Y) + H(Y),$$

$$H(X, Y, Z) = H(X|Y, Z) + H(Z|Y) + H(Y).$$

Since $H(Z|X, Y) \geq 0$, we have

$$\begin{aligned} H(X|Y) &\leq H(X|Y, Z) + H(Z|Y) \\ &\leq H(X|Y, Z) + H(Z). \end{aligned}$$

□

Proposition 9.6 (Fano's inequality)**.** *Let X, Y be random variables taking values in $\mathcal{A}$, $|\mathcal{A}| = m$. Let $p = \mathbb{P}(X \neq Y)$. Then*

$$H(X|Y) \leq H(p) + p \log(m - 1).$$

Proof. Define $Z = 0$ if $X = Y$ and $Z = 1$ if $X \neq Y$. Then $\mathbb{P}(Z = 0) = 1 - p$, $\mathbb{P}(Z = 1) = p$. So $H(Z) = H(p)$. By the previous lemma,

$$H(X|Y) \leq H(p) + H(X|Y, Z). \qquad (*)$$

Since $Z = 0$ implies $X = Y$, $H(X|Y = y, Z = 0) = 0$. Also there are $m - 1$ remaining possibilities for X, so $H(X|Y = y, Z = 1) \leq \log(m - 1)$. Therefore

$$\begin{aligned} H(X|Y) &= \sum_{y,z} \mathbb{P}(Y = y, Z = z) H(X|Y = y, Z = z) \\ &\leq \sum_{y \in \mathcal{A}} \mathbb{P}(Y = y, Z = 1) \log(m - 1) \\ &= \mathbb{P}(Z = 1) \log(m - 1) \\ &= p \log(m - 1). \end{aligned}$$

So by $(*)$, $H(X|Y) \leq H(p) + p \log(m - 1)$. □

Remark. $H(p)$ represents the information needed to decide whether or not there is an error, and $p \log(m - 1)$ represents the information needed to resolve the error assuming the worst possible case.

Definition. Let X, Y be random variables. Then the *mutual information* is $I(X; Y) := H(X) - H(X|Y)$. By Corollary 9.4, $I(X; Y) \geq 0$ with equality if and only if X, Y are independent. Clearly $I(X; Y) = I(Y; X)$.

Suppose we're given a discrete memoryless channel (DMC) with input alphabet $\mathcal{A}$, $|\mathcal{A}| = m$ and output alphabet $\mathcal{B}$. Let X be a random variable taking values in $\mathcal{A}$ and used as input to the channel. Let Y be a random variable output, depending on X and the channel matrix.

Definition. The *(information) capacity* is $\max_X I(X;Y)$.

Remarks.

- The maximum is over all probability distributions $(p_1, \ldots, p_m)$ for X on $\mathcal{A}$.
- The maximum is attained since $(p_1, \ldots, p_m) \mapsto I((p_1, \ldots, p_m); Y)$ is continuous on the compact set

$$\{(p_1, \ldots, p_m) \in [0,1]^m : p_1 + \ldots + p_m = 1\}.$$

- The information capacity depends only on the channel matrix.

Theorem 9.7 (Shannon's Second Coding Theorem)**.** *For a DMC, the operational capacity is equal to the information capacity.*

Remark. We'll show one inequality in general and the other for the BSC only.

Assuming Shannon's Second Coding Theorem, let us compute he capacity of certain channels.

Example 9.2. BSC, error probability p. Input X: $\mathbb{P}(X=0)=\alpha$, $\mathbb{P}(X=1)=1-\alpha$. Output Y: $\mathbb{P}(Y=0)=\alpha(1-p)+(1-\alpha)p$, $\mathbb{P}(Y=1)=(1-\alpha)(1-p)+\alpha p$. Then C is

$$\begin{aligned}\max_{\alpha} I(X;Y) &= \max_{\alpha}\left(H(Y)-H(Y|X)\right)\\ &= \max_{\alpha}\left(H(\alpha(1-p)+(1-\alpha)p)-H(p)\right)\\ &= 1-H(p).\end{aligned}$$

Where the maximum is attained for $\alpha=1/2$. Hence $C=1+p\log p+(1-p)\log(1-p)$.

Remark. We can choose to calculate either $H(Y)-H(Y|X)$ or $H(X)-H(X|Y)$ depending on which is easier.

Example 9.3. BEC, erasure probability p. Input X: $\mathbb{P}(X=0)=\alpha$, $\mathbb{P}(X=1)=1-\alpha$. Output Y: $\mathbb{P}(Y=0)=\alpha(1-p)$, $\mathbb{P}(Y=*)=p$, $\mathbb{P}(Y=1)=(1-\alpha)(1-p)$. Can calculate $H(X|Y=0)=0$, $H(X|Y=1)=0$, $H(X|Y=*)=H(\alpha)$, which gives $H(X|Y)=pH(\alpha)$. So

$$\begin{aligned}C=\max_{\alpha} I(X;Y)=\max_{\alpha}(H(X)-H(X|Y)) &= \max_{\alpha}(H(\alpha)-pH(\alpha))\\ &= (1-p)\max_{\alpha} H(\alpha)\\ &= 1-p.\end{aligned}$$

Where the max is attained for $\alpha=1/2$.

Now model using a channel n times as the *nth extension*, i.e replace alphabets $\mathcal{A},\mathcal{B}$ with $\mathcal{A}^n,\mathcal{B}^n$,

$$\mathbb{P}(y_1,\dots,y_n \text{ received}|x_1,\dots,x_n \text{ sent})=\prod_{i=1}^{n}\mathbb{P}(y_i|x_i).$$

Lemma 9.8. *The nth extension of a DMC with information capacity C has information capacity nC.*

Proof. Take random variable input $X_1,\dots,X_n$ producing output $Y_1,\dots,Y_n$. Since the channel is memoryless,

$$H(Y_1,\dots,Y_n|X_1,\dots,X_n)=\sum_{i=1}^{n}H(Y_i|X_1,\dots,X_n)=\sum_{i=1}^{n}H(Y_i|X_i).$$

Hence

$$\begin{aligned} I(X_1,\ldots,X_n;Y_1,\ldots,Y_n) &= H(Y_1,\ldots,Y_n) - H(Y_1,\ldots,Y_n|X_1,\ldots,X_n) \\ &= H(Y_1,\ldots,Y_n) - \sum_{i=1}^{n} H(Y_i|X_i) \\ &\leq \sum_{i=1}^{n} (H(Y_i) - H(Y_i|X_i)) \\ &= \sum_{i=1}^{n} I(X_i;Y_i) \leq nC. \end{aligned}$$

To finish, we find a distribution for $X_1,\ldots,X_n$ giving equality. Equality is attained by taking $X_1,\ldots,X_n$ independent, each of the same distribution such that $I(X_i;Y_i) = C$. Indeed, if $X_1,\ldots,X_n$ are independent, so are $Y_1,\ldots,Y_n$. □

Proposition 9.9. *For a DMC the (operational) capacity is at most the information capacity. Let C be the information capacity. Suppose reliable transmission is possible at some rate $R > C$, i.e there exist a sequence of codes $(C_n)_{n\geq 1}$ with C_n having length n and size $\lceil 2^{nR} \rfloor$ for all n such that*

$$\lim_{n\to\infty} \rho(C_n) = R \text{ and } \lim_{n\to\infty} \hat{e}(C_n) = 0.$$

(Recall $\hat{e}(C_n) = \max_{c\in C_n} \mathbb{P}(\text{error}|c \text{ sent})$.)

Proof. Define the *average error rate* $e(C_n) = \frac{1}{|C_n|}\sum_{c\in C_n} \mathbb{P}(\text{error}|c \text{ sent})$. Note $e(C_n) \leq \hat{e}(C_n)$ and so $e(C_n) \to 0$ as $n \to \infty$. Take input random variable X equidistributed over C_n (i.e takes all values with equal probability). Let Y be the random variable output when X is transmitted and decoded. So $e(C_n) = \mathbb{P}(X \neq Y)$, define $p := e(C_n)$. Now we have

$$H(X) = \log|C_n| = \log\lceil 2^{nR}\rceil \geq nR - 1 \text{ for sufficiently large } n.$$

And

$$\begin{aligned} H(X|Y) &\leq H(p) + p\log(|C_n| - 1) && \text{(Fano's inequality)} \\ &\leq 1 + pnR. \end{aligned}$$

Recall $I(X;Y) = H(X) - H(X|Y)$. By the previous lemma $nC \geq I(X;Y)$ so

$$nC \geq nR - 1 - (1 + pnR)$$

$$\implies pnR \geq n(R-C) - 2 \implies p \geq \frac{n(R-C)-2}{nR} \to \frac{R-C}{R} \neq 0.$$

Since $R > C$, which contradicts $p \to 0$. Hence we conclude that we cannot transmit reliably at any rate exceeding C. □

To complete the proof of Shannon's SCT for a BSC with error probability p, we need to show that the operational capaity is at least $1-H(p)$ (i.e the information capacity).

Proposition 9.10. *Consider BSC with error probability p. Suppose $R < 1 - H(p)$. Then there is a sequence of codes $(C_n)_{n\geq 1}$ with C_n of length n and size $\lceil 2^{nR} \rceil$ for all n such that*

$$\lim_{n\to\infty} \rho(C_n) = R \text{ and } \lim_{n\to\infty} e(C_n) = 0.$$

Remark. Note that the above proposition deals with the average error rate e, not the error rate $\hat{e}$.

Proof. We'll use the method of 'random coding'. Wlog $p < 1/2$. Take $\varepsilon > 0$ such that

$$p + \varepsilon < 1/2 \text{ and } R < 1 - H(p + \varepsilon).$$

(This is possible by continuity of H). We use minimum distance decoding (in case of a tie, make arbitrary choice). Let $m = \lfloor 2^{nR} \rfloor$. Pick $C = \{c_1, \ldots, c_m\}$ at random from $\mathcal{C} = \{[n, m]\text{-codes}\}$, a set of size $\binom{2^n}{m}$. Choose $1 \leq i \leq m$ at random (i.e each with probability $1/m$). We sent c_i through the channel and get output Y. Then

$$\begin{aligned}\mathbb{P}(Y \text{ not decoded as } c_i) &= \text{average value of } e(C) \text{ as } C \text{ runs through } \mathcal{C} \\ &= \frac{1}{|\mathcal{C}|} \sum_{C\in\mathcal{C}} e(C).\end{aligned}$$

We can pick C_n, an $[n, m]$-code with $e(C_n)$ at most this average. So it is enough to show

$$\mathbb{P}(Y \text{ is not decoded as } c_i) \to 0 \text{ as } n \to \infty.$$

Let $r = \lfloor n(p+\varepsilon) \rfloor$. Then if $B(Y, r) \cap C = \{c_i\}$, Y is correctly decoded as c_i. So

$$\mathbb{P}(Y \text{ is not decoded as } c_i) \leq \mathbb{P}(c_i \notin B(Y, r)) + \mathbb{P}(B(Y, r) \cap C \supsetneq \{c_i\}).$$

Consider the two cases separately, according as $d(c_i, Y) > r$ (case (i)) or $d(c_i, Y) \leq r$ (case (ii)). For (i) we have

$$\begin{aligned}\mathbb{P}(d(c_i, Y) > r) &= \mathbb{P}(\text{BSC makes} > r \text{ errors}) \\ &= \mathbb{P}(\text{BSC makes} > n(p + \varepsilon) \text{ errors}) \\ &\to 0 \text{ as } n \to \infty \text{ by a previous theorem.}\end{aligned}$$

For (ii), if $j \neq i$, $\mathbb{P}(c_j \in B(Y, r) | c_i \in B(Y, r)) = \frac{V(n,r)-1}{2^n-1} \leq \frac{V(n,r)}{2^n}$ (since

$V(n,r) \leq 2^n$). Hence

$$\begin{aligned}
\mathbb{P}(B(Y,r) \cap C \supsetneq \{c_i\}) &\leq \sum_{j \neq i} \mathbb{P}(c_j \in B(Y,r) \text{ and } c_i \in B(Y,r)) \\
&\leq \sum_{j \neq i} \mathbb{P}(c_j \in B(Y,r) | c_i \in B(Y,r)) \\
&\leq (m-1)\frac{V(n,r)}{2^n} \\
&\leq \frac{mV(n,r)}{2^n} \\
&\leq 2^{nR} 2^{nH(p+\varepsilon)2^{-n}} \\
&= 2^{n(R-(1-H(p+\varepsilon)))} \xrightarrow{n\to\infty} 0
\end{aligned}$$

since $R < 1 - H(p+\varepsilon)$. □

Proposition 9.11. *We can replace e by $\hat{e}$ in the previous result.*

Proof. Pick R' with $R < R' < 1 - H(p)$. The previous result constructs a sequence $(C'_n)_{n \geq 1}$ of codes, C'_n of length n and size $\lfloor 2^{nR'} \rfloor$ for all n, and $e(C'_n) \to 0$ as $n \to \infty$.

For each n, order the codewords of C'_n by $\mathbb{P}(\text{error}|c \text{ sent})$ and delete the worst half, to get a code C_n with $\hat{e}(C_n) \leq 2e(C_n)$. So $\hat{e}(C_n) \to 0$ as $n \to \infty$. Since C_n has length n and size $\lfloor 2^{nR'} \rfloor / 2 = floor 2^{nR'-1}$. But $2^{nR'-1} = 2^{n(R'-1/n)} \geq 2^{nR}$ for all n sufficiently large. So can replace C'_n by a code of smaller size $\lfloor 2^{nR} \rfloor$ and still get $\hat{e}(C_n) \to 0$ and $\rho(C_n) \to R$. □

Remarks.

1. A BSC with error probability p has operational capacity $1 - H(p)$. Indeed the above proposition says we can transmit reliably at any rate $R < 1 - H(p)$, wheras 9.9 says the capacity is at most $1 - H(p)$.
2. Shannon's SCT shows that good codes exist, but the proof does not say how to construct them.

10 The Kelly criterion

Let $1 > p > 0$, $u > 0$, $1 > w \geq 0$. Suppose a coin is tossed n times in succession, $\mathbb{P}(\text{H}) = p$. If I pay k ahead of a particular throw, then get back ku if the throw is heads, or get back 0 if throw is tails.

Say we have initial bankroll $X_0 = 1$. If we have bankroll of X_n after nth toss, bet wX_n, retaining $(1-w)X_n$. Then bankroll X_{n+1} after $(n+1)$th throw is

$$X_{n+1} = \begin{cases} X_n(wu + (1-w)) & \text{if } nth \text{ toss is heads} \\ X_n(1-w) & \text{if } nth \text{ toss is tails} \end{cases}.$$

Let $Y_{n+1} = X_{n+1}/X_n$. Then $Y_1, Y_2, \ldots$, is a sequence of iid random variables. Hence $\log Y_1, \log Y_2, \ldots$ is a sequence of iid random variables. Note $\log X_n = \sum_{i=1}^{n} \log Y_i$.

Lemma 10.1. *Let* $\mathbb{E} \log Y_1 = \mu$, $\mathrm{Var}(\log Y_1) = \sigma^2$. *Then if* $a > 0$

(i) $\mathbb{P}\left(\left|\frac{1}{n}\sum_{i=1}^{n} \log Y_i - \mu\right| \geq a\right) \leq \frac{\sigma^2}{na^2}$

(ii) $\mathbb{P}\left(\left|\frac{\log X_n}{n} - \mu\right| \geq a\right) \leq \frac{\sigma^2}{na^2}$

(iii) Given $\varepsilon > 0$ *and* $\delta > 0$ *there exists* $N \in \mathbb{N}$ *such that* $\mathbb{P}\left(z \left|\frac{\log X_n}{n} - \mu\right| \geq \delta\right) \leq \varepsilon$ *for all* $n \geq N$.

Proof. Obvious by Chebyshev. □

Lemma 10.2. *Conside a single toss of a coin,* $\mathbb{P}(H) = p < 1$. *Suppose a bet of* k *on* H *has payout ratio* $u > 0$. *Suppose we have initial bankroll* 1, *and we bet* w *on* H, *retaining* $1 - w$ *(*$0 \leq w < 1$*). If* $Y =$ *fortune after toss then* $\mathbb{E} \log Y = p \log(1 + (u-1)w) + (1-p)\log(1-w)$. *The value of* $\mathbb{E} \log Y$ *is maximised by taking* $w = 0$ *if* $up \leq 1$, *and otherwise taking* $w = \frac{up-1}{u-1}$.

Proof. Not given (see Körner's book). □

Remark. Put $q = 1 - p$¿ If $up > 1$, at the optimum w have

$$\mathbb{E} \log Y = p \log p + q \log q + \log u - q \log(u-1)$$

and $-(p \log p + q \log q)$ is the entropy of a simple probabilistic system.

The suggestion that those making a long sequence of bets should maximise the expectation of the log is called Kelly's criterion.
Exercise: show that if you bet less than the optimum, your bankroll will tend to increase, but more slowly; and if you bet more than some critical proportion, your bankroll will tend to decrease.

11 Algebraic Coding Theory

11.1 Linear codes

Definition. A code $C \subseteq \mathbb{F}_2^n$ is *linear* if

- $0 \in C$;
- whenever $x, y \in C$, we have $x + y \in C$.

i.e C is a $\mathbb{F}_2$-vector subspace of $\mathbb{F}_2^n$.

Definition. The *rank* of a linear code C is its dimension as a $\mathbb{F}_2$-vector space. A linear code of length n, rank k is called a (n, k)-code. If the minimum distance of the code is d, it is a (n, k, d)-code.

Let $v_1, \ldots, v_k$ be a basis of C. Then $C = \{\sum_{i=1}^{k} \lambda_i v_i : \lambda_1, \ldots, \lambda_n \in \mathbb{F}_2\}$. So $|C| = 2^k$. So an (n, k)-code is a $[n, 2^k]$-code. The information rate is k/n.

Definition. The *weight* of $x \in \mathbb{F}_2^n$ is $w(x) = d(x, 0)$.

Lemma 11.1. *The minimum distance of a linear code is the minimum weight of a non-zero code word.*

Proof. Obvious. □

Definition. For $x, y \in \mathbb{F}_2^n$, let $x \cdot y = \sum_{i=1}^{n} x_i y_i \in \mathbb{F}_2$. This operation is symmetric and bilinear, but not positive definite.

Definition. Take some $P \subseteq \mathbb{F}_2^n$. The *parity check code* defined by P is

$$C = \{x \in \mathbb{F}_2^n : p \cdot x = 0, \ \forall p \in P\}.$$

Examples.

- $P = \{(1, \ldots, 1)\}$ gives the simple parity check code;
- $P = \{(1,0,1,0,1,0,1), (0,1,1,0,0,1,1), (0,0,0,1,1,1,1)\}$ gives Hamming's $[7, 16, 3]$-code;
- Show that C^+ and C^- are linear if C is linear. When is C' linear?

Lemma 11.2. *Every parity check code is linear.*

Proof. Obvious. □

Definition. Take $C \subseteq \mathbb{F}_2^n$ a linear code. The *dual code*

$$C^\perp = \{x \in \mathbb{F}_2^n : x \cdot y, \ \forall y \in C\}.$$

This is a parity-check code, so is linear. Note $C \cap C^\perp$ is not necessarily $\{0\}$.

Lemma 11.3. $\dim(C) + \dim(C^\perp) = n$.

Proof. See later. □

Corollary 11.4. *Let C be linear. Then $(C^\perp)^\perp = C$. In particular, C is a parity check code.*

Proof. Clearly $C \subseteq (C^\perp)^\perp$. Equality follows from the previous lemma. □

Definition. Let C be a (n, k)-code.

(i) A *generator matrix* G for C is a $k \times n$ matrix with rows a basis of C;

(ii) A *parity check matrix* H for C is a generator matrix for $C^\perp$. It is a $(n - k) \times n$ matrix.

The codewords of C can be viewed either as:

- Linear combinations of rows of G;
- Linear dependence relations between the columns of H, i.e $C = \{x \in \mathbb{F}_2^n : Hx = 0\}$.

Syndrome decoding

Definition. Let C be a (n,k)-code. The *syndrome* of $x \in \mathbb{F}_2^n$ is Hx. If we recieve $x = c + z$, where c is a codeword and z is the 'error pattern', the syndrome $Hx = H(c+z) = Hc + Hz = Hz$. If C is e-error correcting, we precompute a list of all possible Hz for all z with $w(z) \leq e$. On receiving $x \in \mathbb{F}_2^n$ we look for Hx in our list. If successful, $Hx = Hz$ so $H(x-z) = 0$ and we decode x as $c = x - z \in C$ with $d(x,c) = w(z) \leq e$.

Definition. We say codes $C_1, C_2 \subseteq \mathbb{F}_2^n$ are *equivalent* if reordering each codeword of C_1 using the same permutation gives the codewords of C_2. i.e $\sigma C_1 = C_2$ for some $\sigma \in S_n$. Usually we only consider codes up to equivalence.

Lemma 11.5. *Every (n,k)-linear code is equivalent to one with generator $G = (I_k|B)$ for some $k \times (n-k)$ matrix B.*

Proof. Given a $k \times n$ generator matrix G for C, using row operations, we can transform G into row-echelon form, i.e

$$G_{ij} = \begin{cases} 0 & j \leq l(i) \\ 1 & j = l(i) \end{cases}$$

for some $l(1) < l(2) < \ldots < l(k)$. Permuting the columns replaces code by an equivalent code, so wlog we can assume $l(i) = i$ for all $1 \leq i \leq k$ which gives the required form upon applying further row operations. □

Remark. A message $y \in \mathbb{F}_2^k$, viewed as a row vector is encoded as yG. So if $G = (I_k|B)$, then $yG = (y|yB)$, where y is the message and yB are check digits.

Now we will prove

Lemma 11.6. $\dim(C) + \dim(C^\perp) = n$.

Proof. Wlog C has generator matrix $G = (I_k|B)$. G has k linearly independent columns so the linear map $\gamma : \mathbb{F}_2^n \to \mathbb{F}_2^k$ defined by $x \mapsto Gx$ is surjective and $\mathrm{Ker}(\gamma) = C^\perp$. By the rank-nullity theorem, $\dim \mathbb{F}_2^n = \dim \mathrm{Ker}(\gamma) + \dim \mathrm{Im}(\gamma) = n$, i.e $n = \dim(C) + \dim(C^\perp)$. □

Lemma 11.7. *An (n,k)-linear code with generator matrix $G = (I_k, B)$ has parity check matrix $H = (B^T|I_{n-k})$.*

Proof. $GH^T = (I_k|B)\left(\frac{B}{I_{n-k}}\right) = B + B = 0$. So the rows of H generate a subcode of $C^\perp$. But $\mathrm{rank}(H) = n - k$ and $n - k = \dim C^\perp$. Hence $C^\perp$ has generator matrix H, and H is a parity check matrix for C. □

Lemma 11.8. *Let C be a linear code with parity check matrix H. Then $d(C) = d$ if and only if*

(i) any $(d-1)$ columns of H are linearly independent;

(ii) some set of d columns of H are linearly dependent.

Proof. Exercise. □

12 Examples of codes

Hamming codes

Definition. For $d \geq 1$, let $n = 2^d - 1$. Let H be the $d \times n$ matrix whose columns are the non-zero elements of $\mathbb{F}_2^d$. The *Hamming* $(n, n-d)$*-linear code* is the code with parity matrix H (original code is $d = 3$).

Lemma 12.1. *The Hamming* $(n, n-d)$*-code* C *has minimum distance* $d(C) = 3$*, is perfect and is* 1*-error correcting.*

Proof. Any two columns of H are linearly independent, but there are 3 columns that are linearly dependent, hence $d(C) = 3$ by the pervious lemma. By a previous result, C is $\lfloor \frac{3-1}{2} \rfloor = 1$-error correcting.

Note that $n = 2^d - 1$, $e = 1$ so $V(n, e) = 1 + 2^d - 1 = 2^d$ and $|C| = 2^{n-d} = \frac{2^n}{2^d} = \frac{2^n}{V(n,e)}$ and C is perfect. □

Reed-Muller codes

Take $X = \{p_1, \ldots, p_n\}$, set of size n. There is a correspondence between $\mathcal{P}(X)$ and $\mathbb{F}_2^n$ via $\mathcal{P}(X) \to \{f : X \to \mathbb{F}_2\} \to \mathbb{F}_2^n$ the composition of $A \mapsto \mathbb{1}_A$ and $f \mapsto (f(p_1), \ldots, f(p_n))$. The symmetric difference then corresponds to vector addition, and the intersection $A \cap B$ to the wedge product $x \wedge y = (x_1 y_1, \ldots, x_n y_n)$.

Take $X = \mathbb{F}_2^d$, so $n = 2^d = |X|$. Let $v_0 = (1, 1, 1, \ldots, 1)$. Let $v_i = \mathbb{1}_{H_i}$ where $H_i = \{p \in X : p_i = 0\}$ for $1 \leq i \leq d$.

Definition. Let $0 \leq r \leq d$. The *Reed-Muller code* $\mathrm{RM}(d, r)$ of order r and length 2^d is the linear code spanned by v_0 and all wedge products of r or fewer of the v_i. By convention, the empty wedge product is v_0.

Example 12.1. $d = 3$, $X = \mathbb{F}_2^3 = \{p_1, \ldots, p_8\}$ in binary order.

X	000	001	010	011	100	101	110	111
v_0	1	1	1	1	1	1	1	1
v_1	1	1	1	1	0	0	0	0
v_2	1	1	0	0	1	1	0	0
v_3	1	0	1	0	1	0	1	0
$v_1 \wedge v_2$	1	1	0	0	0	0	0	0
$v_2 \wedge v_3$	1	0	0	0	1	0	0	0
$v_1 \wedge v_3$	1	0	1	0	0	0	0	0
$v_1 \wedge v_2 \wedge v_3$	1	0	0	0	0	0	0	0

Examples.

- RM$(3,0)$ - span of v_0, a repitition code of length 8;
- RM$(3,1)$ - span of v_0, v_1, v_2, v_3, equivalent to a parity check extension of Hamming's $(7,4)$-code;
- RM$(3,2)$ - an (8.7)-code, actually simple parity check code of length 8;
- RM$(3,3) = \mathbb{F}_2^8$, trivial code.

Theorem 12.2.

(i) The vectors $v_{i_1} \wedge \ldots \wedge v_{i_s}$ where $1 \leq i_1 < i_2 < \ldots, i_s \leq d$ and $0 \leq s \leq d$ form a basis for $\mathbb{F}_2^n$;

(ii) $\text{rank}(\text{RM}(d,r)) = \sum_{s=0}^{r} \binom{d}{s}$.

Proof. In (i) we've listed $\sum_{s=0}^{d} \binom{d}{s} = (1+1)^d = 2^d = n$ vectors, so it is enough to show spanning, i.e $\text{RM}(d,d) = \mathbb{F}_2^n$. Let $p \in X$. Put

$$y_i = \begin{cases} v_i & p_i = 0 \\ v_0 + v_i & p_i = 1 \end{cases}.$$

Then $\mathbb{1}_{\{p\}} = y_1 \wedge \ldots \wedge y_d$. Expand this using the distributive law to see that $\mathbb{1}_{\{p\}} \in \text{RM}(d,d)$. Since $\{\mathbb{1}_{\{p\}}\}_{p \in X}$ clearly spans, we are done.

For (ii), RM(d,r) is spanned by $v_{i_1} \wedge \ldots \wedge v_{i_s}$ for $1 \leq i_1 < \ldots < i_s \leq d$ with $0 \leq s \leq r$. These vectors are linearly independent by (i), so a basis. Hence $\text{rank}(\text{RM}(d,r)) = \sum_{s=0}^{r} \binom{d}{s}$. □

Definition. Let C_1, C_2 be linear codes of length n with $C_2 \subseteq C_1$. The *bar product* is $C_1 \mid C_2 = \{(x \mid x+y) : x \in C_1, y \in C_2\}$. It is a linear code of length $2n$.

Lemma 12.3.

(i) $\text{rank}(C_1 \mid C_2) = \text{rank}(C_1) + \text{rank}(C_2)$;

(ii) $d(C_1 \mid C_2) = \min\{2d(C_1), d(C_2)\}$.

Proof. (i) If C_1 has basis $x_1, \ldots, x_k$, C_2 has basis $y_1, \ldots, y_l$, then $C_1 \mid C_2$ has basis $\{(x_i \mid x_i) : 1 \leq i \leq k\} \cup \{(0 \mid y_i) : 1 \leq i \leq l\}$. So $\text{rank}(C_1 \mid C_2) = k + l = \text{rank}(C_1) + \text{rank}(C_2)$;

(ii) Let $0 \neq (x \mid x+y) \in C_1 \mid C_2$. If $y \neq 0$ then $w(x \mid x+y) = w(x) + w(x+y) \geq w(y) \geq d(C_2)$. If $y = 0$ then $w(x \mid x+y) = 2w(x) \geq 2d(C_1)$. So $d(C_1 \mid C_2) \geq \min\{2d(C_1), d(C_2)\}$. But there exists $0 \neq x \in C_1$ with $w(x) = d(C_1)$, so $d(C_1 \mid C_2) \leq w(x \mid x) = 2d(C_1)$. There exists $0 \neq y \in C_2$ with $w(y) = d(C_2)$, so $d(C_1 \mid C_2) \leq w(0 \mid y) = w(y) = d(C_2)$.

□

Theorem 12.4.

(i) $\mathrm{RM}(d,r) = \mathrm{RM}(d-1,r) \mid \mathrm{RM}(r-1,r-1)$*;*

(ii) $\mathrm{RM}(d,r)$ *has minimum distance* 2^{d-r}*.*

Proof.

(i) Exercise, noting $\mathrm{RM}(d-1,r-1) \subseteq \mathrm{RM}(d-1,r)$ so the bar product is defined;

(ii) If $r = 0$ then $\mathrm{RM}(d,0)$ is a repitition code of length $n = 2^d$, having minimum distance 2^{d-0}. If $r = d$, then $\mathrm{RM}(d,d) = \mathbb{F}_2^n$ having minimum distance $1 = 2^{d-d}$. We prove the $0 < r < d$ case by induction on d. Recall from (i) that $\mathrm{RM}(d,r) = \mathrm{RM}(d-1,r) \mid \mathrm{RM}(d-1,r-1)$. By induction, the minimum distance of $\mathrm{RM}(d-1,r)$ is 2^{d-1-r} and the minimum distance of $\mathrm{RM}(d-1,r-1)$ is 2^{d-r}. Hence the minimum distance of $\mathrm{RM}(d,r)$ is $\min\{2 \cdot 2^{d-1-r}, 2^{d-r}\} = 2^{d-r}$.

□

13 Cyclic codes

GRM Recap

Fact. Let F be a field. The rings $\mathbb{Z}$ and $F[X]$ both have a division algorithm: if $a, b \in \mathbb{Z}$, $b \neq 0$, there exist $q, r \in \mathbb{Z}$ such that $a = qb + r$ and $0 \leq r < |b|$. If $f, g \in F[X]$, $g \neq 0$, there exist $q, r \in F[X]$ such that $f = qg + r$ with $\deg r < \deg g$.

Recall ideals I: subsets of a ring R such that whenever $r \in R$, $x \in I$ we have $rx \in I$ (e.g $2\mathbb{Z} \trianglelefteq \mathbb{Z}$). Principle ideals generated by $r \in R$ are $(x) = \{rx : r \in R\}$. By the division algorithm, every ideal in $\mathbb{Z}$ or $F[X]$ is principle.

If $I \trianglelefteq R$ then the set of cosets $R/I = \{x + I : x \in R\}$ is a ring, called the *quotient ring*. In particular, we identify $\mathbb{Z}/n\mathbb{Z}$ and $\{0, 1, \ldots, n-1\}$ and reduce modulo n after each operation.

Similarly, $F[X]/(f(X)) \leftrightarrow \left\{\sum_{i=0}^{n-1} a_i X^i : a_i \in F\right\} \leftrightarrow F^n$, reducing 'modulo f' after each operation using the division algorithm.

Definition. A linear code $C \subseteq \mathbb{F}_2^n$ is *cyclic* if

$$(a_0, a_1, \ldots, a_{n-1}) \in C \implies (a_{n-1}, a_0, a_1, \ldots, a_{n-2}) \in C.$$

Identify $\mathbb{F}_2[X]/(X^n - 1)$ with $\mathbb{F}_2^n$ via

$$a_0 + a_1 X + \ldots + a_{n-1} X^{-1} + (X^n - 1) \leftrightarrow^{\pi} (a_0, \ldots, a_{n-1}).$$

Lemma 13.1. *A code $C \subseteq \mathbb{F}_2^n$ is cyclic if and only if $\pi(C) = \mathcal{C}$ satisfies*

(i) $0 \in \mathcal{C}$;

(ii) $f, g \in \mathcal{C} \Rightarrow f + g \in \mathcal{C}$;

(iii) $f \in \mathbb{F}_2[X]$, $g \in \mathcal{C} \Rightarrow fg \in \mathcal{C}$.

In other words, $\mathcal{C}$ is an ideal of $\mathbb{F}_2[X]/(X^n - 1)$.

Proof. If $g(X) = a_0 + a_1 X + \ldots + a_{n-1} X^{n-1} + (X^n - 1)$ then $Xg(X) = a_{n-1} + a_0 X + \ldots + a_{n-2} X^{n-1} + (X^n - 1)$. So $\mathcal{C}$ is cyclic iff (i) and (ii) hold (linearity) and $g(X) \in \mathcal{C} \Rightarrow Xg(X) \in \mathcal{C}$. Note that closure under mulitplication by X implies closure under multiplication by any $f(X)$, so we have the result. □

From now on, we identify $\pi(C)$ and $\mathcal{C}$.

Basic problem:

$$\begin{aligned} \text{cyclic codes of length } n &\leftrightarrow \text{ ideals in } \mathbb{F}_2[X]/(X^n - 1) \\ &\leftrightarrow \text{ ideals in } \mathbb{F}_2[X] \text{ containing } (X^n - 1) \\ &\leftrightarrow \text{ polynomials } g(X) \in \mathbb{F}_2[X] \text{ with } g(X) \mid (X^n - 1) \end{aligned}$$

Theorem 13.2. *Let $C \trianglelefteq \mathbb{F}_2[X]/(X^n - 1)$ be a cyclic code. Then there exists a unique polynomial $g(X)$ such that*

(i) $C = \{f(X)g(X) + (X^n - 1) : f(X) \in \mathbb{F}_2[X]\} = (g)$;

(ii) $g(X) \mid X^n - 1$.

In particular, $p(X) \in \mathbb{F}_2[X]$ represents a codeword $g(X) \mid p(X)$. We say $g(X)$ is the generator polynomial of C.

Proof.

(i) Let $g(X) \in \mathbb{F}_2[X]$ be of least degree representing a non-zero codeword. Now $\deg g \leq n - 1$. Since C is cyclic, we have $(g) \subseteq C$.

Let $p(X) \in \mathbb{F}_2[X]$ represent a codeword. By the division algorithm, $p(X) = q(X)g(X) + r(X)$ where $\deg r < \deg g$. So $r(X) = p(X) - q(X)g(X) \in C$, contradicting minimality of $\deg g$ unless $r = 0$. Hence $g(X) \mid p(X)$ and so $C \subseteq (g)$, i.e $C = (g)$.

(ii) Take $p(X) = X^n - 1$, giving (ii).

Uniqueness: suppose $g_1(X), g_2(X)$ both satisfy (i) and (ii). Then $g_1(X) \mid g_2(X)$ and $g_2(X) \mid g_1(X)$, i.e $g_1(X), g_2(X)$ are equal up to a unit. But the only unit in $\mathbb{F}_2[X]$ is 1, so $g_1(X) = g_2(X)$. □

Lemma 13.3. *Let C be a cyclic code of length n with generator polynomial $g(X) = a_0 + a_1X + \ldots + a_kX^k$ ($a_k \neq 0$). Then C has a basis given by $g(X), Xg(X), \ldots, X^{n-k-1}g(X)$. In particular,* $\operatorname{rank}(C) = n - k$.

Proof. Trivial check. □

Corollary 13.4. *Given generator polynomial of degree k as above, a generator matrix for C is given by*

$$G = \begin{pmatrix} a_0 & a_1 & \ldots & a_k & 0 & \ldots & 0 \\ 0 & a_0 & \ldots & a_{k-1} & a_k & \ldots & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \\ 0 & 0 & \ldots & a_0 & a_1 & \ldots & a_k \end{pmatrix}$$

a $(n-k) \times n$ matrix.

Corollary 13.5. *If $g(X)h(X) = X^n - 1$ where $g(X)$ is the generator polynomial, then writing $h(X) = b_0 + b_1X + \ldots + b_{n-k}X^{n-k}$ ($b_{n-k} \neq 0$), h is called the parity check polynomial and the parity check matrix is*

$$H = \begin{pmatrix} b_{n-k} & b_{n-k-1} & \ldots & b_0 & 0 & \ldots & 0 \\ 0 & b_{n-k} & \ldots & b_1 & b_0 & \ldots & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \\ 0 & 0 & \ldots & b_{n-k} & b_{n-k-1} & \ldots & b_0 \end{pmatrix}$$

a $k \times n$ matrix.

Proof. Check that the "dot product" of the ith row of G and the jth row of H is the coefficient of $X^{(n-k-i)+j}$ in $g(X)h(X) = X^n - 1$. But $1 \leq i \leq n-k$ and $1 \leq j \leq k$ so $0 < (n-k-i)+j < n$ so all these coefficients are 0. Hence the row of G and H are "orthogonal". Also as $b_{n-k} \neq 0$, $\operatorname{rank}(H) = k = \operatorname{rank}(C^\perp)$, so H is the parity check matrix. □

Remarks. The cyclic code given by $h(X)$ is the dual code $C^\perp$ with coefficients 'reversed'. (Given a polynomial $f(X) = \sum_{i=0}^m f_iX^i$ of degree m in $\mathbb{F}_2[X]$, the reverse polynomial is $\check{f}(X) = \sum_{i=0}^m f_{m-i}X^i = X^{\deg f}f(1/X)$.)

Lemma 13.6. *If n is odd, then $X^n - 1 = f_1(X)\ldots f_t(X)$ with $f_i(X)$ distinct irreducibles in $\mathbb{F}_2[X]$ (false if n is even, e.g $X^2 - 1 = (X-1)^2$). Thus there exist 2^t cyclic codes of length n.*

Proof. Not given (see II Galois Theory). □

14 BCH codes

These are a certain type of cyclic code.

Fact (A)**.** Suppose p prime, $\mathbb{F}_p = \mathbb{Z}/p\mathbb{Z}$ is a field. If $f(X) \in \mathbb{F}_p[X]$ is irreducible, then $K = \mathbb{F}_p[X]/(f)$ is a field of order $p^{\deg f}$.

Fact (B)**.** Let $q = p^\alpha$ for p prime, $\alpha \in \mathbb{N}$. Then there exists a field $\mathbb{F}_q$ of order q and it is unique up to isomorphism.

Fact (C)**.** The multiplicative group $\mathbb{F}_q^\times = \mathbb{F}_q \setminus \{0\}$ is cyclic. A generator of this group is called a *primitive element*.