

Introduction

Example. Suppose that we have a gambler who repeatedly tosses a fair coin, betting £1 on getting a heads for each toss. Let

$$\xi_k = \begin{cases} 1 & \text{heads on } k\text{th toss} \\ -1 & \text{otherwise} \end{cases}$$

so $(\xi_k)_{k \geq 1}$ is an iid Bern(1/2) sequence. Let $X_n = \sum_{k=1}^n \xi_k$ be the net winnings of the gambler and $X_0 = 0$. Note $(X_n)_{n \geq 0}$ is a simple random walk on \mathbb{Z} , hence is a martingale (MG) with respect to $\mathcal{F}_n = \sigma(\xi_1, \dots, \xi_n)$. Suppose that at the m th toss, they bet $\mathcal{L}H_m$ on heads. Then the net winnings at time n are

$$(H \cdot X)_n = \sum_{k=1}^n H_k(X_k - X_{k-1}).$$

Assume $(H_m)_{m \geq 1}$ is deterministic. We claim $H \cdot X$ is an \mathcal{F}_n -MG. Indeed:

- (a) Integrability: obvious;
- (b) Adapted: obvious;
- (c) $\mathbb{E}[(H \cdot X)_{n+1} - (H \cdot X)_n | \mathcal{F}_n] = H_{n+1} \mathbb{E}[X_{n+1} - X_n | \mathcal{F}_n] = 0$.

More generally, the same is true if H_{n+1} is integrable and \mathcal{F}_n measurable for each n . This is called a *previsible process*. As before, $H \cdot X$ gives the winnings of the gambler. This is called a *martingale transform*.

The goal for the first part of the course: extend this reasoning to define

$$(H \cdot X)_t = \int_0^t H_s dX_s \quad (*)$$

where H is previsible and X is a continuous martingale (e.g Brownian motion).

We cannot use the Lebesgue-Stieljes integral to define $(*)$ since this requires X to have finite variation, and the only continuous martingales with finite variation are constant (see later in course). Our strategy to define the Itô integral: set

$$(H \cdot X)_t \text{ " " } \lim_{\varepsilon \rightarrow 0} \sum_{k=1}^{\lfloor t/\varepsilon \rfloor} H_{k\varepsilon} (X_{(k+1)\varepsilon} - X_{k\varepsilon}).$$

However we need to be careful about the type of limit since X in general will be rough (not differentiable), like Brownian motion. To get convergence, we need to take advantage of cancellations. For example, if X is a Brownian motion and

H is a deterministic and continuous process we have

$$\begin{aligned}
& \mathbb{E} \left[\left[\sum_{k=0}^{\lfloor t/\varepsilon \rfloor} H_{k\varepsilon} (X_{(k+1)\varepsilon} - X_{k\varepsilon}) \right]^2 \right] \\
&= \mathbb{E} \left[\sum_{k=0}^{\lfloor t/\varepsilon \rfloor} H_{k\varepsilon}^2 (X_{(k+1)\varepsilon} - X_{k\varepsilon})^2 + \sum_{j \neq k} H_{k\varepsilon} H_{j\varepsilon} (X_{(k+1)\varepsilon} - X_{k\varepsilon}) (X_{(j+1)\varepsilon} - X_{j\varepsilon}) \right] \\
&= \mathbb{E} \left[\sum_{k=0}^{\lfloor t/\varepsilon \rfloor} H_{k\varepsilon}^2 (X_{(k+1)\varepsilon} - X_{k\varepsilon})^2 \right] \\
&= \sum_{k=0}^{\lfloor t/\varepsilon \rfloor} H_{k\varepsilon}^2 \cdot \varepsilon \\
&\xrightarrow{\varepsilon \rightarrow 0} \int_0^t H_s^2 ds.
\end{aligned}$$

The cancellations that make this work come from MG orthogonality and are what makes it possible to define the Itô integral.

After this we will learn about properties of the Itô integral:

- Stochastic analogue of the chain rule;
- Stochastic analogue of integration by parts.

The formulas will look like those in regular calculus, but with an extra term to reflect that X is rough (quadratic variation). We write

$$Y_t = \int_0^t H_s dX_s \iff dY_t = H_t dX_t.$$

Itô's formula tells us how to write $df(Y_t)$ in terms of dY_t for $f \in C^2$. This has many applications, for example

Theorem (Dubins-Schwarz theorem). *Any continuous martingale is a time-change of a Brownian motion.*

Then we will look at Stochastic Differential Equations (SDEs), i.e

$$dX_t = b(t, X_t)dt + \sigma(t, X_t)dB_t$$

where b, σ are “nice” and B is a Brownian motion. For $\sigma = 0$ this is just an ODE. For $\sigma \neq 0$ this corresponds to adding noise depending on the time and state of the system.

Last part of the course: diffusion processes and how they are related to SDEs, as well as how they can be used to solve PDEs involving 2nd order elliptic operators.

0 Preliminaries

Recall that $a : [0, \infty) \rightarrow \mathbb{R}$ is *càdlàg* if it is right-continuous and has left limits. Let $a(x^-) = \lim_{y \rightarrow x^-} a(y)$ and $\Delta a(x) = a(x) - a(x^-)$. Suppose a is non-decreasing, càdlàg, $a(0) = 0$. Then there exists a unique Borel measure da on $[0, \infty)$ such that $d((s, t]) = a(t) - a(s)$ for all $0 \leq s < t$ (see Part II Probability & Measure).

For f measurable and integrable then the *Lebesgue-Stieljes* integral $f \cdot a$ is defined by

$$(f \cdot a)(t) = \int_{(0, t]} f(s) da(s) \quad \forall t \geq 0.$$

Then $(f \cdot a)$ is right-continuous. Moreover if a is continuous then $(f \cdot a)$ is continuous and so we can write

$$\int_{(0, t]} f(s) da(s) = \int_0^t f(s) da(s).$$

We want to integrate against a wider class of functions. Suppose that a^+, a^- are functions satisfying the same conditions as from before (i.e non-decreasing and càdlàg) and set $a = a^+ - a^-$. Define

$$(f \cdot a)(t) = (f \cdot a^+)(t) - (f \cdot a^-)(t)$$

for all f measurable and such that both terms on the RHS are finite. The class of functions which are a difference of càdlàg non-decreasing functions coincides with the class of càdlàg functions of *finite variation*.

Definition. Let $a : [0, \infty) \rightarrow \mathbb{R}$ be càdlàg. For each $n \in \mathbb{N}$, $t \geq 0$, let

$$v^n(t) = \sum_{k=0}^{\lceil 2^n t \rceil - 1} |a((k+1)2^{-n}) - a(k2^{-n})|. \quad (*)$$

Then the limit $v(t) := \lim_{n \rightarrow \infty} v^n(t)$ exists and is called the *total variation* of a on $(0, t]$. If $v(t) < \infty$ then we say that a has *finite variation* on $(0, t]$. If a has finite variation on $(0, t]$ for all $t \geq 0$, we say that a is of *finite variation*.

To see that $\lim_{n \rightarrow \infty} v^n(t)$ exists, fix $t > 0$ and let $t_n^+ = 2^{-n} \lceil 2^n t \rceil$, $t_n^- = 2^{-n} (\lceil 2^n t \rceil - 1)$ so that $t_n^+ \geq t \geq t_n^-$ for all n and

$$v^n(t) = \sum_{k=0}^{2^n t_n^- - 1} |a((k+1)2^{-n}) - a(k2^{-n})| + |a(t_n^+) - a(t_n^-)|.$$

The triangle inequality implies that the sum is non-decreasing in n , so converges. The càdlàg property tells us that the second term on the RHS converges to $|\Delta a(t)|$, so $v^n(t)$ does indeed converge.

Lemma. Let a be a càdlàg function of finite variation. Then v is càdlàg of finite variation with $\Delta v(t) = |\Delta a(t)|$ for all $t \geq 0$, and v is non-decreasing. In particular, if a is continuous then v is also continuous.

Proof. See Example Sheet. \square

Proposition. A càdlàg function can be written as a difference of two right-continuous non-decreasing if and only if it has finite variation.

Proof. First assume $a = a^+ - a^-$ for a^+, a^- càdlàg and non-decreasing. We show a has finite variation. Note

$$|a(t) - a(s)| \leq (a^+(t) - a^+(s)) + (a^-(t) - a^-(s)) \quad \forall 0 \leq s < t.$$

Plugging this into (*) and using the fact the sum telescopes for monotone functions to get

$$v^n(t) \leq (a^+(t_n^+) - a^+(0)) + (a^-(t_n^+) - a^-(0)).$$

Since a^+, a^- are right-continuous, the RHS converges to $(a^+(t) - a^+(0)) + (a^-(t) - a^-(0))$.

Now we show the reverse direction. Assume a has finite variation $v(t) < \infty$ for all $t > 0$. Set $a^+ = \frac{1}{2}(v + a)$ and $a^- = \frac{1}{2}(v - a)$. Then $a = a^+ - a^-$ and a^+, a^- are càdlàg since v, a are càdlàg (by the above lemma). We show a^+, a^- are non-decreasing. For $0 \leq s < t$ define t_n^+, t_n^- as before and s_n^+, s_n^- analogously. Then

$$\begin{aligned} & a^+(t) - a^+(s) \\ &= \lim_{n \rightarrow \infty} \frac{1}{2} (v^n(t) - v^n(s) + a(t) - a(s)) \\ &= \lim_{n \rightarrow \infty} \frac{1}{2} \left[\sum_{k=2^n s_n^+}^{2^n t_n^- - 1} (|a((k+1)2^{-n}) - a(k2^{-n})| + a((k+1)2^{-n}) - a(k2^{-n})) \right. \\ &\quad \left. + |a(t_n^+) - a(t_n^-)| + (a(t_n^+) - a(t_n^-)) \right] \\ &\geq 0. \end{aligned}$$

The same argument works for a^- . \square

Random integrators: now we discuss integration against random functions of finite variations. Let $(\Omega, \mathcal{F}, (\mathcal{F}_t)_{t \geq 0}, \mathbb{P})$ be a filtered probability space. Recall a stochastic process $X : \Omega \times [0, \infty) \rightarrow \mathbb{R}$ is *adapted* if $X_t = X(\cdot, t)$ is \mathcal{F}_t -measurable for all $t \geq 0$. We also say X is *càdlàg* if $X(\omega, \cdot)$ is càdlàg for all $\omega \in \Omega$.

Definition. Given a càdlàg adapted process $A : \Omega \times [0, \infty) \rightarrow \mathbb{R}$, its *total variation process* $V : \Omega \times [0, \infty) \rightarrow \mathbb{R}$ is defined pathwise by setting $V(\omega, \cdot)$ to be the total variation of $A(\omega, \cdot)$.

Lemma. *If A is càdlàg, adapted and of finite variation, then V is càdlàg, adapted and non-decreasing.*

Proof. We just need to show V is adapted (the rest follows by previous results). For $t \geq 0$, set as before $t_n^- = 2^{-n}(\lceil 2^n t \rceil - 1)$. Then define

$$\tilde{V}_t^n = \sum_{k=0}^{2^n t_n^- - 1} |A_{(k+1)2^{-n}} - A_{k2^{-n}}|$$

so \tilde{V}^n is adapted for all n as $t_n^- \leq t$. Then

$$V_t = \lim_{n \rightarrow \infty} \tilde{V}_t^n + |\Delta A(t)|$$

is \mathcal{F}_t -measurable as a limit/sum of \mathcal{F}_t -measurable functions. □

Recall that a discrete time process $(H_n)_{n \geq 0}$ is *previsible* with respect to $(\mathcal{F}_n)_{n \geq 0}$ if H_{n+1} is \mathcal{F}_n -measurable for all $n \geq 0$.

Definition. The *previsible* σ -algebra \mathcal{P} on $\Omega \times (0, \infty)$ is generated by sets of the form $E \times (s, t]$ for $E \in \mathcal{F}_s$ and $s < t$. A process $H : \Omega \times (0, \infty) \rightarrow \mathbb{R}$ is *previsible* if it is \mathcal{P} -measurable.

Examples.

1. $H(\omega, t) = Z(\omega) \mathbb{1}_{(t_1, t_2]}(t)$ for $t_1 < t_2$ and Z being \mathcal{F}_{t_1} -measurable;
2. $H(\omega, t) = \sum_{k=0}^{n-1} Z_k(\omega) \mathbb{1}_{(t_k, t_{k+1}]}(t)$ for $0 = t_0 < \dots < t_n$ and Z_k \mathcal{F}_{t_k} -measurable. H of this form is called a *simple process* and will be important for constructing the Itô integral.

Remark. Simple processes are left-continuous and adapted. It turns out that \mathcal{P} is the smallest σ -algebra on $\Omega \times (0, \infty)$ such that all left-continuous adapted processes are measurable.

In general, a càdlàg process is not previsible, but their left-continuous modification is.

Proposition. Let X be a càdlàg adapted process and let $H_t = X_{t-}$, $t \geq 0$. Then H is previsible.

Proof. Since X is càdlàg and adapted, it is clear that H is left-continuous and adapted. For each n set

$$H_t^n = \sum_{k=0}^{\infty} H_{k2^{-n}} \mathbb{1}_{(k2^{-n}, (k+1)2^{-n}]}(t).$$

Then H_t^n is previsible for all n . By left continuity of H we have $\lim_{n \rightarrow \infty} H_t^n = H_t$ for all t . So H is previsible as the limit of previsible functions. \square

Remark. The above proposition shows that continuous and adapted processes are previsible.

Proposition. If H is previsible then H_t is $\sigma(\mathcal{F}_s : s < t) = \mathcal{F}_{t-}$ -measurable for all t .

Proof. See Example Sheet. \square

Remark. The Poisson process $(N_t)_{t \geq 0}$ is not previsible since N_t is not \mathcal{F}_{t-} -measurable for $(\mathcal{F}_t)_{t \geq 0}$ the natural filtration for N .

We will not show that integrating a previsible process against a càdlàg process which is adapted and has finite variation yields an adapted càdlàg process of finite variation.

Theorem. Let $A : \Omega \times (0, \infty) \rightarrow \mathbb{R}$ be a càdlàg process which is adapted and has finite variation V . Let H be a previsible process with

$$\int_{(0,t]} |H(\omega, s)| dV(s) < \infty \quad \forall t > 0, \omega \in \Omega. \quad (1)$$

Then the process $H \cdot A : \Omega \times [0, \infty) \rightarrow \mathbb{R}$ given by

$$(H \cdot A)(\omega, t) = \int_{(0,t]} H(\omega, s) dA(\omega, s), \quad (H \cdot A)(\omega, 0) = 0 \quad (2)$$

is càdlàg adapted and of finite variation.

Proof. The integral in (2) is well-defined due to (1). Indeed, let H^+, H^- be the positive/negative parts of H respectively and let $A^\pm = \frac{1}{2}(V \pm A)$. Then $H = H^+ - H^-$, $A = A^+ - A^-$ and

$$(H \cdot A) = (H^+ - H^-) \cdot (A^+ - A^-) = H^+ \cdot A^+ - H^- \cdot A^+ - H^+ \cdot A^- + H^- \cdot A^-$$

and all terms on the RHS are finite by assumption (1).

We need to show $H \cdot A$ is (1) càdlàg, (2) adapted and (3) of finite variation.

Step 1: note $\mathbb{1}_{(0,s]} \rightarrow \mathbb{1}_{(0,t]}$ as $s \downarrow t$ and $\mathbb{1}_{(0,s]} \rightarrow \mathbb{1}_{(0,t)}$ as $s \uparrow t$. By definition $(H \cdot A)_t = \int H_s \mathbb{1}(s \in (0, t]) dA_s$ so

$$\begin{aligned} (H \cdot A)_t &= \int H_s \lim_{r \downarrow t} \mathbb{1}(s \in (0, r]) dA_s \\ &= \lim_{r \downarrow t} \int H_s \mathbb{1}(s \in (0, r]) dA_s \\ &= \lim_{r \downarrow t} (H \cdot A)_r \end{aligned} \quad (\text{DCT})$$

so $H \cdot A$ is right-continuous. An analogous argument shows $H \cdot A$ has left-limits, so is càdlàg. Also $\Delta(H \cdot A)_t = \int H_s \mathbb{1}(s = t) dA_s = H_t \Delta A_s$.

Step 2: we'll use a "monotone class" style argument. Suppose $H = \mathbb{1}_{B \times (s, u]}$ where $B \in \mathcal{F}_s$ and $s < u$. Then $(H \cdot A)_t = \mathbb{1}_B(A_{t \wedge u} - A_{t \wedge s})$ which is \mathcal{F}_t -measurable. Let $\mathcal{A} = \{C \in \mathcal{P} : \mathbb{1}_C \cdot A \text{ is adapted}\}$. We want to show $\mathcal{A} = \mathcal{P}$. Let $\Pi = \{B \times (s, u] : B \in \mathcal{F}_s, s < u\}$ so $\Pi \subseteq \mathcal{A}$ and Π is a π -system generating \mathcal{P} by definition. Not difficult to see that \mathcal{A} is a d -system, implying $\mathcal{A} = \mathcal{P}$ by Dynkin's lemma.

Now suppose $H \geq 0$ is previsible. Set

$$\begin{aligned} H_n &= (2^{-n} \lfloor 2^n H \rfloor) \wedge n \\ &= \sum_{k=0}^{2^n-1} 2^{-nk} \underbrace{\mathbb{1}(H \in [2^{-n}k, 2^{-n}(k+1)))}_{\in \mathcal{P}} + \underbrace{\mathbb{1}(H \geq n)}_{\in \mathcal{P}} \end{aligned}$$

so H_n is a finite linea combination of functions of the form $\mathbb{1}_C$ for $C \in \mathcal{P}$. Thus $(H^n \cdot A)$ is adapted for all n . By the MCT $(H^n \cdot A)_t \rightarrow (H \cdot A)_t$ so $H \cdot A$ is itself adapted. For general previsible H we write $H = H^+ - H^-$ as usual.

Step 3: we have

$$H \cdot A = (H^+ \cdot A^+ + H^- \cdot A^-) - (H^- \cdot A^+ + H^+ \cdot A^-)$$

which is a difference of non-decreasing functions. □

Local Martingales

Let $(\Omega, \mathcal{F}, (\mathcal{F}_t)_{t \geq 0}, \mathbb{P})$ be a filtered probability space.

Definition. Say that $(\mathcal{F}_t)_{t \geq 0}$ satisfies the *usual conditions* if

- \mathcal{F}_0 contains all \mathbb{P} -null sets;
- $(\mathcal{F}_t)_{t \geq 0}$ is right-continuous, i.e $\mathcal{F}_t = \mathcal{F}_{t+} = \bigcap_{s > t} \mathcal{F}_s$ for all $t \geq 0$.

Throughout we assume that (\mathcal{F}_t) satisfies the usual conditions.

For T a stopping time, set $\mathcal{F}_T = \{E \in \mathcal{F} : E \cap \{t \leq T\} \in \mathcal{F}_t \forall t \geq 0\}$. Then X_T is \mathcal{F}_T -measurable. If X is a martingale then $X^T = X_{T \wedge t}$ is also a martingale. Recall:

Theorem. *Optional stopping theorem* Let X be an adapted, càdlàg, integrable process. Then the following are equivalent

1. X is a martingale;
2. X^T is a martingale for all stopping times T ;
3. For all bounded stopping times $S \leq T$, we have

$$\mathbb{E}[X_T | \mathcal{F}_S] = X_S \text{ almost-surely};$$

4. For all bounded stopping times T , we have that

$$\mathbb{E}[X_T] = \mathbb{E}[X_0].$$

Definition. A càdlàg adapted process X is called a *local martingale* if there exists a sequence $(T_n)_{n \geq 1}$ of stopping times with $T_n \uparrow \infty$ almost-surely such that the stopped process X^{T_n} is a martingale for all $n \geq 1$. In this case, we say that $(T_n)_{n \geq 1}$ *reduces* X .

Note that a martingale is always a local martingale as any deterministic sequence $T_n \uparrow \infty$ will reduce it.

Example. Let B be a standard Brownian motion in \mathbb{R}^3 and let $M_t = \frac{1}{|B_t|}$. In Example Sheet 4 of Part III Advanced Probability we have seen that

- (i) M is bounded in L^2 ;
- (ii) $\mathbb{E}M_t \rightarrow 0$ as $t \rightarrow \infty$;
- (iii) M is a supermartingale.

M cannot be a martingale as otherwise its expectation would vanish by (ii). Now we will show that M is a local martingale. For each $n \geq 1$ set $T_n = \inf\{t \geq 1 : |B_t| < 1/n\} = \inf\{t \geq 1 : |M_t| > n\}$. We want to show:

- (i) $(M_t^{T_n})_{t \geq 1}$ is a martingale for all n ;
- (ii) $T_n \uparrow \infty$ as $n \rightarrow \infty$ almost-surely.

Note that $n \leq M_1$ implies $T_n = 1$ and $n > M_1$ implies $T_n > 1$. Since $|B_t|$ cannot hit $1/n$ before hitting $1/(n+1)$, we see T_n is non-decreasing.

In Advanced Probability we saw that for $f \in C_b^2(\mathbb{R}^3)$ (C^2 with bounded derivatives) we have

$$f(B_t) - f(B_0) - \frac{1}{2} \int_0^t \Delta f(B_s) ds$$

is a martingale. Note that $f(x) = 1/|x|$ is harmonic in $\mathbb{R}^3 \setminus \{0\}$. Let $(f^n)_{n \geq 1}$ be a sequence of $C_b^2(\mathbb{R}^3)$ functions with $f^n(x) = 1/|x|$ on $\{|x| \geq 1/n\}$. If $0 < |B_1| < 1/n$ then $T_n = 1$ and $M_t^{T_n} = M_1$ is a martingale. Since $B_1 \neq 0$ almost-surely, we have $|B_1| > 1/n$ for all n sufficiently large in which case $f(B_{t \wedge T_n}) = f^n(B_{t \wedge T_n})$. Thus

$$\begin{aligned} M_{t \wedge T_n} &= f(B_{t \wedge T_n}) - f(B_1) + f(B_1) \\ &= \left(f(B_{t \wedge T_n}) - f(B_1) - \frac{1}{2} \int_1^{t \wedge T_n} \Delta f(B_s) ds \right) + f(B_1) \\ &= \left(\underbrace{f^n(B_{t \wedge T_n}) - f^n(B_1) - \frac{1}{2} \int_1^{t \wedge T_n} \Delta f^n(B_s) ds}_{\text{martingale}} \right) + f^n(B_1) \end{aligned}$$

so $M^{T_n} = (M_{t \wedge T_n})_{t \geq 1}$ is a martingale. Now we show $T_n \uparrow \infty$ almost-surely as $n \rightarrow \infty$. Since $T_n \leq T_{n+1}$ it suffices to show $T_n \rightarrow \infty$. For each R let $S_R = \inf\{t \geq 1 : |B_t| > R\} = \inf\{t \geq 1 : M_t < 1/R\}$. Then $S_R \rightarrow \infty$ as $R \rightarrow \infty$. We have

$$\begin{aligned} \mathbb{P}(\lim_n T_n < \infty) &\leq \mathbb{P}(\exists R : T_n < S_R \ \forall n) \\ &= \lim_{R \rightarrow \infty} \lim_{n \rightarrow \infty} \mathbb{P}(T_n < S_R). \end{aligned}$$

The OST says that $\mathbb{E}[M_{T_n \wedge S_R}] = \mathbb{E}[M_1] := \mu \in (0, \infty)$. Also

$$\begin{aligned} \mathbb{E}[M_{T_n \wedge S_R}] &= n\mathbb{P}(T_n < S_R) + \frac{1}{R}\mathbb{P}(S_R \leq T_n) \\ &= n\mathbb{P}(T_n < S_R) + \frac{1}{R}(1 - \mathbb{P}(T_n < S_R)) \\ &= \mu \end{aligned}$$

so $\mathbb{P}(T_n < S_R) = \frac{n-1/R}{n-1/R} \rightarrow 0$ as $n \rightarrow \infty$. Therefore M is a non-negative local martingale but not a martingale. It is also a super martingale and bounded in L^2 .

We actually have:

Proposition. If X is a local martingale and X is non-negative then X is a supermartingale.

Proof. Let (T_n) be a reducing sequence for X . Then for any $s \leq t$ we have that

$$\begin{aligned}\mathbb{E}[X_t|\mathcal{F}_s] &= \mathbb{E}[\lim_n X_{t \wedge T_n}|\mathcal{F}_s] \\ &\leq \liminf_{n \rightarrow \infty} \mathbb{E}[X_{t \wedge T_n}|\mathcal{F}_s] && \text{(Fatou)} \\ &= \lim_{n \rightarrow \infty} X_{s \wedge T_n} \\ &= X_s \text{ almost-surely.}\end{aligned}$$

□

We often work with local martingales instead of martingales because we want to avoid having to worry about integrability.